Structural Properties and Signed Boundaries of 2-Gaps in Sieve Sequences

Thiago Henrique Ramos da Mata
Independent Researcher
thiago.henrique.mata@gmail.com
ORCID: 0009-0002-7366-939X
https://thiagomata.com/
https://github.com/thiagomata
DOI: https://doi.org/10.5281/zenodo.22955786

Proof status: The sequence foundation is Stainless-verified in the companion Sieve Sequence article; the theorems introduced here are proved mathematically in this article.

Abstract

This article studies how 2-gaps evolve under successive prime filters and sharpens the boundary between complete-period survival and square-window placement. Complete-period CRT arguments prove that 2-gaps persist globally: an incoming odd prime removes exactly two copy classes of each old 2-gap. Those counts do not force a survivor into a particular interval $[Q, Q^2)$.

The full argument reduces survival to a single signed quantity. A weighted conservation law and its Cauchy–Schwarz corollary give a sharp terminal threshold on the harmful-excess energy. A closed exhaustion argument—capacity envelopes, native-period Bessel, fixed and moving cuts, and the stability-gap repair—shows that no unsigned-capacity route can clear that threshold. The subsequent analysis is signed and local. At filter 7, exact residue order gives the sharp interval bound $|b_7| \le 18/7$, replacing a capacity estimate that grows quadratically with the window scale. More generally, when an old period is copied through an incoming prime r, the centered harmful excess in copy block j is exactly $B_j = d_t + d_{t-2}$ for two entries of the centered old-start histogram modulo r. Consequently,

$$\sum_{j=0}^{r-1} B_j^2 = 2V_r + 2 \sum_{t \bmod r} d_t d_{t-2} \le 4V_r.$$

This composes residue energy with the weighted harmful-excess survival criterion over complete old-period blocks. It does not control the two partial boundary fragments of an arbitrary interval, nor does it prove a relative residue-energy estimate across an unbounded filter chain. Those are now the precise remaining arithmetic obligations, framed by a scale conflict that limits fixed-seed averaging and a Type-II barrier that any prime-producing sieve must overcome.

Licensed under the Creative Commons Attribution 4.0 International License (CC BY 4.0):
https://creativecommons.org/licenses/by/4.0/.

1 Introduction

The Sieve Sequence represents an infinite accepted-value stream using one finite cyclic gap list. That representation makes complete-period identities exact, but the twin-prime application asks for a

surviving 2-gap start in the eligible square-safe window of one future head. The distinction between those scopes is the organizing principle of this article.

Version 1 established the copy-or-merge and complete-period/local-window boundary. This version retains that theorem spine and adds the signed localization results that emerged from the later quadratic investigation. The article develops the following properties:

1. exact complete-period 2-gap count — Section 3.1;
2. exact copy-index filter frequency — Section 3.2;
3. exact batched survival — Section 3.3;
4. exact complete-period `(2,4,2)` cluster count — Section 3.4;
5. rotation invariance of cyclic counts — Section 3.5;
6. stable global absence of 2-gaps — Section 3.6;
7. square-safe certification — Section 4.1;
8. post-filter-3 isolation — Section 4.2;
9. exact accepted strikes and the sharp one-transition threshold — Sections 4.3–4.4;
10. why the capacity envelope is exhausted — Section 6;
11. exact weighted deletion conservation and its terminal quadratic corollary — Sections 5.1–5.2;
12. the exact filter-7 interval saving — Section 7;
13. the live frontier (accepted-boundary discrepancy and residue-collision energy estimates) — Section 8;
14. the copy-block residue-energy bridge — Section 9;
15. the classified routes and the live program — Section 10; and
16. the fixed-seed scale conflict (Section 10.1) and the Type-II barrier (Section 11.1).

The complete-period counts and gap-cycle mechanisms have earlier sources in Holt and Rudd [2]: their Section 1.2 gives the two-gap product, Lemma 2.1 gives cycle expansion and closure, and Lemma 3.1 treats intact constellation copies. We rederive these established ingredients in the present notation. Their primorial includes the last installed prime; our stage is indexed by the next head, so the cyclic counts agree after this change of index and origin.

The weighted identities and energy bounds apply standard telescoping, Cauchy–Schwarz and cyclic reindexing to the stated sieve observables. Historical priority of the particular signed/local formulations and sharp fixed-wheel calculation is not established here. A proof given in this article does not by itself make its conclusion a new mathematical discovery.

The final sections separate the proved complete-block control from the open partial-boundary and cross-layer estimates. No complete-period theorem is used as a short-window lower bound.

2 Preliminaries And Evidence Boundary

For a prime stage head p, let

$$M_p = \prod_{q<p} q$$

and retain the integers coprime to M_p. Their accepted residues repeat modulo M_p; adjacent differences in one complete period form a cyclic gap list. The verified companion article establishes accepted-value completeness, strict increase, period shift, exact survivor count, and the copy-or-merge transition rule.

This article takes those verified construction facts as inputs. It then studies mathematical properties of 2-gap populations and signed interval discrepancies. The division is deliberate: the companion article fixes and machine-checks the sieve object; this article expands the mathematical theory of that object without claiming that every result proved here has a Scala encoding.

2.1 Roles, Populations, And Scopes

The notation follows the shared research vocabulary:

- Q is a fixed **future head** used for square-safe certification;
- r is one **incoming prime**, and r_i is filter i in a complete conditioned chain to Q;
- a **complete period** contains one full cyclic sieve pattern;
- a **local window** is an explicitly stated bounded integer interval;
- a **2-gap-start population** counts starts x, not all accepted values; and
- a **square-safe certificate** requires both x and $x + 2$ to lie strictly below Q^2 after every missing prime below Q has been installed.

For a conditioned chain, S_i denotes the local 2-gap starts immediately before filter r_i, and $N_i = |S_i|$. The final survivor population is S_m with size N_m. Complete-period populations are named separately and are never substituted for N_i without an explicit localization theorem.

2.2 Evidence Status

The following status convention is used:

- **Verified foundation:** supported by maintained Stainless source through the companion article.
- **Open:** the required estimate is stated but not proved.

Every other result proved in this article is proved mathematically in the text itself; the derivations and appendices are the authority for the article's mathematical claims.

No theorem in this article claims infinitely many twin primes.

Each property section states its population, scope, and quantifier, then proves the result. Verified construction inputs link to maintained Scala contracts.

3 Complete-Period 2-Gap Properties

The first property group concerns cyclic 2-gap starts in complete periods. These identities are exact because every installed prime contributes a full residue system. They are not statements about a local square window.

3.1 Exact Global 2-Gap Count

The product count below already appears in Holt and Rudd's Section 1.2 [2]; their Corollary 6.3 also uses its one-filter recurrence. The CRT proof here recalls that established complete-period result.

After filter 2 is installed, every accepted value is odd, so accepted endpoints x and $x+2$ are consecutive survivors: the intermediate value is even. For every prime stage after filter 2, this property counts all such cyclic 2-gap starts in one complete period of that stage, without constructing the gap list.

Let p be the stage head and

$$M_p = \prod_{q<p} q,$$

where q ranges over installed primes. A residue x represents a cyclic 2-gap exactly when

$$\gcd(x(x+2), M_p) = 1.$$

The exact count is

$$G_2(p) = \prod_{\substack{3\le q<p \\ q \text{ prime}}} (q-2).$$

For the filter 2, the start must be odd, leaving one allowed residue. For each installed odd prime q, the two endpoints fail precisely in the two classes

$$x \equiv 0 \pmod{q}, \qquad x \equiv -2 \pmod{q}.$$

They are distinct because q is odd, so exactly $q-2$ residues remain allowed at each installed odd prime q.

The installed primes are pairwise coprime, so CRT gives a bijection between one allowed choice at every installed prime and one residue modulo M_p. Consequently,

$$\begin{aligned} G_2(p) &= 1 \cdot \prod_{\substack{3\le q<p \\ q \text{ prime}}} (q-2) \quad \text{[Chinese Remainder Theorem]} \\ &= \prod_{\substack{3\le q<p \\ q \text{ prime}}} (q-2). \qquad \blacksquare \text{ [Q.E.D.]} \end{aligned}$$

The empty odd-prime product is 1, so the statement includes the first odd stage. The result proves complete-period global presence, not placement in any specified local window.
The derivation above proves the exact product count.

3.2 Exact Filter Frequency Across Repeated Copies

Expansion distributes every old 2-gap through equally spaced copies before filtering, and one incoming prime cannot choose arbitrary copies: the two endpoint strikes occur at two exact copy-index phases. This holds for lifted copies of one fixed old cyclic 2-gap, across every complete or finite run of copy indices satisfying the stated coprimality precondition.

Let M be the old period, let $(a, a+2)$ be one old cyclic 2-gap, and let $r > 2$ be an incoming prime with $\gcd(M, r) = 1$. Copy j has endpoints

$$a + jM, \qquad a + 2 + jM.$$

Let M^{-1} denote the inverse of M modulo r. The left endpoint is deleted exactly when

$$\begin{aligned} a + jM &\equiv 0 \pmod{r} && \text{[Left Endpoint Strike]} \\ j &\equiv -aM^{-1} \pmod{r}. && \text{[Multiply By } M^{-1}\text{]} \end{aligned}$$

Likewise, the right endpoint is deleted exactly when

$$\begin{aligned} a + 2 + jM &\equiv 0 \pmod{r} && \text{[Right Endpoint Strike]} \\ j &\equiv -(a+2)M^{-1} \pmod{r}. && \text{[Multiply By } M^{-1}\text{]} \end{aligned}$$

If the two copy-index classes were equal, subtracting would give $2M^{-1} \equiv 0 \pmod{r}$, hence $2 \equiv 0 \pmod{r}$, impossible for $r > 2$. Therefore the classes are distinct.

Every residue class modulo r occurs at most $\lceil L/r \rceil$ times in a run of L consecutive copy indices. Thus

$$K_r(J) \le 2 \left\lceil \frac{L}{r} \right\rceil.$$

For one complete run of r consecutive copy indices, each forbidden class occurs exactly once. Hence

$$K_r = 2, \qquad G_{\text{survive}} = r - 2.$$

■ [Q.E.D.]

This is exact distribution across copies before and under one filter. It does not say that the $r-2$ survivors are evenly distributed after filtering or that one lies in a chosen numerical window. The derivation above proves this exact copy-index law.

3.3 Exact Batched 2-Gap Survival

The one-filter copy law composes exactly. Applying several future filters as one batch counts intersections through CRT, so no survivor is subtracted twice and no intermediate floor or density approximation is needed. This holds for every finite set of distinct incoming odd primes coprime to the old period, over one complete period after the whole batch, for lifted copies of every old cyclic 2-gap in one complete combined period.

Let M be the old period and let $(a, a+2)$ be one old cyclic 2-gap. Choose a finite set of distinct odd primes

$$\mathcal{R} = \{r_1, r_2, \ldots, r_k\}, \qquad \gcd\left(M, \prod_{r \in \mathcal{R}} r\right) = 1,$$

and define

$$R = \prod_{r \in \mathcal{R}} r.$$

The complete combined period modulo MR contains R lifted copies of the old pair. Exactly

$$\prod_{r \in \mathcal{R}} (r - 2)$$

survive every filter in the batch. Therefore, if the old complete period has G cyclic 2-gaps, the new complete period has

$$G_{\text{after}} = G \prod_{r \in \mathcal{R}} (r - 2).$$

For each $r \in \mathcal{R}$, the filter-frequency theorem (Section 3.2) identifies two distinct forbidden copy-index classes: one where r divides the left endpoint and one where it divides the right endpoint. Hence there are exactly $r - 2$ allowed choices modulo r. Because the primes in $\mathcal{R}$ are pairwise coprime, CRT makes these choices independent. Thus, for one old 2-gap,

$$G_{\text{one old gap}} = \prod_{r \in \mathcal{R}} (r - 2). \quad \text{[Chinese Remainder Theorem; Copy-Index Filter Frequency]}$$

The lifted copy sets belonging to distinct old cyclic starts are counted as distinct starts in the new complete period. Summing the same exact count over the G old starts gives

$$\begin{aligned} G_{\text{after}} &= \sum_{g=1}^{G} \prod_{r \in \mathcal{R}} (r - 2) \quad \text{[Sum Over Old 2-Gap Starts]} \\ &= G \prod_{r \in \mathcal{R}} (r - 2). \quad \blacksquare \text{ [Simplification; Q.E.D.]} \end{aligned}$$

The formula is independent of the conceptual order of the filters and counts overlapping strikes correctly. Its scope is nevertheless the full combined period MR. A shorter interval can omit all allowed CRT classes, so this theorem cannot place a survivor in an eligible square-safe window. The derivation above proves the finite-batch formula and its complete-period limitation.

3.4 Exact Global `(2,4,2)` Cluster Count

Holt and Rudd [2], Lemma 3.1, give the intact-copy mechanism for a constellation of length three and explicitly discuss `(2,4,2)`. The four closures below are that known mechanism. The total-count recurrence also requires the no-new-occurrence argument supplied below.

The label $(2, 4, 2)$ names the sequence of consecutive gaps inside a four-term cluster

$$a, \qquad a + 2, \qquad a + 6, \qquad a + 8 :$$

gap 2 from a to $a + 2$, gap 4 from $a + 2$ to $a + 6$, gap 2 again from $a + 6$ to $a + 8$. This is the classic prime-quadruplet shape, realized for example by $11, 13, 17, 19$ or $101, 103, 107, 109$: two twin-prime pairs sharing their two middle values four apart. Counting how often this exact pattern survives is the natural next question once single 2-gaps (Section 3.1) are understood, since it asks whether 2-gaps keep arriving in tight pairs rather than only individually. It is also where the pattern comes from in the first place: at the earliest stage, right after filters 2 and 3 are installed, the entire accepted sequence already repeats as $\ldots, 2, 4, 2, 4, 2, \ldots$ forever, so $(2, 4, 2)$ starts out as the universal shape of the sequence, not a rare event — this section tracks how much of that shape survives as larger primes filter it down.

The cluster's two 2-gaps are endpoint-disjoint and lie in a total span of 8. Expansion creates r copies of every old cluster; the incoming filter strikes exactly four copies and preserves the other $r - 4$ intact.

Let C_M be the cyclic cluster count at modulus M. Copy j of a cluster has endpoints

$$a + jM + \{0, 2, 6, 8\}, \qquad 0 \le j < r.$$

For each offset $h \in \{0, 2, 6, 8\}$, the endpoint $a + jM + h$ is removed exactly in the copy-index class

$$j \equiv -(a + h)M^{-1} \pmod{r}.$$

If two offsets produced the same class, their difference would be divisible by r. The nonzero pairwise differences are $2, 4, 6, 8$. Neither 5 nor 7 divides any applicable difference, and every prime $r \geq 11$ exceeds all four. Thus the four classes are distinct. Out of the r total copies created by expansion, let N_{struck} be the number of copies struck by the incoming filter (one per distinct endpoint class) and N_{intact} the number left intact. Then

$$\begin{aligned} N_{\text{struck}} &= 4 && \text{[Four Distinct Endpoint Classes]} \\ N_{\text{intact}} &= r - N_{\text{struck}} = r - 4. && \text{[Subtraction]} \end{aligned}$$

It remains to exclude newly created clusters. After filter 2, all old gaps are positive and even. A merged gap is a sum of at least two old gaps, so it cannot equal 2. It can equal 4 only as $2 + 2$. But two consecutive 2-gaps would require all three values $x, x + 2, x + 4$ to be accepted, whereas one of them is divisible by 3. Since filter 3 is installed, that is impossible. Therefore every new gap of size 2 or 4 is a copied old gap, and every new $(2, 4, 2)$ occurrence is one of the intact copies counted above. Hence

$$C_{rM} = (r - 4)C_M.$$

■ [No New Occurrences; Q.E.D.]

This recurrence holds for every complete-period stage with modulus M divisible by 6 and every incoming prime $r \geq 5$ with $r \nmid M$, counting cyclic occurrences of the gap word $(2, 4, 2)$ in one complete period after filters 2 and 3 are installed.

The wheel modulo 6 has cyclic gap word $(4, 2)$: accepted residues $1, 5 \pmod 6$ repeat forever as $\ldots, 4, 2, 4, 2, 4, 2, \ldots$ before any prime beyond 3 is installed. At this base stage the sequence is *nothing but* copies of $(2, 4, 2)$ back to back, so there is exactly one cyclic $(2, 4, 2)$ occurrence per period, $C_6 = 1$ — the recurrence below tracks how that initially-universal pattern gets diluted, but never destroyed, as larger primes are installed. Iterating the recurrence over a finite installed-prime set $\mathcal{P}$ containing 2 and 3 gives

$$C(\mathcal{P}) = \prod_{\substack{p \in \mathcal{P} \\ p \geq 5}} (p - 4).$$

The absolute cluster population is positive at every stage and grows whenever an incoming prime exceeds 5. Its proportion among accepted positions is multiplied by $(r - 4)/(r - 1) < 1$, so global growth does not imply that a chosen short window contains a cluster.
The derivation above establishes the recurrence, no-creation result, closed product, and localization boundary.

3.5 Rotation Preserves Cyclic Gap Counts

Rotation chooses a new origin for the same cyclic list. It neither filters an accepted value nor merges adjacent gaps, so it preserves the number of entries having every gap value, including 2, for every nonempty finite cyclic gap list, every nonnegative rotation offset, and every gap value d. The rotation invariance itself is proved.

Let

$$G = (g_0, g_1, \ldots, g_{T-1}), \qquad T \geq 1,$$

and define rotation by offset j through

$$\mathrm{rot}_j(G)_i = g_{(i+j) \bmod T}.$$

For every value d,

$$\#\{i : \mathrm{rot}_j(G)_i = d\} = \#\{i : g_i = d\}.$$

Define the index map $\varphi_j(i) = (i + j) \bmod T$. Addition by j modulo T is invertible, with inverse $\varphi_{-j}(i) = (i - j) \bmod T$. Hence φ_j is a bijection of $\{0, 1, \ldots, T - 1\}$. Therefore

$$\begin{aligned}
\#\{i : \mathrm{rot}_j(G)_i = d\} &= \#\{i : g_{\varphi_j(i)} = d\} && \text{[By Definition Of Rotation]} \\
&= \#\{k : g_k = d\} && \text{[Reindex By The Bijection } k = \varphi_j(i)\text{]} \\
&= \#\{i : g_i = d\}. && \blacksquare \text{ [Rename The Bound Index; Q.E.D.]}
\end{aligned}$$

Taking $d = 2$ proves exact invariance of the complete cyclic 2-gap count. Rotation may change which gap follows the displayed head and whether one gap crosses the end-to-start boundary of a linear rendering. It cannot imply that an absolute coordinate window such as $[Q, Q^2)$ contains the same number of 2-gaps, because that window is not defined only by cyclic indices.

Scala Verification Foundation And Pending Multiplicity Theorem

The maintained list implementation defines rotation as `back ++ front`. Its verified properties preserve membership in both directions and preserve list size:

```
def assertRotateContainsForward(
  list: List[BigInt], index: BigInt, x: BigInt
): Boolean = {
  require(index >= 0)
  require(list.contains(x))
  // ... verified proof ...
  ListUtils.rotateAt(list, index).contains(x)
}.holds

def assertRotateContainsBackward(
  list: List[BigInt], index: BigInt, x: BigInt
): Boolean = {
  require(index >= 0)
  require(ListUtils.rotateAt(list, index).contains(x))
  // ... verified proof ...
  list.contains(x)
}.holds

def assertRotateSameSize(
  list: List[BigInt], index: BigInt
): Boolean = {
  require(index >= 0)
  // ... verified proof ...
  ListUtils.rotateAt(list, index).size == list.size
}.holds
```

These foundations are verified in RotationProperties::assertRotateContainsForward, assertRotateContainsBackward, and assertRotateSameSize. They establish the rotation operation and its same-elements/size behavior, but membership alone does not count duplicate entries. The exact multiplicity theorem above is

3.6 Absence Of 2-Gaps Is Stable

Later filtering can copy an old gap or merge consecutive old gaps, but it cannot create a smaller positive gap. Thus, once the complete cyclic population contains no 2-gap, no later filter can recreate one — a fact that holds for every later filter transition, and therefore every finite chain of later transitions, over every gap in one complete cyclic post-filter-2 gap list.

Let the old cyclic gaps be $g_0, \ldots, g_{T-1}$. Because filter 2 is already installed, every g_i is positive and even. Under the no-2 hypothesis,

$$\forall i, \qquad g_i \neq 2 \quad \Longrightarrow \quad g_i \geq 4.$$

For two values that remain consecutive after the next filter, their new gap h has one of two forms. If no intermediate accepted value was removed, it copies one old gap. If one or more intermediate accepted values were removed, it is the sum of at least two consecutive old gaps. Therefore

$$h = g_i \Longrightarrow h \geq 4 \qquad \text{[Copied Gap]}$$

$$h = \sum_{j=u}^{v} g_j, \quad v \geq u+1 \Longrightarrow h \geq 4+4 = 8. \quad \text{[Merged Gaps]}$$

In either case $h \neq 2$, so

$$(\forall i,\ g_i \neq 2) \Longrightarrow (\forall j,\ h_j \neq 2)\,.$$

■ [Copy-Or-Merge Exhaustion; Q.E.D.]

Applying the same implication inductively yields

$$G_s(2) = 0 \Longrightarrow G_t(2) = 0 \qquad \text{for every later stage } t \geq s,$$

where $G_s(2)$ is the complete-period 2-gap count at stage s. This is a one-way extinction theorem. A positive global count does not force a 2-gap into a chosen short window, so it must not be used as a localization result.

The copy-or-merge argument above proves the inductive consequence and its global/local boundary.

These are complete-period statements. They explain global growth but do not locate any surviving copy in a prescribed short interval.

The verified Scala representation of the underlying repetition, filtering, and next-stage reconstruction is documented in [3]. This article does not duplicate those maintained source proofs.

Sections 3.1–3.6 above are exact complete-period identities; none of them measure a short local window, and Section 3.6 explicitly warns against reading a positive global count as a local guarantee. The figure below shows the quantity those theorems deliberately do not bound: the fraction of one stage's own gaps that equal 2, per stage, over a real generated dataset (200 stages, the first 100,000 gaps of each). It falls sharply from 100% (stage 1, where every odd-only gap is 2) toward roughly 10% by head ≈ 1200 — consistent with $G_2(p)$'s product formula thinning the *complete-period* count, but shown here at the *local, per-stage* scale that Section 3's theorems are careful not to claim.

The figure above summarizes each stage to one number; the figure below shows the same underlying data at full per-row detail, over the same 200-stage dataset. Every 2-gap renders in a fixed green; everything between one 2-gap and the next is collapsed into a single colored unit representing that distance. A row that is mostly green is a stage whose 2-gaps are still nearly back to back; a row with long colored runs is a stage where 2-gaps have spread apart. This is the row-by-row texture behind the declining-frequency curve above, not a new claim.

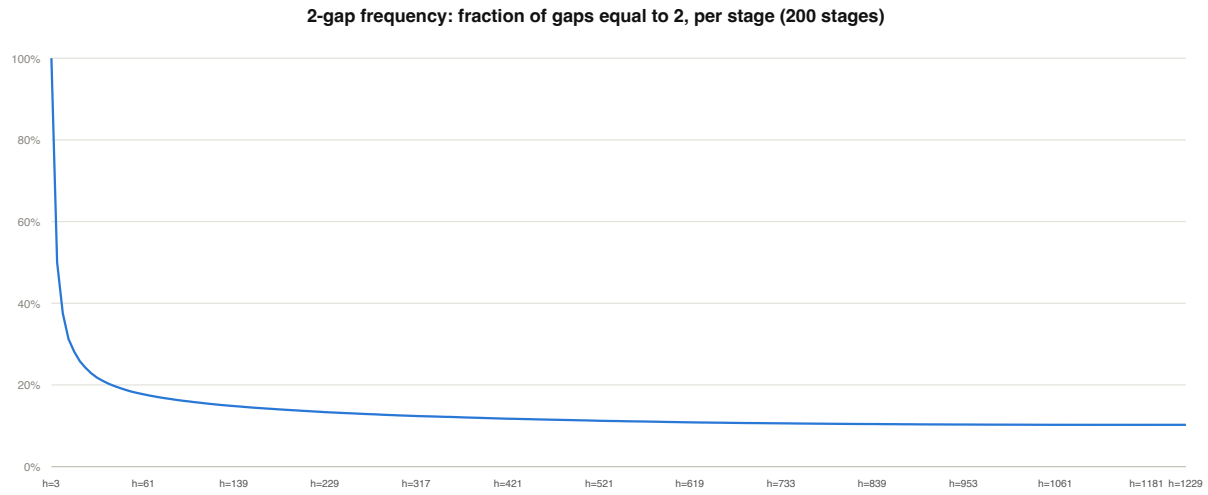


Figure 1: Fraction of a stage's gaps equal to 2, per stage, declining from 100% toward roughly 10%.

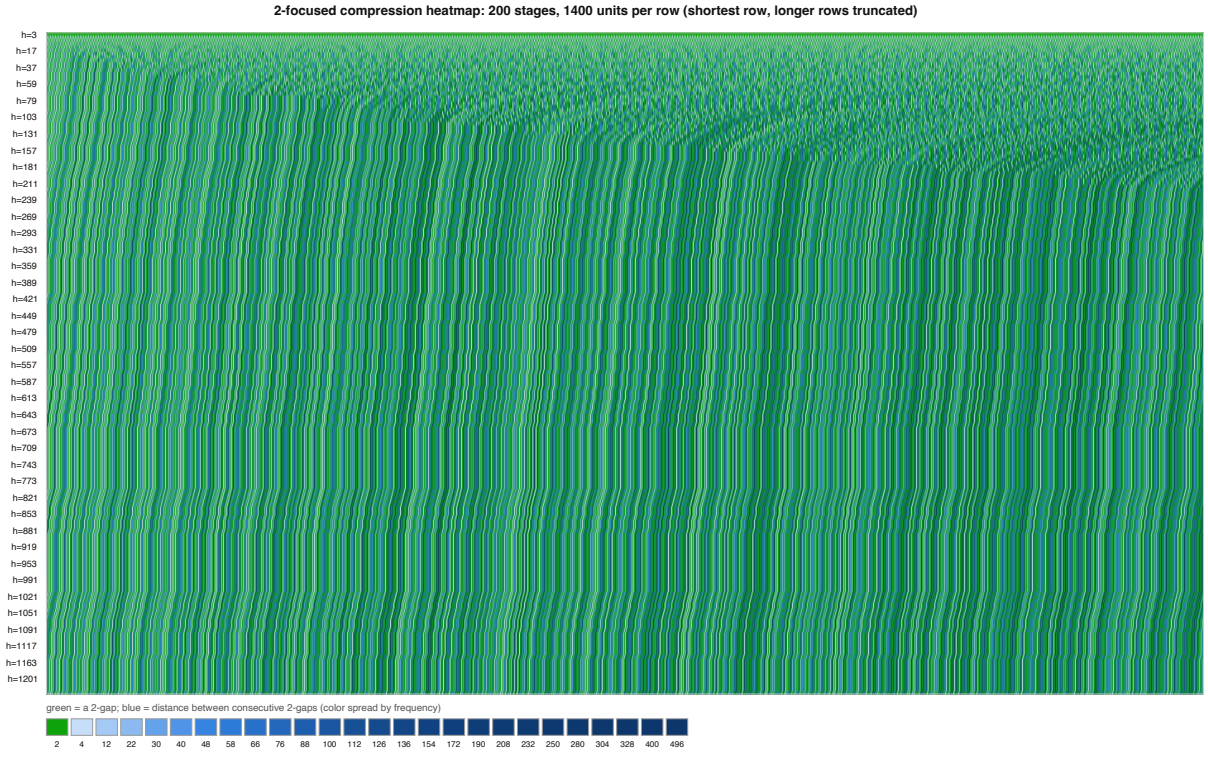


Figure 2: 2-focused compression heatmap: every 2-gap in green, the collapsed distance to the next one colored by frequency, one row per stage.

4 Local Certification And One-Transition Survival

Complete-period growth becomes relevant to twin primes only after a theorem places a surviving 2-gap inside an eligible square-safe window. This section first proves what such a survivor certifies, then isolates the exact local attrition conditions needed to retain one.

4.1 Safe-Window 2-Gaps Certify Twin Primes

The square bound turns acceptance into primality. Any composite below Q^2 has a prime divisor below Q, but every such divisor has already been installed as a filter. Therefore, for every prime future head Q and every accepted integer n with $Q \le n < Q^2$ after all primes below Q are installed, a surviving accepted pair at distance 2 is not merely a putative pair; it is a twin-prime pair.

Let

$$P_Q = \prod_{r<Q} r,$$

where r ranges over primes. If

$$Q \le n < Q^2, \qquad \gcd(n, P_Q) = 1,$$

then n is prime. Indeed, suppose that n were composite. It has a prime divisor $r \le \sqrt{n}$. The strict square bound gives $r \le \sqrt{n} < Q$, so r is one of the factors of P_Q. Consequently $r \mid n$ and $r \mid P_Q$, contradicting $\gcd(n, P_Q) = 1$:

$$\begin{aligned}
n \text{ composite} &\implies \exists r \text{ prime} : r \mid n \ \wedge \ r \le \sqrt{n} && \text{[Small Prime Divisor]} \\
&\implies r < Q && \text{[Strict Bound } n < Q^2\text{]} \\
&\implies r \mid P_Q && \text{[By Definition]} \\
&\implies \gcd(n, P_Q) > 1 && \text{[Since } r \mid n\text{]} \\
&\implies \bot. && \text{[Contradiction]}
\end{aligned}$$

Hence n is prime. Applying this result independently to both endpoints gives

$$\begin{aligned}
&Q \le x, \qquad x + 2 < Q^2, \qquad \gcd(x(x+2), P_Q) = 1 \\
&\implies x \text{ and } x + 2 \text{ are prime} \qquad \blacksquare \text{ [Q.E.D.]}
\end{aligned}$$

The right-endpoint inequality must be strict: Q^2 is composite but has no prime divisor below Q. The theorem certifies one eligible survivor; it does not prove that any such survivor exists.
The derivation above proves the theorem and its endpoint discipline.

4.2 Isolation Of 2-Gaps After Filter 3

After filter 3, two 2-gaps cannot share an endpoint, for every sieve stage whose modulus M is divisible by 6. This improves the sharp destruction capacity of one later filter strike: for every later deletion of an accepted value, removing one accepted value can destroy at most one existing 2-gap.

Suppose both $(x, x+2)$ and $(x+2, x+4)$ were accepted 2-gaps. The three values $x, x+2, x+4$ occupy all three residue classes modulo 3, so exactly one is divisible by 3. Because $3 \mid M$, that value cannot be accepted. Thus

$$\begin{aligned}
&(x, x+2) \text{ and } (x+2, x+4) \text{ both accepted} \\
&\implies 3 \nmid x, \ 3 \nmid (x+2), \ 3 \nmid (x+4) && \text{[Acceptance]} \\
&\implies \bot && \text{[Complete Residue System Modulo 3].}
\end{aligned}$$

Equivalently, parity and filter 3 force every 2-gap start into one residue:

$$\begin{aligned} x, x+2 \text{ accepted} &\Longrightarrow x \equiv 1 \pmod 2 \quad \text{[Filter 2]} \\ &\Longrightarrow x \equiv 5 \pmod 6. \quad \text{[Filter 3; CRT]} \end{aligned}$$

Any destroyed 2-gap must contain a removed accepted value as an endpoint. A value could belong to two 2-gaps only in the forbidden overlapping configuration above. Let $N_{\text{destroyed}}$ be the number of destroyed 2-gaps and N_{removed} the number of removed accepted values. Therefore

$$N_{\text{destroyed}} \leq N_{\text{removed}}.$$

■ [Q.E.D.]

Isolation limits destruction efficiency. It does not imply that a chosen window contains a 2-gap before or after the filter.

The isolation argument above forces $x \equiv 5 \pmod 6$: the smallest gap that can separate two consecutive 2-gaps is therefore the classic prime-quadruplet pattern $[2, 4, 2]$, a distance of exactly 4, never less. The figure below measures this directly on the same 200-stage generated dataset used in Section 3: the average and maximum distance between consecutive 2-gaps (summed runs, not individual raw gaps), per stage. The average grows steadily ($4 \to \sim 125$) while the maximum grows much faster ($\sim$1450 by the last stage), diverging further as the head grows — but the average's floor sits at exactly 4 across every stage, never lower, matching the isolation bound proved above rather than merely approximating it.

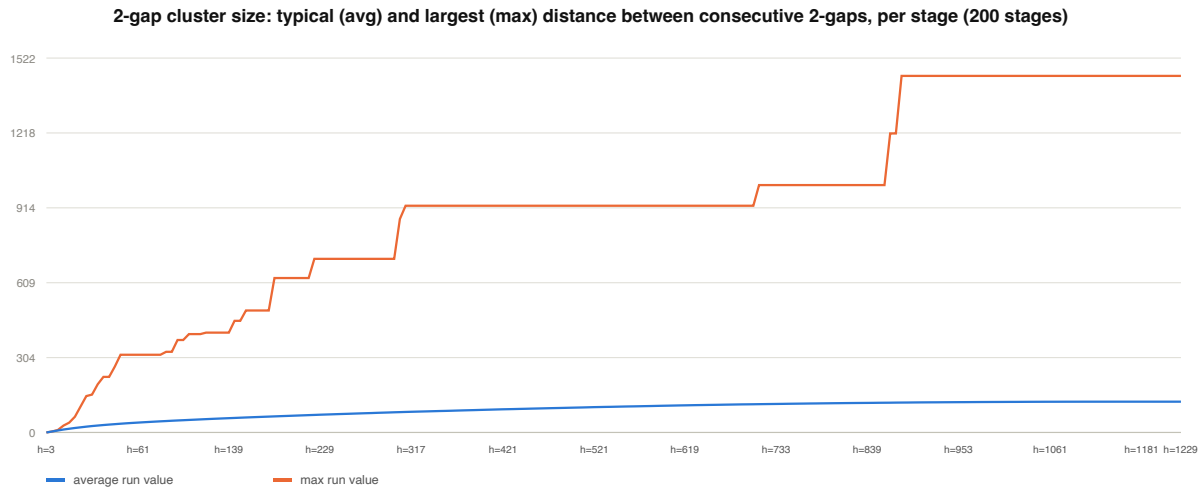


Figure 3: Average and maximum distance between consecutive 2-gaps, per stage; the average floor is exactly 4.

The overlap argument above proves its filtering consequence.

4.3 Exact Accepted Local Filter Strikes

Counting every multiple of r overstates local destruction because most multiples have already been removed by smaller filters. For every incoming prime $r \geq 5$ and its next prime future head Q, the

remaining multiples in that square-safe window admit an exact prime-multiplier description, proved here using Bertrand's postulate as an external dependency.

Let

$$M_r = \prod_{s<r} s, \qquad K = \left\lfloor \frac{Q^2 - 1}{r} \right\rfloor,$$

where s ranges over primes and Q is the next prime after r. Before filter r is applied, the accepted multiples of r in $[Q, Q^2)$ are exactly

$$r\ell \qquad \text{with } \ell \text{ prime and } r \le \ell \le K.$$

Consequently their number is

$$A(r, Q) = \pi(K) - \pi(r - 1).$$

Every multiple of r in the window has the form $r\ell$ with $2 \le \ell \le K$, because Bertrand's postulate gives $r < Q < 2r$. Since $\gcd(r, M_r) = 1$, this value survived the earlier filters exactly when $\gcd(\ell, M_r) = 1$.

Bertrand's bound also gives

$$\begin{aligned} K &< \frac{Q^2}{r} && \text{[By Definition Of } K\text{]} \\ &< 4r && \text{[Bertrand: } Q < 2r\text{]} \\ &\le r^2. && \text{[Since } r \ge 5\text{]} \end{aligned}$$

If an accepted multiplier $\ell < K + 1 \le r^2$ were composite, its least prime divisor would be at most $\sqrt{\ell} < r$ and would divide M_r, contradicting $\gcd(\ell, M_r) = 1$. Thus every accepted multiplier is prime and at least r. Conversely, every prime $\ell \ge r$ is coprime to M_r. Therefore

$$\begin{aligned} &\#\{n \in [Q, Q^2) : r \mid n,\ \gcd(n, M_r) = 1\} \\ &= \#\{\ell : r \le \ell \le K,\ \ell \text{ prime}\} && \text{[Accepted Multiplier Characterization]} \\ &= \pi(K) - \pi(r - 1) && \text{[Prime-Counting Definition]} \\ &= A(r, Q). && \blacksquare \text{ [Q.E.D.]} \end{aligned}$$

This is the exact count of accepted values struck by filter r in the stated window. A struck value need not be a 2-gap endpoint, so $A(r, Q)$ is only an upper bound on the number of destroyed local 2-gaps.

The derivation above proves the exact characterization, including its Bertrand dependency.

4.4 Sharp Local 2-Gap Survival Threshold

The proved theorem is the conditional implication below; its local-abundance antecedent remains open. It holds for every incoming prime $r \ge 5$, its next prime future head Q, and the single transition that installs filter r, over pre-filter and post-filter 2-gap starts whose two endpoints remain inside one eligible square-safe window.

The exact accepted-strike count becomes a sharp deterministic survival test. If the eligible window initially contains more 2-gaps than filter r has accepted values to remove, endpoint isolation forces at least one gap to survive.

Let $G_{\text{local}}(r, Q)$ count the pre-filter 2-gaps $(x, x + 2)$ satisfying

$$Q \le x, \qquad x + 2 < Q^2,$$

and let $G_{\text{surviving}}(r, Q)$ count those still present after filter r. With

$$K = \left\lfloor \frac{Q^2 - 1}{r} \right\rfloor, \qquad A(r, Q) = \pi(K) - \pi(r - 1),$$

let $G_{\text{destroyed}}(r, Q)$ count the destroyed eligible 2-gaps. The 2-gap isolation and accepted local strikes properties give

$$\begin{aligned} G_{\text{destroyed}}(r, Q) &\leq \#\{\text{accepted values struck in } [Q, Q^2)\} && \text{[2-Gap Isolation]} \\ &= A(r, Q). && \text{[Accepted Local Strikes]} \end{aligned}$$

Therefore

$$\begin{aligned} G_{\text{surviving}}(r, Q) &= G_{\text{local}}(r, Q) - G_{\text{destroyed}}(r, Q) && \text{[Population Accounting]} \\ &\geq G_{\text{local}}(r, Q) - A(r, Q). && \text{[Substitution]} \end{aligned}$$

Both counts are integers, so

$$G_{\text{local}}(r, Q) > A(r, Q) \Longrightarrow G_{\text{surviving}}(r, Q) > 0 \qquad \blacksquare \text{ [Integer Positivity; Q.E.D.]}$$

Equivalently, $A(r, Q) + 1$ eligible pre-filter gaps suffice. The theorem does not prove this local-abundance antecedent. Iterating it through later filters would require a fresh eligible population bound at every transition.
The derivation above proves the conditional theorem and its exact boundary.

Local Harmful-Excess Notation

Let Q be a future prime head and define the eligible start window

$$W_Q = \{x \in \mathbb{Z} : Q \leq x \ \wedge \ x + 2 < Q^2\}.$$

For a fixed filter r, let $K_r(I)$ count the incoming 2-gap starts in an interval I destroyed because r divides one endpoint. If $N_r(I)$ is the incoming population, define its centered harmful excess by

$$b_r(I) = K_r(I) - \frac{2N_r(I)}{r}.$$

The density term $2N_r(I)/r$ is not itself an upper bound. The sign and size of $b_r(I)$ encode the interval-order information discarded by separate capacity estimates.

5 Weighted Harmful-Excess Survival

Sections 3 and 4 count and certify 2-gap starts; whether a certified start survives is decided by one signed quantity, the harmful excess. This section derives the exact conservation law that makes that quantity bookkeepable across a whole conditioned chain (Section 5.1) and the terminal criterion that turns survival into a single quadratic inequality (Section 5.2).

5.1 Weighted Deletion Conservation

The shooting-versus-cluster question becomes an exact signed conservation law. Each layer has a multiplicative main term and a harmful excess. For one fixed eligible 2-gap-start population followed through every filter in one complete conditioned chain $5 \le r_0 < r_1 < \cdots < r_{m-1} < Q$ to a future head Q, with exact layer populations N_i, the weighted sum of those excesses is not an approximation: it is exactly the predicted final population minus the realized final population.

Let N_i be the number of eligible 2-gap starts immediately before filter r_i, and let N_m be the number surviving the entire conditioned chain. Define

$$a_i = 1 - \frac{2}{r_i}, \qquad A_{u,v} = \prod_{j=u}^{v-1} a_j, \qquad w_i = A_{i+1,m}, \qquad w_{-1} = A_{0,m}.$$

Let

$$T = N_0 A_{0,m}, \qquad b_i = a_i N_i - N_{i+1},$$

Moreover, $a_i w_i = A_{i,m} = w_{i-1}$ and $w_{m-1} = 1$. Multiplying the definition of b_i by w_i therefore gives

$$\begin{aligned}
\sum_{i=0}^{m-1} w_i b_i &= \sum_{i=0}^{m-1} (w_i a_i N_i - w_i N_{i+1}) && \text{[By Definition Of } b_i\text{]} \\
&= \sum_{i=0}^{m-1} (w_{i-1} N_i - w_i N_{i+1}) && \text{[Identity } a_i w_i = w_{i-1}\text{]} \\
&= w_{-1} N_0 - w_{m-1} N_m && \text{[Telescoping]} \\
&= N_0 A_{0,m} - N_m && \text{[Boundary Weights]} \\
&= T - N_m. && \text{[By Definition Of } T\text{]}
\end{aligned}$$

This is an identity, not an independent upper bound: the condition $\sum_i w_i b_i < T$ is exactly equivalent to $N_m > 0$.
The derivation above proves the recurrence, telescoping identity, and per-gap interpretation.

5.2 Terminal Survival Criterion

The strict energy inequality is sufficient for survival, for every nonempty chain $5 \le r_0 < \cdots < r_{m-1} < Q$, using its exact realized populations, over the same fixed eligible 2-gap-start population and conditioned filter chain as the conservation law of Section 5.1. Proving that it holds for infinitely many future heads remains open (the terminal survival estimate).

Define

$$E_b = \sum_{i=0}^{m-1} w_i \frac{r_i}{2(r_i - 2)} b_i^2, \qquad W_- = \sum_{i=0}^{m-1} w_{i-1}.$$

Every a_i, w_i, and w_{i-1} is positive, so $W_- > 0$.

Put $c_i = r_i/[2(r_i-2)]$. Weighted Cauchy–Schwarz yields

$$\begin{aligned}
(T-N_m)^2 &= \left(\sum_i w_i b_i\right)^2 && \text{[Exact Conservation]}\\
&= \left(\sum_i \sqrt{w_i c_i}\, b_i \sqrt{\frac{w_i}{c_i}}\right)^2 && \text{[Factorization]}\\
&\le \left(\sum_i w_i c_i b_i^2\right)\left(\sum_i \frac{w_i}{c_i}\right) && \text{[Cauchy–Schwarz]}\\
&= E_b\left(\sum_i 2w_i \frac{r_i-2}{r_i}\right) && \text{[By Definition Of } c_i\text{]}\\
&= E_b\left(2\sum_i w_i a_i\right) && \text{[Substitution]}\\
&= 2W_- E_b. && \text{[Identity } a_i w_i = w_{i-1}\text{]}
\end{aligned}$$

Hence the actual chain always satisfies

$$E_b \ge \frac{(T-N_m)^2}{2W_-}.$$

If the final population were extinct, $N_m = 0$, this lower bound would become $E_b \ge T^2/(2W_-)$. Therefore

$$E_b < \frac{T^2}{2W_-} \Longrightarrow N_m > 0 \qquad \blacksquare \text{ [Contradiction At Extinction; Q.E.D.]}$$

Safe-window certification (Section 4.1) then certifies the surviving eligible start as a twin-prime pair. The implication is a theorem for every complete conditioned chain. The open terminal survival estimate is the arithmetic statement that the strict energy inequality holds for infinitely many future heads Q. Complete-period density and separate one-layer capacity bounds do not establish that inequality.
The derivation above proves the sharp lower bound and terminal classification.

6 Why The Capacity Envelope Is Exhausted

This section presents a closed exhaustion argument rather than a single new theorem: it shows why every unsigned-capacity route to the terminal survival threshold $E_b < T^2/(2W_-)$ fails, leaving signed residue information as the only remaining ingredient. It holds for every nonempty conditioned chain $5 \le r_0 < \cdots < r_{m-1} < Q$, over the same fixed eligible 2-gap-start population and conditioned filter chain as the weighted deletion conservation and terminal survival criterion properties above. This article proves the load-bearing exhaustion steps self-contained in Appendix C.

The terminal-energy theorem of Section 5.2 reduces survival to a single inequality on the weighted harmful-excess energy E_b. The most natural way to upper-bound E_b is to maximize each layer's contribution b_i^2 separately, using only the arithmetic every residue histogram must satisfy. This produces a *capacity envelope*. The next six paragraphs walk through the results that build, refine, and ultimately exhaust that envelope.

6.1 The separate capacity envelope

Each harmful residue class can hold at most B_i incoming 2-gap starts, and the two harmful classes together can hold at most $2B_i$. The sharp capacity-envelope theorem (Appendix C.1) proves that maximizing b_i^2 over every histogram compatible with these class capacities gives the sharp separate-layer envelope

$$E_b \le \mathcal{U}_{\mathrm{cap}} = \sum_i \alpha_i X_i, \qquad X_i = \max\big((\ell_i - \mu_i)^2, (u_i - \mu_i)^2\big),$$

with $\mu_i = 2N_i/r_i$ and ℓ_i, u_i the feasible harmful-count endpoints. The capacity-stability enlargement adds the repair Γ_{cap}, enlarging the certificate threshold to $T^2/(2W_-) + \Gamma_{\mathrm{cap}}$. The capacity envelope is correct and sharp *for the information it retains*; the question is whether that information is enough.

6.2 Why capacity alone gives no positive floor

The width-floor theorem (Appendix C.2) proves the explicit per-layer floor

$$X_i \ge \frac{1}{4}\min(N_i, 2B_i, r_iB_i - N_i)^2.$$

The same property proves the obstruction: this floor vanishes at both $N_i = 0$ and $N_i = r_iB_i$. Consequently no theorem using only r_i and B_i can force a positive envelope, because the fully occupied profile $N_i = r_iB_i$ is positive and has zero capacity envelope. Progress requires either keeping realized populations away from both extremes, or replacing the capacity interval by localized residue information.

6.3 Native-period Bessel refines but does not finish

The native-period hybrid envelope intersects Bessel's inequality over the native prefix with the coordinate capacities via an exact greedy linear program. This gives the hybrid envelope $\mathcal{U}_{\mathrm{hyb}} \le \mathcal{U}_{\mathrm{cap}}$, with strict gain exactly when a normalized prefix-capacity box exceeds an interval remainder. The overflow quantification (Appendix C.2) measures the gain at cut k by the normalized overflow e_k, which the width-floor theorem lower-bounds by population slack. The envelope improves, but the next three subsections show that no cut in the chain can clear the original survival threshold.

6.4 Fixed cuts fail

The fixed-seven-cut theorem (Appendix C.3) proves that the fixed cut immediately after filter 7 fails: under the seven-layer density floor at the next untouched layer (filter 11), the hybrid envelope satisfies

$$\mathcal{U}_2^{\mathrm{hyb}} > \frac{T^2}{2W_-} \qquad \text{whenever } m \ge 37.$$

The constant 37 comes from the exact integer check $29403 \cdot 275^2 = 2{,}223{,}601{,}875 < 37 \cdot 847 \cdot 269^2 = 2{,}267{,}721{,}379$. The arbitrary-cut theorem (Appendix C.4) generalizes this to every fixed cut k:

$$m > P_k(r_k - 2)^2\left(1 + \frac{6}{D}\right)^2 \quad \Longrightarrow \quad \mathcal{U}_k^{\mathrm{hyb}} > \frac{T^2}{2W_-}.$$

For any fixed k the right side is bounded as Q grows, while m tends to infinity along any unbounded family of heads. Therefore no fixed cut can certify survival on unbounded chains.

6.5 Moving cuts lose their complete native blocks

A cut that moves outward with the head can in principle avoid the fixed-cut obstruction. The moving-cut theorem (Appendix C.5) proves the opposite pressure: a threshold-clearing cut with at least one complete native block forces

$$m < \frac{3}{7}\left(1+\frac{6}{D}\right)^2\left(\frac{2\log H}{c}-2\right)^2,$$

where H is the window length and c is a Chebyshev-type lower-bound constant for ϑ. This bound is $O(\log^2 H)$. But the prime number theorem gives $m = \pi(Q)-3 \sim Q/\log Q$, so $m/\log^2 H \sim Q/(4\log^3 Q) \to \infty$.

The prime number theorem is an explicit external dependency here; the exact logarithmic-squared inequality holds without it under five stated finite hypotheses. The asymptotic conclusion requires PNT.

Therefore, for all sufficiently large Q, any threshold-clearing cut has $M_k > H$: no complete native block remains. The incomplete-block theorem (Appendix C.6) then proves that the single incomplete block left at that moving-prime scale contributes $e_k = 0$—the native-capacity route gives back exactly $\mathcal{U}_{\mathrm{cap}}$.

6.6 The stability-gap repair is negligible

The capacity-stability enlargement had enlarged the certificate threshold by Γ_{cap}. The stability-gap theorem closes that escape: under the seven-layer density floor at filter 7,

$$\Gamma_{\mathrm{cap}} \le \frac{25P_m}{18}\left(\frac{2}{5}+\frac{3N_0}{5S}\right)^2, \qquad \mathcal{U}_{\mathrm{cap}} \ge \frac{P_m D^2}{1080}.$$

Prime Mertens and PNT make the stability gap eventually positive but negligible relative to the $P_m D^2/1080$ floor. The capacity-relaxed threshold therefore cannot rescue the separate capacity envelope on an unbounded family.

6.7 Verdict: signed information is required

The capacity-plus-native-Bessel envelope cannot certify the terminal survival threshold under the full seven-layer density floor on unbounded chains. Every route that discards residue *signs and order*—maximizing each layer over sign-compatible histograms—has been exhausted. The shared reason is that unsigned capacity forgets exactly the interval-order information that controls actual harmful excess.

This sets up the next two sections. Section 7 computes the first genuine localized saving at filter 7 by restoring exact residue order. Section 9 then bridges that saving to the residue-collision energy of Section 8.2 across complete old-period copy blocks. The full self-contained proofs of the exhaustion steps summarized in Sections 6.1–6.6 appear in Appendix C, so the reader can verify each constant without leaving the article.

7 Exact Filter-Seven Localization

At the first nontrivial conditioned layer, exact residue order replaces the quadratic capacity envelope by a constant boundary discrepancy, for every finite integer interval I, after filters 2, 3, and 5 have

been installed and immediately before filter 7, over actual pre-filter-7 2-gap starts in an arbitrary integer interval. This is an arithmetic localization theorem, not a density estimate.

The incoming starts occupy exactly three classes modulo 30. Define

$$F_7(x) = \mathbf{1}_{\{11,17,29\} \bmod 30}(x), \qquad h_7(x) = \mathbf{1}_{\{0,5\} \bmod 7}(x),$$

where h_7 marks the two endpoint-strike classes. The centered observable and its interval sum are

$$g_7(x) = F_7(x)\left(h_7(x) - \frac{2}{7}\right), \qquad b_7(I) = \sum_{x \in I} g_7(x).$$

Because 30 and 7 are coprime, g_7 has period 210. Its 21 admissible start residues in increasing order are

$$11, 17, 29, 41, 47, 59, 71, 77, 89, 101, 107,$$
$$119, 131, 137, 149, 161, 167, 179, 191, 197, 209.$$

Clearing the denominator gives weight 5 at a harmful residue and -2 at a harmless residue. In the order above the exact sequence is

$$-2, -2, -2, -2, 5, -2, -2, 5, 5, -2, -2, 5, 5, -2, -2, 5, -2, -2, -2, -2, -2.$$

It contains six harmful and fifteen harmless terms, so

$$6 \cdot 5 + 15 \cdot (-2) = 0 \qquad \text{[Complete-Period Cancellation]}$$

Starting from zero, its cumulative sums are

$$0, -2, -4, -6, -8, -3, -5, -7, -2, 3, 1, -1,$$
$$4, 9, 7, 5, 10, 8, 6, 4, 2, 0.$$

Their minimum is -8 and maximum is 10. Every non-wrapping consecutive subsum is the difference of two cumulative sums. Every wrapping subsum is the negative of its non-wrapping complement because the full sum is zero. Hence every cyclic subsum C satisfies

$$\begin{aligned} |C| &\le 10 - (-8) && \text{[Cumulative-Sum Range]} \\ &= 18. && \text{[Simplification]} \end{aligned}$$

Now partition any integer interval I into consecutive complete blocks of length 210 and one remainder of length less than 210. The complete blocks contribute zero, and the remainder selects one cyclic subsum. Therefore

$$\begin{aligned} |7b_7(I)| &\le 18 && \text{[Complete Blocks Cancel]} \\ |b_7(I)| &\le \frac{18}{7}. && \text{[Divide By 7]} \quad \blacksquare \text{ [Q.E.D.]} \end{aligned}$$

The interval from residue 47 through residue 161 attains $18/7$, so the constant is sharp. For the

complete conditioned chain with $r_0 = 5$ and $r_1 = 7$, write $P_m = A_{0,m}$. Then

$$\begin{aligned} w_1 &= A_{2,m} && \text{[By Definition]} \\ &= \frac{P_m}{a_0 a_1} && \text{[Factor } A_{0,m}\text{]} \\ &= \frac{P_m}{(3/5)(5/7)} && \text{[Substitute } r_0 = 5,\ r_1 = 7\text{]} \\ &= \frac{7P_m}{3}, && \text{[Simplification]} \\ \alpha_1 &= w_1 \frac{r_1}{2(r_1 - 2)} && \text{[Energy Coefficient]} \\ &= \frac{49P_m}{30}. && \text{[Substitution]} \end{aligned}$$

Therefore

$$\begin{aligned} \alpha_1 b_7^2 &\le \frac{49P_m}{30}\left(\frac{18}{7}\right)^2 && \text{[Sharp Interval Bound]} \\ &= \frac{54}{5}P_m. && \text{[Simplification]} \end{aligned}$$

The saving comes from the exact ordered residue pattern. It controls one fixed early layer and does not give a uniform bound for the growing family of later coefficients.
The derivation above proves the exact certificate and arbitrary-interval bound.

8 Open Estimates And Their Proven Reductions

The filter-seven excess bound and copy-block excess control properties are the last two steps of a longer argument, and the article's conclusion (Section 10, Section 12) identifies two open estimates at the live twin-prime frontier. This section introduces them so that the article's conclusion is readable on its own. It states each estimate's exact reduction, open bound, and relationship to what the article has already proved. Further algebraic reductions (activation shells, CRT lift indices, and Gram matrices) exist beyond the article's scope. This section states the exact reductions needed to follow the article's own arc and keeps the remaining estimate open.

8.1 Accepted-Boundary Discrepancy Estimate

The reduction below is proved for every nonempty conditioned chain $5 \le r_0 < \cdots < r_{m-1} < Q$, over accepted anchor values in one interval, followed through the conditioned filter chain; the terminal signed-mean-square estimate is open.

Recall from Section 5.1 that the harmful excess at layer i is $b_i = a_i N_i - N_{i+1}$. The accepted-boundary model isolates the *strike-density error* $\varepsilon_i = H_i/A_i - 1/r_i$, where H_i is the count of accepted anchors struck by filter r_i and A_i is the accepted-anchor population. Its central result is an exact bridge: the expected bulk density $1/r_i$ is already exact, and the harmful excess reduces to a difference of signed Möbius boundary sums. Specifically, with $E_i = A_i - \ell\varphi(P_i)/P_i$ the centered inclusion–exclusion discrepancy at layer i,

$$H_i - \frac{A_i}{r_i} = \left(1 - \frac{1}{r_i}\right) E_i - E_{i+1}, \qquad \varepsilon_i = \frac{(1 - 1/r_i)E_i - E_{i+1}}{A_i}.$$

The signed discrepancies telescope under one-anchor survival weights, but the weighted-energy budget requires a *weighted sum of their squares*:

$$\sum_i w_i \frac{r_i}{2(r_i-2)} \left(\left(1 - \frac{1}{r_i}\right) E_i - E_{i+1} \right)^2 .$$

Bounding the divisor summands independently gives only $|E_P| < 2^{\omega(P)} - 1$, exponentially too large. Success requires signed cancellation, correlation between consecutive layers' boundary sums, or cross-layer averaging — a genuinely new arithmetic input. The properties from Strike Divisor-Activation Kernel through First-Deletion Reindexing catalog the activation-shell, CRT-lift, summatory, Gram, and first-deletion reductions; each returns the original energy after exact algebra, so the remaining input is signed arithmetic, not another coordinate rewrite.

Why this is the article's general coefficient. The filter-7 calculation of Section 7 is the one-layer, one-prime instance of this same discrepancy: $|b_7| \le 18/7$ came from exact residue *order*, and the general layer coefficient b_i is exactly the two-residue boundary discrepancy studied here.

8.2 Residue-Collision Energy Estimate

For every incoming prime $r \ge 5$ and its actual conditioned layer population, over 2-gap-start residues modulo one incoming prime in one conditioned layer, the reduction below is proved; the relative four-point correlation estimate is open.

The residue-collision energy is the input that the copy-block bridge (Section 9) consumes. Let c_t count the incoming 2-gap starts in residue class $t \bmod r$, so $N_r = \sum c_t$. The centered deviation is $d_t = c_t - N_r/r$, and the residue-collision energy is

$$V_r = \sum_{t \bmod r} d_t^2 .$$

The two harmful classes contain $c_0 + c_{-2}$ starts. Their centered excess reduces exactly to the histogram second moment and its autocorrelation:

$$C_r = \sum_{t \bmod r} c_t^2 = N_r + 2 \sum_{h \ge 1} A_r(6rh),$$

where $A_r(\cdot)$ is the four-point autocorrelation of the start indicator at the given shift. The open estimate is the *relative* bound

$$C_r \le N_r + \frac{N_r^2}{r}, \qquad \text{equivalently} \qquad V_r \le \frac{N_r^2}{r}.$$

An absolute upper-bound-sieve estimate is insufficient until its normalization by the actual N_r is justified independently. Minimal falsifying histograms exist at small scale ($3+2+1$ at $(r, N) = (5, 6)$; $2 + 2$ at $(7, 4)$), but exact conditioned-layer search through $Q \le 251$ found none.

How the two compose. The copy-block bridge of Section 9 proves that the complete-block harmful excess $B_j = d_{t_j} + d_{t_j - 2}$ satisfies $\sum_j B_j^2 \le 4V_r$. Therefore a relative bound for the residue-collision energy V_r immediately controls the complete-block portion of the terminal harmful excess. The frontier is consequently a single composition: a relative residue-energy estimate feeding the signed boundary discrepancy through the copy-block bridge, with the two partial old-period boundary fragments still open. This is exactly what the article's conclusion (Section 10) names.

9 Copy-Block Harmful Excess And Residue Energy

The harmful excess of a copy block is not an arbitrary scalar. It is exactly the sum of two centered entries of the old start histogram modulo r, for every old period $M \geq 1$, every incoming prime $r \geq 5$ coprime to M, every run of complete copy blocks, and every finite integer interval in the copied stream, over lifted copies of one complete old-period 2-gap-start set, grouped into old-period copy blocks. This turns residue-collision energy into a quantitative bound for the complete block portion of a local interval.

Let $S \subset [0, M)$ be the old 2-gap starts and put $N = |S|$. For $t \bmod r$, define

$$c_t = \#\{a \in S : a \equiv t \pmod r\}, \qquad d_t = c_t - \frac{N}{r}, \qquad V_r = \sum_{t \bmod r} d_t^2.$$

Copy block j is $[jM, (j+1)M)$. Let K_j count starts $a + jM$ destroyed by filter r and define its centered harmful excess

$$B_j = K_j - \frac{2N}{r}.$$

Set $t_j \equiv -jM \pmod r$. The two endpoint strikes are disjoint because $r > 2$, and they are equivalent to

$$a \equiv t_j \pmod r, \qquad a \equiv t_j - 2 \pmod r.$$

Consequently,

$$\begin{aligned} K_j &= c_{t_j} + c_{t_j - 2} && \text{[Two Endpoint Classes]} \\ B_j &= d_{t_j} + d_{t_j - 2}. \ \blacksquare && \text{[Centering; Q.E.D.]} \end{aligned}$$

Because $\gcd(M, r) = 1$, the map $j \mapsto -jM \pmod r$ permutes all residues. Also $\sum_t d_t = \sum_t c_t - N = 0$. Hence a complete run of r blocks has

$$\begin{aligned} \sum_{j=0}^{r-1} B_j &= \sum_{t \bmod r} (d_t + d_{t-2}) && \text{[Residue Permutation]} \\ &= 2 \sum_{t \bmod r} d_t && \text{[Cyclic Reindexing]} \\ &= 0. && \text{[Centered Histogram]} \end{aligned}$$

The same permutation gives the exact energy identity

$$\begin{aligned} \sum_{j=0}^{r-1} B_j^2 &= \sum_{t \bmod r} (d_t + d_{t-2})^2 && \text{[Residue Permutation]} \\ &= 2V_r + 2 \sum_{t \bmod r} d_t d_{t-2}. && \text{[Expansion]} \end{aligned}$$

The autocorrelation may be negative. Discarding its sign with $2xy \leq x^2 + y^2$ gives

$$\begin{aligned} 2 \sum_t d_t d_{t-2} &\leq \sum_t d_t^2 + \sum_t d_{t-2}^2 && \text{[Termwise Quadratic Bound]} \\ &= 2V_r, && \text{[Cyclic Reindexing]} \\ \sum_{j=0}^{r-1} B_j^2 &\leq 4V_r. \ \blacksquare && \text{[Substitution; Q.E.D.]} \end{aligned}$$

For any k consecutive blocks with $0 \le k < r$, Cauchy–Schwarz yields

$$\begin{aligned}\left|\sum_{j\in J} B_j\right|^2 &\le k\sum_{j\in J} B_j^2 && \text{[Cauchy–Schwarz]}\\ &\le k\sum_{j=0}^{r-1} B_j^2 && \text{[Nonnegative Terms]}\\ &\le 4kV_r. && \text{[Complete-Block Energy]}\end{aligned}$$

Thus $|\sum_{j\in J} B_j| \le 2\sqrt{kV_r}$. For a longer run, remove complete groups of r blocks by the zero-sum identity and take k to be the number of remaining blocks.

An arbitrary integer interval I has a left partial old-period block, a run of complete blocks, and a right partial block. Each partial block contains at most one copy of every start in S. Since $r \ge 5$ and the strike indicator is either 0 or 1,

$$\left|\mathbf{1}_{r|x(x+2)} - \frac{2}{r}\right| \le 1 - \frac{2}{r}.$$

Each partial contribution is therefore at most $N(1-2/r)$ in absolute value. Combining both fragments with the complete-block bound gives

$$|b_r(I)| \le 2N\left(1-\frac{2}{r}\right) + 2\sqrt{kV_r}, \qquad 0 \le k < r \qquad \blacksquare \text{ [Triangle Inequality; Q.E.D.]}$$

Finally, if $C_r = \sum_t c_t^2$ is the residue collision count, direct expansion gives

$$\begin{aligned}V_r &= \sum_t \left(c_t - \frac{N}{r}\right)^2 && \text{[By Definition]}\\ &= \sum_t c_t^2 - \frac{2N}{r}\sum_t c_t + \frac{N^2}{r} && \text{[Expansion]}\\ &= C_r - \frac{N^2}{r}. && \text{[Since } \sum_t c_t = N\text{]}\end{aligned}$$

Therefore a relative collision-energy theorem for the actual conditioned starts would control the complete-block contribution to Section 5's harmful excess. The reduction remains incomplete: no suitable relative bound for V_r is proved across the growing layers, two partial fragments remain, and the layer bounds must still compose under the weights. When M exceeds the whole square-safe window there may be no complete block, and the boundary term dominates.
The derivation above proves the exact identities, energy bound, and arbitrary-interval boundary.

10 Routes That Are Now Classified

Complete-period counting, native-period Bessel bounds, and separate capacity envelopes do not resolve the late short-window problem. Accepted-anchor recursion also returns the existing summatory coprime discrepancy after exact CRT cancellation. These facts do not refute the quadratic survival condition; they identify which additional information it must use.

The live twin-prime program is now narrow:

control V_r relatively in the actual short window,

control partial blocks, then compose the signed layers.

More optimization of unsigned capacity or complete-period norms cannot supply those missing facts.

10.1 The Fixed-Seed Scale Conflict

The exhaustion argument of Section 6 shows that no unsigned-capacity envelope can certify survival. There is a deeper, structural reason complete-period counting cannot place a survivor in a chosen window: the primorial modulus outgrows the window itself.

Constrain the consecutive prime chain to end below the square horizon of its seed prime p, so that the future head Q satisfies $Q < p^2$. This keeps the chain numerically short, but it conflicts with using many repeated copies of one fixed seed residue inside the final safe window $[Q, Q^2)$. The seed period is the primorial

$$M_p = \prod_{r<p} r,$$

and by the prime number theorem in Chebyshev-theta form,

$$\log M_p = \sum_{r<p} \log r \sim p.$$

Hence $M_p = \exp((1 + o(1))p)$. Meanwhile $Q < p^2$ implies $Q^2 < p^4$. Therefore

$$\begin{aligned} &\frac{M_p}{Q^2} > \frac{\exp((1 + o(1))p)}{p^4} \longrightarrow \infty, && \text{[Asymptotic Limit]} \\ &M_p > Q^2 && \text{[Eventually]}. \end{aligned}$$

For all sufficiently large scenarios satisfying $Q < p^2$, one fixed residue class modulo M_p occurs at most once in $[Q, Q^2)$. Its exact global repetition frequency therefore cannot force local placement.

This is not a disproof of survival. It says that a proof cannot simultaneously rely on a short chain under p^2 and on many local copies of one fixed seed. A viable averaging argument must instead range over all seed residues at once, use a much earlier seed with a longer filter chain, average over final heads Q, or introduce an additional factorization or additive structure that creates a bilinear variable. The prime number theorem is an explicit external dependency in this argument.

This scale conflict is the geometric counterpart of the exhaustion verdict: complete-period identities are exact globally but become locally uninformative once the primorial exceeds the window. The signed local estimates of Sections 7–9 are precisely the response to this obstruction.

11 A Distinct Almost-Prime Program

Requiring both endpoints of a square-safe 2-gap to be prime reaches the twin-prime boundary. A separate program relaxes the second endpoint to have at most two prime factors. That program has different local factors and a different Type-I/Type-II formulation; it is developed separately.

Its success would not prove a surviving 2-gap or infinitely many twin primes.

11.1 Why This Is Hard: The Type-II Barrier

The exhaustion argument and the fixed-seed scale conflict together explain why this project's complete-period algebra cannot finish the twin-prime program. The deeper reason the program is hard at all comes from the contemporary theory of prime-producing sieves.

Ford and Maynard study nonnegative sequences whose difference from a comparison model satisfies Type-I and Type-II estimates. Type-I controls divisibility averages over many factor scales; Type-II controls bilinear sums against arbitrary bounded coefficient sequences. Their framework proves that a *substantial Type-II range is genuinely necessary* to guarantee a nontrivial prime lower bound: very strong Type-I information alone can still be consistent with a sequence containing no primes.

This project's exact residue and CRT formulas are algebraic input for a possible Type-I analysis. They are not yet a Type-I theorem over the required short intervals, because the accumulated discrepancy norm has not been bounded—the fixed-seed scale conflict of Section 10.1 is one face of this. No arbitrary-coefficient Type-II estimate is currently proved for the sieve-sequence weights.

The relaxed almost-prime program of Section 11 makes the Type-II barrier concrete. Its scalar-centered final weight retains a nonprincipal character mode modulo 3 on the complete reduced wheel: a bounded product coefficient can correlate perfectly with the full relaxed survivor count. This refutes the shortcut that scalar-density centering alone creates Type-II orthogonality—the exact obstruction Ford–Maynard's framework predicts. The refutation, the implication from relaxed positivity to prime-plus-P_2 production, and the program's exact divisor and bilinear character properties are proved in Relaxed Almost-Prime Production in Sieve Sequences.

Green and Sawhney's work on prime values of p^2+nq^2 demonstrates a modern successful Type-II strategy using extra algebraic variables, number-field factorization, and Gowers-norm machinery. Those structural inputs do not automatically exist for the affine pair $(x, x+2)$. Their result is therefore a methodological guide, not a theorem that transfers.

The actionable consequence is that any route beyond the exhaustion boundary of Section 6 must either prove a genuine short-window signed Type-I estimate for the residue-energy / accepted-boundary quantities of Section 8, or introduce a new bilinear variable that supplies the Type-II cancellation the affine pair lacks.

12 Conclusion

Complete-period sieve algebra gives exact global 2-gap and `(2,4,2)` cluster counts. One old 2-gap has $r-2$ surviving lifts under one incoming prime, each old cluster has $r-4$ intact lifts, finite batches compose by CRT, rotation preserves cyclic multiplicity, and global 2-gap extinction is stable. Writing C_M for the cyclic cluster count at modulus M:

$$
\begin{aligned}
G_2(p) &= \prod_{3\le q<p}(q-2), && \text{[Exact Global Count]}\\
G_{\text{after}} &= G\prod_{r\in\mathcal{R}}(r-2), && \text{[Exact Batched Survival]}\\
C_{rM} &= (r-4)C_M, && \text{[Exact Cluster Recurrence]}\\
\#\{i : \text{rot}_j(G)_i = 2\} &= \#\{i : g_i = 2\}. && \text{[Rotation Bijection]}\\
G_s(2) = 0 &\Longrightarrow G_t(2) = 0 \quad (t \ge s). && \text{[Stable Global Absence]}
\end{aligned}
$$

The local theorems identify both the certificate and the sharp one-transition condition:

$$
\begin{aligned}
Q \le x,\ x+2 < Q^2,\ \gcd(x(x+2), P_Q) = 1 &\Longrightarrow x, x+2 \text{ prime}, && \text{[Square-Safe Certification]}\\
G_{\text{surviving}}(r,Q) &\ge G_{\text{local}}(r,Q) - A(r,Q), && \text{[Exact Accepted Strikes]}\\
G_{\text{local}}(r,Q) > A(r,Q) &\Longrightarrow G_{\text{surviving}}(r,Q) > 0. && \text{[Sharp Local Threshold]}
\end{aligned}
$$

For a complete conditioned chain, the signed conservation law (Section 5.1) and its weighted Cauchy–Schwarz corollary (Section 5.2) prove

$$\sum_i w_i b_i = T - N_m, \qquad \text{[Exact Conservation]}$$

$$E_b \geq \frac{(T - N_m)^2}{2W_-}, \qquad \text{[Weighted Cauchy–Schwarz]}$$

$$E_b < \frac{T^2}{2W_-} \Longrightarrow N_m > 0. \qquad \text{[Terminal Implication]}$$

The properties from Filter-Seven Excess Bound through Copy-Block Excess Control then add exact local arithmetic:

$$|b_7(I)| \leq \frac{18}{7}, \qquad \text{[Sharp Filter-7 Boundary]}$$

$$B_j = d_{t_j} + d_{t_j - 2}, \qquad \text{[Exact Copy-Block Formula]}$$

$$\sum_{j=0}^{r-1} B_j^2 = 2V_r + 2\sum_t d_t d_{t-2} \leq 4V_r. \qquad \text{[Residue-Energy Bridge]}$$

The proved results therefore move the question beyond global density and unsigned capacity. They do not complete the twin-prime program. The remaining theorem must control relative residue energy for the actual conditioned populations, control the two partial old-period boundary fragments, and make those estimates beat the weighted terminal threshold through an unbounded family of future heads.

References

[1] G. H. Hardy and E. M. Wright. *An Introduction to the Theory of Numbers*. Oxford University Press, 6th edition, 2008. https://doi.org/10.1093/oso/9780199219858.001.0001.

[2] Fred B. Holt and Helgi Rudd. Eratosthenes sieve and the gaps between primes. arXiv:1408.6002, v1, 2014. Section 1.2, Lemmas 2.1 and 3.1, and Corollary 6.3. https://arxiv.org/abs/1408.6002v1.

[3] T. H. Mata. Formal Verification of Sieve Sequence Stages and Their Transitions. https://doi.org/10.5281/zenodo.22955782, 2026. Zenodo:10.5281/zenodo.22955782.

A Evidence And Verification Status

The table separates the verified construction foundation from the mathematical results discussed here and the open frontier. Proof status is distinct from historical priority: established results rederived here are credited above, and priority of the specific signed/local formulations remains unestablished. The rotation operation's underlying membership and size behavior is Stainless-verified, but its exact gap-multiplicity theorem belongs to the mathematical layer.

Claim group	Status	Article location
Sieve construction and transition mechanics	Stainless-verified foundation	Companion Sieve Sequence article
Complete-period, local, weighted, capacity-exhaustion, filter-seven, and copy-block results	Mathematically proved	Sections 3–7, Section 9, and Appendix C
Accepted-boundary discrepancy and residue-collision energy	Exact reductions proved; required estimates open	Section 8

Classical number-theoretic inputs enter as follows: Bertrand's postulate supports the exact accepted-local-strikes characterization of Section 4.3 and the asymptotic corollaries in Appendix C, and the Prime Number Theorem (with Mertens estimates) supports the asymptotic stability-gap discussion of Section 6.6. Hardy and Wright [1] provide a modern reference for these classical inputs, not an attribution of their original discovery.

The operational Sieve Sequence construction used by these mathematical properties is Stainless-verified separately in Formal Verification of Sieve Sequence Stages and Their Transitions.

B Self-Contained Proofs For The Exhaustion Chain

This appendix gives the full proofs of the capacity-exhaustion steps summarized in Section 6, so the article is self-contained. Each entry states the property's population, scope, derivation, and boundary. Its proofs are the authority for the exhaustion claims made in this article.

Notation shared across the entries: $D = Q^2 - Q - 3$, $a_i = 1 - 2/r_i$, $P_i = \prod_{j<i} a_j$, $P_m = \prod_{j<m} a_j$, $T = N_0 P_m$, and $W_- = \sum_{i<m} P_m/P_i$. The per-layer squared harmful-excess envelope is written X_i throughout (renamed from M_i to avoid collision with the native modulus M_k).

B.1 Sharp Harmful-Capacity Excess Envelope

Fix one filter layer's residue histogram, restricted to the two harmful classes: for every incoming prime $r > 2$, every common residue-class capacity $B \geq 0$, and every population $0 \leq N \leq rB$, let the residue counts c_a ($a \bmod r$) satisfy $0 \leq c_a \leq B$ and $\sum c_a = N$. The two harmful classes hold $N_{\text{harm}} = c_0 + c_{-2}$ starts, and the signed harmful excess is $b = N_{\text{harm}} - 2N/r$.

The total N_{harm} is constrained by two facts. First, each harmful class holds at most B, so $N_{\text{harm}} \leq 2B$. Also $N_{\text{harm}} \leq N$. Second, the other $r-2$ classes hold at most $(r-2)B$ of the N starts, forcing at least $N - (r-2)B$ into the harmful pair. The sixfold-capacity envelope proves both endpoints are attainable, giving the exact feasible interval

$$\ell \leq N_{\text{harm}} \leq u, \qquad \ell = \max(0, N - (r-2)B), \qquad u = \min(N, 2B).$$

Because $b^2 = (N_{\text{harm}} - 2N/r)^2$ is convex in N_{harm}, its maximum over $[\ell, u]$ occurs at an endpoint:

$$b^2 \leq X_{r,N,B} := \max\left\{ \left(\ell - \frac{2N}{r}\right)^2, \left(u - \frac{2N}{r}\right)^2 \right\}.$$

Both endpoints are attainable, so this bound is sharp. The conditioned-chain upper envelope follows by summing the per-layer sharp bounds with the energy coefficients $\alpha_i = w_i r_i/[2(r_i - 2)]$:

$$E_b \leq \mathcal{U}_{\text{cap}} := \sum_i \alpha_i X_i \qquad \blacksquare \text{ [Q.E.D.]}$$

The bound is sharp one layer at a time but need not be sharp over a chain, because the histograms attaining each X_i separately may not co-arise from one nested survivor sequence. A cross-layer CRT restriction could lower the true aggregate energy below $\mathcal{U}_{\mathrm{cap}}$. The property does not prove $\mathcal{U}_{\mathrm{cap}} < T^2/(2W_-) + \Gamma_{\mathrm{cap}}$; it reduces the capacity-only route to that explicit inequality.

B.2 Capacity-Envelope Width Floor Needs Population Slack

For every $r \geq 5$, $B \geq 0$, $0 \leq N \leq rB$, over one filter layer's feasible harmful-count interval, this property extracts the explicit lower bound on $X_{r,N,B}$ supplied by the *width* of the feasible interval $[\ell, u]$. Write $c = (\ell + u)/2$, $h = (u - \ell)/2$. The farther endpoint from $2N/r$ has distance

$$\max\left(\left|\frac{2N}{r} - (c-h)\right|, \left|\frac{2N}{r} - (c+h)\right|\right) = h + \left|\frac{2N}{r} - c\right| \geq h.$$

Therefore

$$X_{r,N,B} \geq \frac{(u-\ell)^2}{4}.$$

The width has the closed form

$$u - \ell = \min(N, 2B, rB - N).$$

This follows by splitting the feasible range into three parts. If $0 \leq N \leq 2B$, then $u = N$, $\ell = 0$, and $u - \ell = N$. If $2B \leq N \leq (r-2)B$, then $u = 2B$, $\ell = 0$, and $u - \ell = 2B$. If $(r-2)B \leq N \leq rB$, then $u = 2B$, $\ell = N - (r-2)B$, and $u - \ell = rB - N$. The formulas agree at the shared endpoints, proving the minimum formula. Combining,

$$X_{r,N,B} \geq \frac{1}{4}\min(N, 2B, rB - N)^2 \qquad \blacksquare \text{ [Q.E.D.]}$$

Zero characterization. Because $X_{r,N,B}$ is the maximum of two squares, $X_{r,N,B} = 0$ iff both endpoints equal $2N/r$, which requires $u = \ell$. By the width formula this is equivalent to $\min(N, 2B, rB - N) = 0$. Since $2B > 0$ and $0 \leq N \leq rB$, this occurs exactly at

$$X_{r,N,B} = 0 \iff N \in \{0, rB\}.$$

This is the obstruction: the envelope vanishes at both the empty and the fully-occupied populations. Theorem using only r and B cannot force a positive floor, because the fully-occupied profile $N = rB$ is positive and has zero capacity envelope.

B.3 Fixed Seven Cut Cannot Clear The Original Threshold

Consider the cut immediately after filter 7, so $k = 2$, for every chain $r_0 = 5, r_1 = 7, r_2 = 11, \ldots$ with $Q \geq 17$ and $m \geq 37$, examining the suffix capacity term at the first untouched layer (filter 11) of a conditioned chain and assuming the seven-layer density floor's local-count threshold at filter 11. The native-period hybrid envelope leaves every coordinate $i \geq 2$ under its separate capacity bound, giving

$$\mathcal{U}_2^{\mathrm{hyb}} \geq \alpha_2 X_2.$$

The first three multiplicative factors are $a_0 = 3/5$, $a_1 = 5/7$, $a_2 = 9/11$, so $P_3 = a_0 a_1 a_2 = 27/77$. Since $w_2 = P_m/P_3$ and $\alpha_2 = w_2/(2a_2)$,

$$\alpha_2 = \frac{P_m}{2a_2 P_3} = \frac{P_m}{2 \cdot (9/11) \cdot (27/77)} = \frac{847}{486} P_m.$$

The filter-11 capacity floor gives $X_2 \geq B_{11}^2$ with $B_{11} = \lfloor D/66 \rfloor + 1 \geq D/66$. Therefore

$$\mathcal{U}_2^{\text{hyb}} \geq \frac{847}{486} P_m \left(\frac{D}{66}\right)^2.$$

For the threshold upper bound, $T = N_0 P_m$ and $W_- = \sum_{i<m} P_m/P_i \geq m P_m$ (since $P_i \leq 1$). Before filter 5 every 2-gap start is 5 mod 6, so $N_0 \leq \lfloor D/6 \rfloor + 1 \leq D/6 + 1$. Hence

$$\frac{T^2}{2W_-} \leq \frac{P_m}{2m} \left(\frac{D}{6} + 1\right)^2.$$

The suffix exceeds the threshold whenever

$$\frac{847}{486} \left(\frac{D}{66}\right)^2 > \frac{1}{2m} \left(\frac{D}{6} + 1\right)^2,$$

equivalently $m > \frac{29403}{847}(1 + 6/D)^2$. For $Q \geq 17$, $D \geq 269$, so the right side is at most $\frac{29403}{847}(275/269)^2$. The integer comparison

$$29403 \cdot 275^2 = 2{,}223{,}601{,}875 < 2{,}267{,}721{,}379 = 37 \cdot 847 \cdot 269^2$$

shows this is strictly below 37. Therefore $m \geq 37$ proves

$$\mathcal{U}_2^{\text{hyb}} > \frac{T^2}{2W_-} \qquad \blacksquare \text{ [Q.E.D.]}$$

This proves the *envelope* cannot certify survival through this fixed cut. It does not bound the actual energy E_b. It does not address later optimized cuts ($k \geq 3$), the capacity-relaxed threshold, or localized residue information.

B.4 Every Fixed Native Cut Fails The Original Threshold

Generalize C.3 to an arbitrary cut k, for every chain $5 \leq r_0 < \cdots < r_{m-1} < Q$ with $Q \geq 17$ and $2 \leq k < m$, over the suffix capacity term at the first untouched layer of an arbitrary native cut. The seven-layer density floor at layer r_k gives $X_k \geq B_k^2$ with $B_k = \lfloor D/(6r_k) \rfloor + 1 \geq D/(6r_k)$, and $\alpha_k = P_m/(2P_k a_k^2)$. Therefore

$$\mathcal{U}_k^{\text{hyb}} \geq \frac{P_m D^2}{72 P_k a_k^2 r_k^2}.$$

The threshold upper bound is unchanged: $T^2/(2W_-) \leq (P_m/2m)(D/6+1)^2$. The suffix exceeds the threshold whenever

$$\frac{P_m D^2}{72 P_k a_k^2 r_k^2} > \frac{P_m}{2m} \left(\frac{D}{6} + 1\right)^2.$$

Canceling $P_m > 0$ and rearranging gives $m > P_k a_k^2 r_k^2 (1 + 6/D)^2$. Since $a_k r_k = r_k - 2$,

$$m > P_k (r_k - 2)^2 \left(1 + \frac{6}{D}\right)^2 \quad \Longrightarrow \quad \mathcal{U}_k^{\text{hyb}} > \frac{T^2}{2W_-}.$$

Every fixed cut eventually fails: for fixed k, P_k and r_k are constants while $(1 + 6/D)^2 \to 1$ as Q grows. Along any family with $m \to \infty$, the fixed cut violates the necessary condition. Thus a cut capable of clearing the original threshold must move: $k = k(Q) \to \infty$.

A parameter-free lower bound on the cut prime follows from $k \geq 2$, giving $P_k \leq P_2 = (3/5)(5/7) = 3/7$. Threshold clearance would require

$$m \leq \frac{3}{7}(r_k - 2)^2 \left(1 + \frac{6}{D}\right)^2,$$

equivalently

$$\mathcal{U}_k^{\text{hyb}} < \frac{T^2}{2W_-} \quad \Longrightarrow \quad r_k \geq 2 + \frac{\sqrt{7m/3}}{1 + 6/D} \qquad \blacksquare \text{ [Q.E.D.]}$$

Recovery of C.3. For the cut after filter 7, $k = 2$, $r_2 = 11$, $P_2 = 3/7$, and $P_2(r_2 - 2)^2 = (3/7) \cdot 81 = 243/7 = 29403/847$, recovering C.3's constant exactly.

The bound $r_k \geq 2 + \sqrt{7m/3}/(1 + 6/D)$ uses no estimate for the distribution of primes. It is a necessary condition on the cut prime, not sufficient: it does not prove such an r_k exists in the chain. Converting the prime bound into a cut-index bound requires the prime number theorem (Appendix C.5).

B.5 Moving Cut Loses Complete Native Blocks

For every chain $5 \leq r_i < Q$, cut $2 \leq k < m$, over a native cut that moves outward with the future head together with the prime-counting function, the exact logarithmic-squared inequality below holds under five stated finite hypotheses. The asymptotic corollary uses Bertrand's postulate and the prime number theorem as explicit external dependencies; the exact theorem itself is finite, and only the asymptotic corollary is external.

The native modulus at cut k is $M_k = \prod_{p<r_k} p = 2 \cdot 3 \prod_{i<k} r_i$. Since r_{k-1} is the prime immediately before r_k,

$$\log M_k = \vartheta(r_{k-1}).$$

Assume the following five hypotheses:

1. the seven-layer density floor holds at layer r_k;
2. the cut clears the original threshold, $\mathcal{U}_k^{\text{hyb}} < T^2/(2W_-)$;
3. the native modulus fits the interval, $M_k \leq H$ where $H = D + 1 = Q^2 - Q - 2$;
4. for some constant $c > 0$, $\vartheta(r_{k-1}) \geq c r_{k-1}$; and
5. Bertrand's inequality $r_k < 2r_{k-1}$.

From C.4, hypothesis 2 forces $r_k \geq 2 + \sqrt{7m/3}/(1 + 6/D)$. From hypotheses 3–5,

$$\log H \geq \log M_k = \vartheta(r_{k-1}) \geq c r_{k-1} > \frac{c}{2} r_k.$$

Combining the lower and upper requirements on r_k and rearranging,

$$m < \frac{3}{7}\left(1 + \frac{6}{D}\right)^2 \left(\frac{2\log H}{c} - 2\right)^2.$$

This is the exact finite theorem: a threshold-clearing cut with at least one complete native block forces the chain length to be at most $O(\log^2 H)$.

Prime-number-theorem corollary (external). The asymptotics $\vartheta(x) \sim x$ and $\pi(x) \sim x/\log x$ are external to the project's verification. For the actual full chain, $m = \pi(Q) - 3 \sim Q/\log Q$, whereas $\log^2 H \sim 4\log^2 Q$. Therefore

$$\frac{m}{\log^2 H} \sim \frac{Q}{4\log^3 Q} \longrightarrow \infty.$$

The exact logarithmic-squared necessary condition fails for all sufficiently large Q. Hence, under the seven-layer density floor at the first suffix layer,

$$\mathcal{U}_k^{\mathrm{hyb}} < \frac{T^2}{2W_-} \quad \Longrightarrow \quad M_k > H$$

for every sufficiently large Q and every cut k. There are no complete native blocks to cancel.

When $M_k > H$, the native-period hybrid envelope still constrains the single incomplete interval block. C.6 addresses whether that constraint can supply the missing gain. The asymptotic conclusion depends explicitly on PNT and Bertrand; the exact inequality remains valid without them under the five stated hypotheses.

B.6 Incomplete-Block Bessel Excludes No Capacity

Under the hypotheses of C.4–C.5, for every cut $2 \le k < m$ with $M_k > H$, the native-period hybrid envelope's interval remainder on the single incomplete native block left is $s_k = H$: the whole window is one incomplete block (the asymptotic scale uses Bertrand and PNT as external dependencies). This entry proves the normalized capacity box then fits inside that budget, so the overflow e_k vanishes and the hybrid envelope collapses back to the all-capacity envelope.

The native-period hybrid envelope's exact norm at coordinate $i < k$ is

$$q_{i,k} = \frac{M_k P_i (r_i - 2)}{3r_i^2}.$$

The capacity numerator obeys $X_i \le N_i^2$. Since conditioned populations decrease, $N_i \le N_0$, and before filter 5 every 2-gap start is 5 mod 6, so $N_0 \le \lfloor D/6 \rfloor + 1 \le D/5$ (using $D \ge 269 > 30$). Therefore

$$X_i \le \frac{D^2}{25}.$$

For the denominator, the function $x \mapsto (x-2)/x^2$ is decreasing for $x \ge 4$, so for $i < k$,

$$\frac{r_i - 2}{r_i^2} \ge \frac{r_k - 2}{r_k^2}, \qquad q_{i,k} \ge \frac{M_k P_k (r_k - 2)}{3r_k^2}.$$

Summing the k prefix coordinates,

$$\sum_{i<k} \frac{X_i}{q_{i,k}} \le \frac{3kD^2 r_k^2}{25 M_k P_k (r_k - 2)}.$$

The overflow quantification defines $e_k = (\sum_{i<k} X_i/q_{i,k} - s_k)_+$. When $M_k > H$, $s_k = H$. The normalized sum is at most H—giving $e_k = 0$—whenever

$$M_k P_k \ge \frac{3kD^2 r_k^2}{25H(r_k - 2)}.$$

Prime-number-theorem scale (external). The prefix product satisfies $M_k P_k = 6 \prod_{i<k}(r_i - 2) \geq M_k/2^k$. Using PNT and Bertrand externally, C.5 forces $r_k \gg \sqrt{Q/\log Q}$ for any potentially successful cut, hence $\log(M_k P_k) \sim r_{k-1} \gg \sqrt{Q/\log Q}$. The logarithm of the right side of the zero-overflow criterion is only $O(\log Q)$, since $k < Q$, $D < H < Q^2$, and $r_k < Q$. The criterion therefore holds for every sufficiently large Q, giving

$$e_k = 0, \qquad \mathcal{U}_k^{\mathrm{hyb}} = \mathcal{U}_{\mathrm{cap}}.$$

Exhaustion of the original native hybrid. Combining C.3–C.6: fixed cuts are excluded by C.3–C.4; any potentially successful moving cut has $M_k > H$ (C.5) and then $\mathcal{U}_k^{\mathrm{hyb}} = \mathcal{U}_{\mathrm{cap}}$ (C.6). Since C.3 gives $\mathcal{U}_{\mathrm{cap}} \geq \mathcal{U}_2^{\mathrm{hyb}} > T^2/(2W_-)$, the capacity-plus-native-Bessel envelope satisfies

$$\mathcal{U}_{\mathrm{hyb}} \geq \frac{T^2}{2W_-}$$

for every sufficiently large complete chain satisfying the seven-layer density floor. The current capacity-plus-native-Bessel envelope cannot certify the terminal survival threshold on an unbounded family.

This is a method obstruction, not a refutation of the seven-layer density floor or the terminal survival estimate. It does not address the capacity-relaxed threshold $T^2/(2W_-) + \Gamma_{\mathrm{cap}}$ (handled by the stability-gap theorem, summarized in Section 6.6) or a localized upper bound for the actual E_b.