Survival Frontiers in Balanced 2-Gap Companion Processes

Thiago Henrique Ramos da Mata
Independent Researcher
thiago.henrique.mata@gmail.com
ORCID: 0009-0002-7366-939X
https://thiagomata.com/
https://github.com/thiagomata
DOI: https://doi.org/10.5281/zenodo.22955798

Proof status: The companion-process identities are proved exactly; the asymptotic theorems are conditional on the premises stated with each result (Section 1.1). Stainless verification is pending. No result is claimed for the real sieve.

Abstract

This article examines companion processes that reproduce the exact global growth of 2-gaps but change where each filter removes them. Every parent produces r copies and exactly two are removed, as in the sieve sequence. The companion may place those two removals randomly, protect a chosen target, or direct them toward it. This construction separates the number of surviving 2-gaps from their location near the head. A filter's local destruction fraction f_r is compared with the random rate $2/r$ through $w_r = rf_r/2$. Under the stated blind-placement premise, the cumulative product proves that every fixed finite value of w_r preserves square-window 2-gaps; with the additional availability and mixing conditions, it produces head 2-gaps infinitely often. The first boundary occurs when $w_r = 1 + c \log r$: square windows survive for $c < 1$, while head recurrence survives for $c < 1/2$. Under the same spatial premises, a separate exact-quota random-location model gives the same local frontiers when its normalized quotas satisfy the CRT-rate cumulative sum and summable finite-population error conditions derived below. It preserves an accepted-strike count, rather than the balanced per-parent recurrence; retaining one CRT strike count alone is insufficient. None of the spatial, availability, or mixing premises used by these conditional theorems is proved for the real sieve. Therefore, proving that the real sieve remains below the head frontier, together with persistent availability and the deterministic discrepancy bound stated in Section 10, would establish the twin-prime conjecture.

Licensed under the Creative Commons Attribution 4.0 International License (CC BY 4.0):
https://creativecommons.org/licenses/by/4.0/.

1 Introduction

Begin with the positive integers and remove multiples of prime numbers one prime at a time. At a stage headed by the prime p, the accepted values are precisely the integers not divisible by any prime smaller than p. These survivors repeat periodically. If e_i and e_{i+1} are consecutive survivors, then $e_{i+1} - e_i$ is their gap; the gaps across one complete period form a finite cycle.

This periodic accepted-value object is the **Sieve Sequence**, introduced and formally verified in Formal Verification of Sieve Sequence Stages and Their Transitions [3]; its finite gap cycle is the representation used throughout this article.

After multiples of 2 have been removed, every survivor is odd. Every gap is therefore even, and 2 is the smallest possible gap. A **2-gap** is a pair of consecutive survivors x and $x + 2$. At the stage headed by a prime Q, all primes below Q have been installed as filters. A surviving 2-gap whose endpoints lie in the square-safe window $[Q, Q^2)$ is therefore a twin-prime pair: any composite number below Q^2 has a prime divisor below Q and would already have been removed.

When the sieve later reaches the stage headed by x, the same pair appears as the first gap after the head. Infinitely many such head 2-gaps would give infinitely many twin-prime pairs. But survival can be asked at three different spatial scales:

1. some 2-gap exists somewhere in the complete period at every layer;
2. a 2-gap occurs in each sufficiently large square-safe window; or
3. the first gap after the distinguished head equals 2 infinitely often.

The first statement is global and purely combinatorial. The second is local but benefits from a window whose length grows quadratically. The third concerns one position and therefore has no window-size reserve. Mixing these meanings hides the actual threshold.

The following two empirical views make the distinction visible in the deterministic Sieve Sequence. They use the same 200 stages and the same 2-focused compression: every 2-gap receives its own green cell, while each maximal run of consecutive non-2 gaps is collapsed into one colored cell equal to its sum. An internal colored cell therefore measures the total distance between two consecutive 2-gaps. Both views display 1,400 compressed units from every row; they differ only in where those rows are placed horizontally.

View A — independent compressed snapshots. Every row begins at column zero, so its horizontal coordinate counts compressed units from that stage's own head. This is the original view. It honestly shows the texture within each stage, but its curved or noisy-looking vertical drift must not be read as a 2-gap changing before the square boundary: equal columns in adjacent rows need not represent the same surviving values.

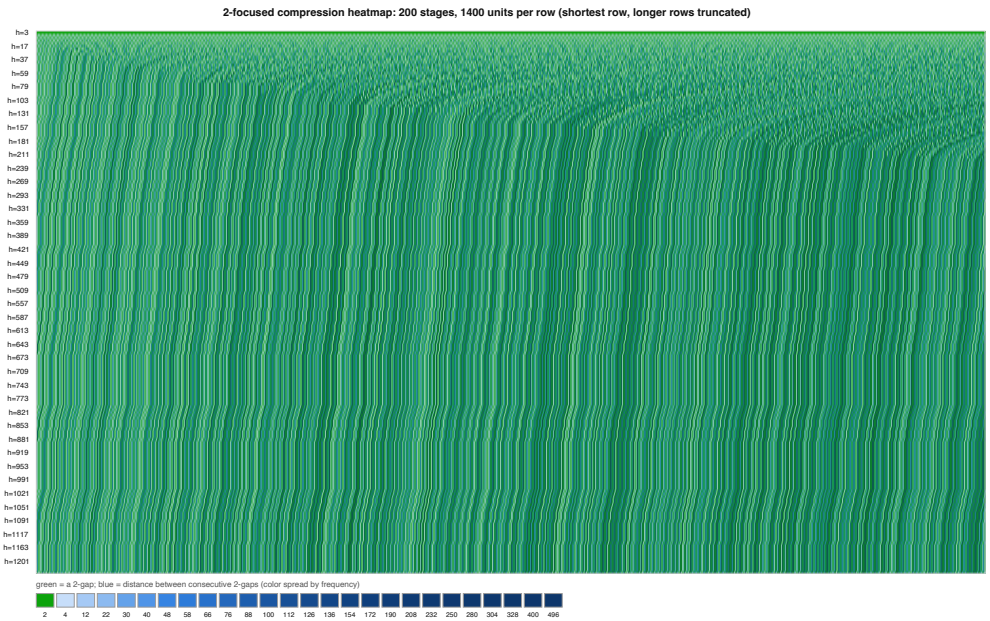


Figure 1: Independent 2-focused snapshots of 200 Sieve Sequence stages: every row restarts at its own head

View B — shared-safe-2 alignment. For each pair of adjacent rows, let h be the previous stage's head. A safe anchor is a 2-gap with the same raw starting value in both rows and with both endpoints strictly below h^2. The alignment calculation compares the compressed indices of every such shared safe 2-gap and shifts the later row by their common difference. Across the 199 observed transitions, the 200-stage dataset exhibits 118 differences of zero and 81 of one; every safe anchor within each transition agrees with that row's difference. Accumulating those differences produces the alternative view below.

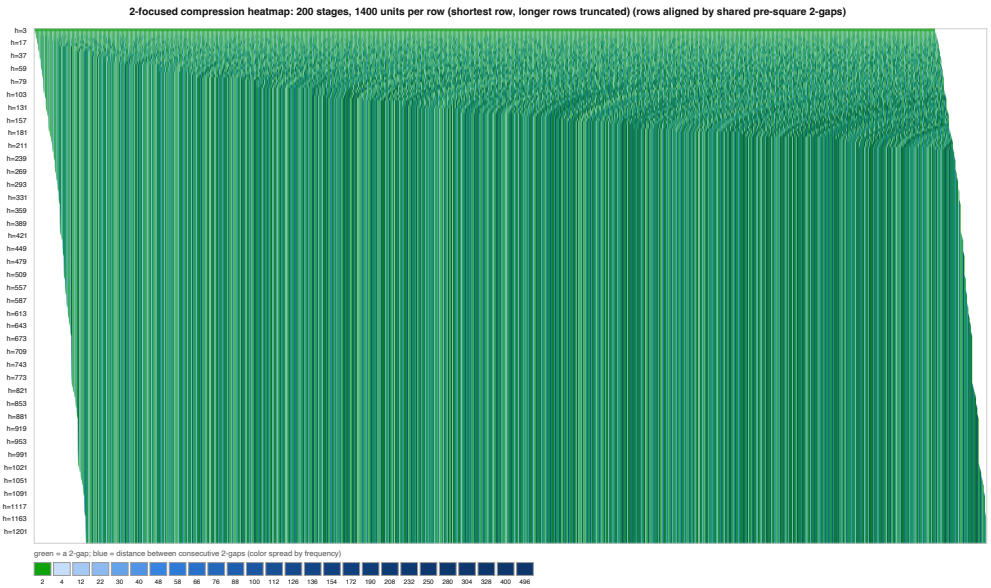


Figure 2: Shared-safe-2 aligned compression of 200 Sieve Sequence stages: unchanged pre-square 2-gaps form vertical green lines

The left white wedge is intentional padding created by the cumulative offsets. A zero shift occurs when advancing the head merely shortens the leading collapsed non-2 run; a one-cell shift occurs when that step removes an entire compressed run or a standalone 2-gap cell. Thus the straight green lines in View B and the curved texture in View A describe the same data under different coordinates.

The sieve does not rotate this compressed row as an atomic list. Within the safe prefix, it advances the head through one raw gap; only afterward does the visualization collapse the remaining consecutive non-2 gaps. Thus a raw prefix $[4, 6, 8, 2, \ldots]$ appears over successive rows as $[18, 2, \ldots] \to [14, 2, \ldots] \to [8, 2, \ldots] \to [2, \ldots]$. The long blue feature is not one unchanged merged gap reused in several sequences; it is the decreasing suffix of one raw run. Rotating one compressed cell per row would show that block only once, but it would define a different dynamics from the Sieve Sequence.

View	Horizontal coordinate	What vertical comparison supports
Independent snapshots	Compressed units from each row's own head	Within-stage density and spacing texture; not cell-by-cell lineage
Shared-safe-2 aligned	Cumulative columns fixed by shared 2-gaps below the previous h^2	Exact alignment of the observed safe prefix; not full lineage beyond it

These figures provide empirical context for the placement problem; neither is evidence for

the companion-model survival thresholds derived below. They are generated by the gap-heatmap calculation from a dataset containing the first 100,000 gaps of each stage.

The central question is therefore not what fixed percentage of behavior may be called adversarial. It is how large the realized local destruction may be relative to the random benchmark $2/r$, and how that damage is allocated near the head.

The balanced companions are designed to separate them. They retain the real sieve's exact number of descendants but replace the arithmetic rule selecting which descendants die. This makes global survival identical in every companion while allowing local behavior to range from maximally protective, through position-blind random, to maximally hostile.

We establish:

- the exact allocation-independent global recurrence $N_{k+1} = (r_k - 2)N_k$;
- the cumulative local-hazard law $P(Q) = e^{-D(Q)}$ and its fixed-factor and logarithmic survival frontiers;
- the distinct square-window and head thresholds for the two mixture models;
- sharp finite allocation bounds separating adversarial-label budget from positional information; and
- exact-quota and biased exact-quota companions that preserve CRT strike counts while randomizing their locations.

1.1 Scope and Evidence

We prove exact identities for the finite companion processes and conditional theorems for their asymptotic local behavior. Whenever a result needs spatial uniformity, head availability, or cross-layer mixing, we state that premise in the property itself before using it. The comparison with the real sieve then shows what additional arithmetic information would transfer the companion result.

The evidence in the article has four different levels:

1. **Exact companion identities.** Recurrences such as $N_{k+1} = (r_k - 2)N_k$ are finite model identities. They do not use an asymptotic approximation.
2. **Conditional asymptotic derivations.** The survival frontiers follow from the stated local-hazard, placement, availability, and mixing premises. The formulas explain what those premises imply; they do not prove the premises for the real sieve.
3. **Empirical context.** The heatmaps and finite comparisons illustrate the placement problem in computed sieve data. They motivate the questions but are not used as proof of the asymptotic thresholds.
4. **Formal-verification status.** The companion-process derivations are mathematical proof records in this article. Stainless verification for this article remains pending, unlike the earlier sieve-sequence and gap-dynamics source properties it cites.

The companion theorems below are proved mathematically under their stated premises; Stainless verification remains pending and is outside this article's scope.

2 Preliminaries and Companion Models

Let $\mathcal{G}_k$ be the 2-gap descendants before installing prime r. Each parent $g \in \mathcal{G}_k$ produces the indexed copies

$$(g, 0), (g, 1), \ldots, (g, r-1).$$

Exactly two distinct indices are harmful. Each parent receives one of three policies. A **random parent** draws the harmful pair uniformly from the two-element subsets of $\mathbb{Z}/r\mathbb{Z}$. An **adversarial parent** places a deletion on its target child whenever possible. A **protective parent**, defined fully in Section 5.2, places both deletions away from the target whenever possible. Every policy leaves exactly $r-2$ children. The companions therefore change the location of the deletions, not the population size. For example, let $r = 5$ and suppose child index 1 is the target. A random parent may remove any pair, such as $\{0, 4\}$. An adversarial parent chooses a pair containing 1, such as $\{1, 4\}$. A protective parent chooses both indices outside the target, again allowing $\{0, 4\}$. The three parents make different local choices, but each leaves exactly three children. This simple example is the distinction used throughout the article: global reproduction is fixed, while local placement changes. The random, adversarial, and protective companion definitions above are the definitions used throughout this article. The corresponding real modular pair is derived in Gap Dynamics §6.1 [4]. The exact-quota process introduced in Section 7 is a separate conditional random-location experiment. It draws a fixed number of strikes from an eligible population and can strike zero, one, or two descendants of a given 2-gap parent. It therefore does not satisfy the balanced per-parent $r-2$ recurrence; its role is to compare conditional local survival at a fixed strike quota.

2.1 Notation

We use the following notation throughout:

Symbol	Meaning
r	incoming filter prime
Q	target prime head
N_k	complete-period 2-gap population at layer k
f_r	conditional destruction probability of an eligible tracked lineage
$\widehat{f}_r$	observed destruction fraction in a specified finite population
$w_r = r f_r/2$	destruction relative to the random benchmark $2/r$
$D(Q) = \sum_{r<Q} -\log(1-f_r)$	cumulative local hazard
α_r	scheduled absolute adversarial share in a mixture
$A(Q) = \sum_{r<Q} -\log(1-\alpha_r)$	cumulative adversarial-share hazard
$J_r/N_r = u_r$	exact-quota strike fraction
β_r	raw endpoint preference in a biased quota
κ_r	effective destruction skew, equal to w_r when measured from f_r

We write $F(Q) \asymp G(Q)$, sometimes read as $F(Q) \simeq G(Q)$ up to constant factors, when there are constants $0 < c \le C < \infty$ and a threshold Q_0 such that

$$cG(Q) \le F(Q) \le CG(Q) \qquad \text{for all } Q \ge Q_0.$$

Thus $\asymp$ describes the eventual growth or decay scale, not an exact numerical equality at a fixed value of Q. Unless a different range is displayed, filter products and sums run over primes $r_0 \le r < Q$, with a fixed $r_0 \ge 5$. Every finite prefix has strictly positive survival factors. Tail schedules are

imposed only once their hazards lie in $[0, 1)$. A lethal prefix cannot be absorbed into a positive constant. For probability statements, all candidate events are defined on one probability space. For each candidate, f_r is its deterministic conditional probability of destruction given initial eligibility and survival through the earlier filters. The common-hazard model assumes these probabilities are the same for the candidates being compared. The chain rule then gives a candidate survival probability $P(Q)$; independence between candidates is a separate issue. Window applications assume $B(Q) \asymp Q^2$ eligible candidate histories, each with survival probability $P(Q)$. Head applications assume an eligibility probability b_Q bounded below by a positive constant and the same survival law conditional on eligibility. Consequently,

$$\begin{aligned} \lambda_Q := \mathbb{E}[X_Q] = B(Q)P(Q) \quad & \text{[Linearity Of Expectation]}, \\ \Pr(H_Q) = b_Q P(Q) \quad & \text{[Conditional Probability]}. \end{aligned}$$

These marginal assumptions are additional to the balanced branching rule. They are not supplied by its global count. Different window and head experiments need not describe one nested random sequence of integers. For the prime-indexed head events H_Q, define

$$S(X) := \sum_{\substack{Q \le X \\ Q \text{ prime}}} \Pr(H_Q).$$

Throughout this article, **adequate cross-layer mixing** means that whenever $S(X) \longrightarrow \infty$,

$$\sum_{\substack{P,Q \le X \\ P,Q \text{ prime}}} \Pr(H_P \cap H_Q) = (1 + o(1))S(X)^2.$$

Mutual independence is a sufficient special case: its diagonal correction is $O(S(X)) = o(S(X)^2)$. The Kochen–Stone form of the second Borel–Cantelli lemma then gives $\Pr(H_Q \text{ infinitely often}) = 1$ [2]. When $S(X)$ converges, the first Borel–Cantelli lemma gives only finitely many head events without any mixing premise. Square-window applications use one further named premise. **Blind placement** means that the surviving starts are placed in the window so that the empty-window bound

$$\Pr(X_Q = 0) \le e^{-\lambda_Q}$$

holds, where λ_Q is the expected surviving population of that window (for example λ_Q^{mix} in Section 4.1). This is an assumption about the joint placement distribution: it holds for independent uniform placement and is not derived in this article for dependent allocators such as exact quotas, whole-filter coins, or block balance. Whenever a square-window result uses it, the result says so; the same caveat as Section 9 applies — allocators with different dependence structures cannot inherit one another's almost-sure conclusions.

2.2 Mathematical Foundation

The companion construction uses three exact sieve-sequence results proved in Gap Dynamics [4]: - the exact complete-period 2-gap count; - the two harmful copy-index classes; and - the exact accepted-strike count. The relative local-damage normalization is defined directly in Section 3.2, and its allocation refinement is defined in Section 6.2.

3 Relative Hazard and Survival Frontiers

This section is the article's main calculation spine. It proceeds in four steps:

1. keep the global population count separate from local placement;
2. normalize local damage by the random benchmark $2/r$;
3. multiply one-lineage survival factors and rewrite the product as $P(Q) = e^{-D(Q)}$;
4. compare the resulting survival scale with either a quadratic square-window supply or a single distinguished head position.

The same hazard calculation supports both the fixed-factor result and the logarithmic thresholds. The difference between the square-window and head frontiers comes from available supply: a square window has about Q^2 possible starts, while the head has only one distinguished position.

3.1 Global Persistence Is Independent of Allocation

We begin with the property shared by every balanced companion. No choice of the harmful pair changes the number of surviving descendants: random, adversarial, protective, and mixed parents all leave the same global population. Local extinction must therefore come from placement rather than from exhausting the complete-period supply. Let $N_k = |\mathcal{G}_k|$ and assume $N_0 > 0$. Installing r_k gives

$$\begin{aligned} N_{k+1} &= \sum_{g \in \mathcal{G}_k} (r_k - 2) && \text{[Exactly Two Copies Removed Per Parent]} \\ &= (r_k - 2) N_k. && \text{[Simplification]} \end{aligned}$$

Consequently,

$$\begin{aligned} N_k &= N_0 \prod_{i<k} (r_i - 2) && \text{[Iteration]} \\ &> 0 && [N_0 > 0;\ r_i \geq 5] \\ &\longrightarrow \infty. && \text{[Every Factor Is At Least 3]} \end{aligned}$$

Thus

global 2-gap persistence holds for every adversarial schedule. [Q.E.D.]

The complete proof record appears in Appendix A.1. The corresponding real-sieve count is proved in Gap Dynamics §5.2 [4].

3.2 Local Destruction Relative to Random

Two quantities must be distinguished. If $L_r > 0$ gaps are present in a specified finite population and H_r are destroyed, its observed fraction is

$$\widehat{f}_r := \frac{H_r}{L_r}, \qquad \widehat{w}_r := \frac{r \widehat{f}_r}{2}.$$

This first display is only a measurement convention: it says how to score a finite dataset after the fact. The next displays define the model quantity used in the proofs. The point is to compare any conditional destruction rate with the neutral random rate for the same incoming filter.

The probability model instead uses the conditional lineage hazard f_r defined in Section 2.1. Balanced random selection has conditional destruction rate

$$d_r := \frac{2}{r}.$$

The dimensionless worse-than-random factor is

$$w_r := \frac{f_r}{d_r} = \frac{r f_r}{2}.$$

So w_r is not a new probability. It is a ratio: $w_r = 1$ means "exactly as destructive as random at this filter," while $w_r > 1$ means "worse than random" by that factor.

This is the meaningful adversariality scale because the benchmark itself shrinks as filters grow:

$$\begin{aligned} w_r = 0 &\iff f_r = 0 \quad \text{[Protective Endpoint]}, \\ w_r = 1 &\iff f_r = 2/r \quad \text{[Random Benchmark]}, \\ w_r = r/2 &\iff f_r = 1 \quad \text{[Complete Local Destruction]}. \end{aligned}$$

The display proves the calibration of the scale. Zero means perfect local protection, one means the random benchmark, and $r/2$ is the largest possible score because it corresponds to certain destruction.

The range $0 \le w_r \le r/2$ compares conditional damage with the neutral law. The analogous observed score $\widehat{w}_r$ can fluctuate above one even for a random filter. An observed fraction does not by itself establish the conditional probability of destruction of a particular child.

Absolute Adversarial/Random Share as a Specialization

Fix a target: either a square-safe window or one distinguished head position. At filter r, let

$$0 \le \alpha_r \le 1$$

be the adversarial share. The remaining share $1 - \alpha_r$ uses the balanced random choice. This can be interpreted parent by parent or as a marginal mixture for one locally relevant lineage. For the cumulative product studied below, mixture choices for a tracked lineage are independent from one filter to the next. The one-lineage calculation at each filter is then the same:

- under adversarial selection, its target child is destroyed;
- under balanced random selection, that child survives with probability $1 - 2/r$.

The model therefore assumes that the adversarial branch is strong enough to identify and kill the locally relevant child. That is exactly what makes it an adversarial comparison rather than a description of the real filter.

The next block computes the effective hazard of that mixture. It has two pieces: the adversarial share contributes certain destruction, and the remaining random share contributes the neutral $2/r$ destruction rate.

$$\begin{aligned} f_r &= \alpha_r + (1 - \alpha_r)\frac{2}{r}, \\ w_r &= \frac{r}{2} f_r \\ &= 1 + \frac{r-2}{2}\alpha_r. \end{aligned}$$

This proves why the absolute-share parameter is dangerous. A fixed $\alpha_r = \alpha > 0$ does not mean "a fixed amount worse than random"; after normalization by $2/r$, it makes w_r grow linearly like $\alpha r/2$. This is why the fixed-share model is asymptotically fatal for an essentially trivial reason; the nontrivial question is how rapidly w_r itself may grow.

3.3 The General Cumulative Local-Hazard Law

Follow one eligible candidate through successive filters. By the conditional hazard definition, its next survival probability is

$$s_r = 1 - f_r = 1 - \frac{2w_r}{r}.$$

Assume $f_r < 1$ for every filter in the tracked chain. The purpose of the next definition is to turn a product of many small survival losses into a sum. That sum is the cumulative local hazard:

$$D(Q) := \sum_{r<Q} -\log(1 - f_r) = \sum_{r<Q} -\log\left(1 - \frac{2w_r}{r}\right).$$

With this definition, the complete survival factor is exactly the exponential of minus the accumulated hazard:

$$\begin{aligned} P(Q) &= \prod_{r<Q} (1 - f_r) && \text{[Conditional Probability Chain Rule]} \\ &= \exp\left(\sum_{r<Q} \log(1 - f_r)\right) && \text{[Product To Sum]} \\ &= e^{-D(Q)}. && \text{[Definition Of } D(Q)\text{]} \end{aligned} \qquad \text{[Q.E.D.]}$$

This proves the reusable translation used throughout the article: survival products become $P(Q) = e^{-D(Q)}$. The identity is exact inside the stated conditional model; all later threshold arguments come from estimating the size of $D(Q)$.

There is also a deterministic version. For a fixed nested cohort with no births or immigration, $L_{r^+} = L_r(1 - \widehat{f}_r)$ telescopes to the ratio of final to initial cohort size. This is an observed ratio, not a marginal probability. Changing the target window or introducing new descendants breaks that telescope unless an additional accounting identity is proved. The Prime Number Theorem, the prime harmonic estimate, and their partial-summation consequences used here and below are classical; we use Hardy and Wright [1]. For the random benchmark $w_r = 1$,

$$\begin{aligned} D_{\text{random}}(Q) &= \sum_{r<Q} -\log\left(1 - \frac{2}{r}\right) \\ &= 2\log\log Q + O(1), \end{aligned}$$

and therefore

$$P_{\text{random}}(Q) \asymp \frac{C}{(\log Q)^2}.$$

So the random benchmark has survival on the order of $(\log Q)^{-2}$. This is the baseline against which the later worse-than-random schedules are compared.

The absolute adversarial/random mixture from Section 3.2 is recovered because

$$1 - f_r = (1 - \alpha_r)\left(1 - \frac{2}{r}\right).$$

If

$$A(Q) := \sum_{r<Q} -\log(1 - \alpha_r),$$

then

$$D_{\text{adversarial/random}}(Q) = D_{\text{random}}(Q) + A(Q),$$
$$P_{\text{adversarial/random}}(Q) \asymp \frac{C}{(\log Q)^2} e^{-A(Q)}.$$

Thus the earlier $A(Q)$ is an excess hazard created by one particular policy mixture. The primary quantity is $D(Q)$, which also applies when no policy label α_r exists. If one filter has $f_r = 1$, local extinction is immediate and the cumulative hazard is infinite from that point. The complete proof record appears in Appendix A.2.

3.4 Every Fixed Finite Worsening Factor Survives

The random destruction rate shrinks like $2/r$. We first ask what happens when the local filter is a fixed number of times worse than that benchmark. With a quadratic supply of eligible starts and blind placement, every fixed factor still leaves occupied square windows. If head candidates remain available and successive layers mix adequately, the head also returns to a 2-gap infinitely often. Let $w \geq 0$ be fixed and suppose

$$f_r = \frac{2w}{r}$$

for all sufficiently large filters. The calculation below asks only for the growth scale of the hazard; any finite prefix is absorbed into a positive constant. Since the quadratic error terms are summable over primes,

$$\begin{aligned} D_w(Q) &= \sum_{r<Q} -\log\left(1 - \frac{2w}{r}\right) && \text{[Definition Of } D(Q)\text{]} \\ &= 2w \sum_{r<Q} \frac{1}{r} + O(1) && \text{[Taylor Expansion; Summable Remainder]} \\ &= 2w \log\log Q + O(1). && \text{[Prime Harmonic Sum]} \end{aligned}$$

Therefore

$$P_w(Q) \asymp \frac{C_w}{(\log Q)^{2w}}.$$

This is the first important frontier calculation. A fixed factor w changes only the power of $\log Q$; it does not introduce any power of Q.

For a square window with $B(Q) \asymp C_0 Q^2$ eligible lineages,

$$\lambda_w(Q) \asymp C_0 \frac{Q^2}{(\log Q)^{2w}} \longrightarrow \infty$$

for every finite w. The quadratic supply beats every fixed logarithmic loss. The growth is strong enough to make the standard empty-window bound summable, so only finitely many square windows are empty almost surely under the blind-placement premise. For a distinguished head with baseline availability bounded below,

$$\Pr(H_Q) \asymp \frac{C_w}{(\log Q)^{2w}}.$$

The sum of this probability over prime heads diverges for every finite w. Under adequate mixing, head 2-gaps therefore recur infinitely often almost surely. Thus

there is no finite constant-factor maximum worse than random. [Q.E.D.]

A filter that is twice, ten times, or one million times worse than the random rate still lies in the same asymptotic survival class once r is sufficiently large. The nontrivial transition begins only when w_r grows with r. The complete proof record appears in Appendix A.3. The fixed-factor conclusion is not only asymptotic; it is visible in the square-window occupancy itself. The figure below plots $\log_{10} \lambda_w(Q)$ against $\log_{10} Q$ for the fixed factors $w = 1, 3, 6, 10$, a constant 1% adversarial share, and the $c = 1$ frontier $w_r = 1 + \log r$. Every fixed finite w climbs without bound – $w = 6$ and $w = 10$ visibly dip first, because Q^2 must first outgrow $(\log Q)^{2w}$ – while the constant share collapses rapidly and the exact $c = 1$ boundary declines only logarithmically: $\lambda_1(Q) \asymp C/(\log Q)^2 \to 0$. This is the failure-side boundary derived in Section 3.5, not a surviving curve.

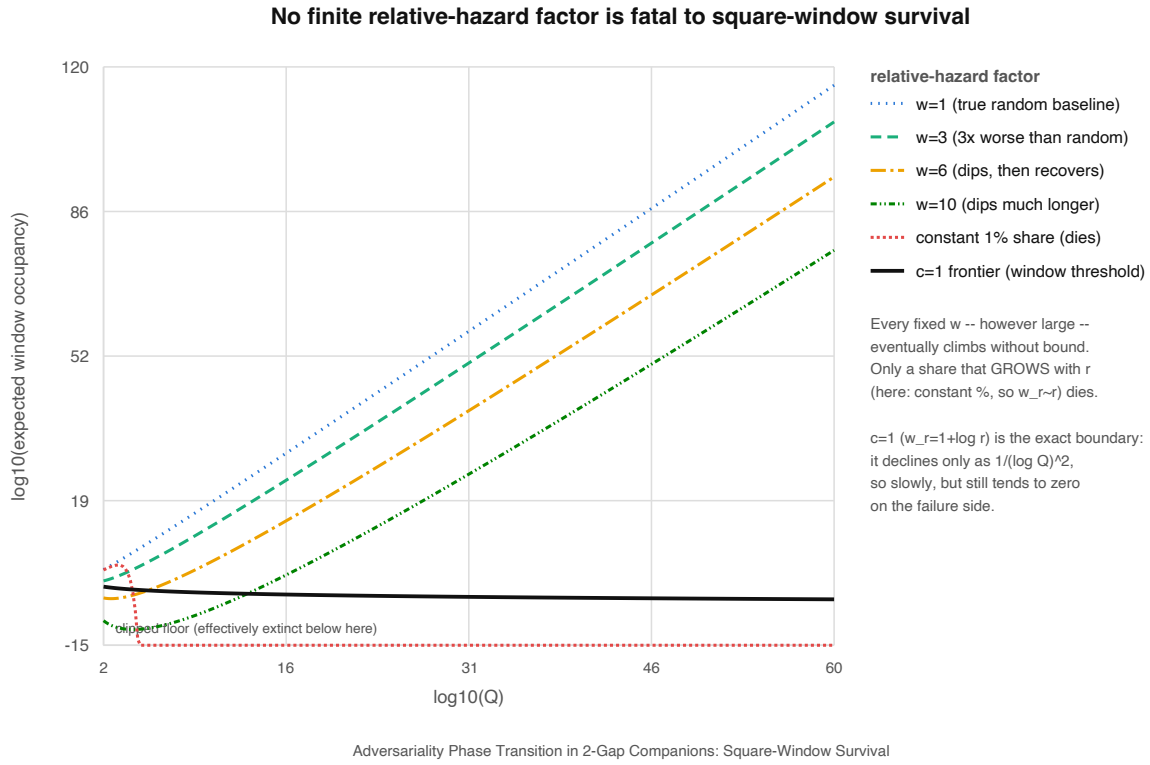


Figure 1: Square-window expected occupancy log10(lambda(Q)) on a log scale: every fixed relative-hazard factor w=1,3,6,10 eventually climbs without bound, a constant 1% adversarial share collapses rapidly, and the exact c=1 boundary declines slowly to zero

3.5 Logarithmically Growing Worsening Has Two Thresholds

The first genuine transition appears when the worsening factor grows with the filter. Using the same supply, availability, and mixing premises as Section 3.4, we let the factor grow logarithmically and compare the reserve supplied by a square window with the much thinner reserve at one distinguished head.

$$w_r = 1 + c \log r, \qquad c \geq 0.$$

The total local destruction rate is

$$f_r = \frac{2w_r}{r} = \frac{2}{r} + 2c\frac{\log r}{r}.$$

The random contribution supplies the first term, while the second is the growing excess. Prime summation gives

$$\begin{aligned} D_c(Q) &= \sum_{r<Q} -\log(1-f_r) && \text{[Definition Of } D(Q)\text{]} \\ &= 2\sum_{r<Q}\frac{1}{r} + 2c\sum_{r<Q}\frac{\log r}{r} + O(1) && \text{[Substitution; Summable Remainder]} \\ &= 2\log\log Q + 2c\log Q + O(1). && \text{[Prime-Sum Asymptotics]} \end{aligned}$$

Hence

$$P_c(Q) \asymp \frac{C_c}{Q^{2c}(\log Q)^2}. \qquad \text{[Exponentiation And Simplification]}$$

For a quadratic square-window supply,

$$\lambda_c(Q) \asymp C_0\frac{Q^{2-2c}}{(\log Q)^2}.$$

Therefore

$$\begin{aligned} c < 1 &\Longrightarrow \text{eventually nonempty square windows almost surely,} \\ c \geq 1 &\Longrightarrow \text{square-window expectation tends to zero.} \end{aligned}$$

For the head,

$$\Pr(H_Q) \asymp \frac{C_c}{Q^{2c}(\log Q)^2}.$$

Summing over prime heads has the same convergence behavior as

$$\int^{\infty} \frac{dx}{x^{2c}(\log x)^3}.$$

Thus

$$\begin{aligned} c < \frac{1}{2} &\Longrightarrow \text{infinitely many head events almost surely, with mixing,} \\ c \geq \frac{1}{2} &\Longrightarrow \text{only finitely many head events almost surely.} \end{aligned}$$

The threshold is the Borel-Cantelli decision rule for the head, and the figure below evaluates it directly. It plots the cumulative sum of $\Pr(H_Q)$ over real enumerated primes up to Q, for $w_r = 1 + c\log r$ at $c = 0.0, 0.1, 0.3, 0.5, 0.7, 1.0$. Below the threshold the sum keeps climbing – $c = 0.0$ and $c = 0.1$ clearly, $c = 0.3$ more slowly but provably – so there are infinitely many head events with mixing. At and above the threshold the sum flattens: $c = 0.5$ only very slowly (it is the boundary itself), $c = 0.7$ and $c = 1.0$ quickly – so there are only finitely many, almost surely.

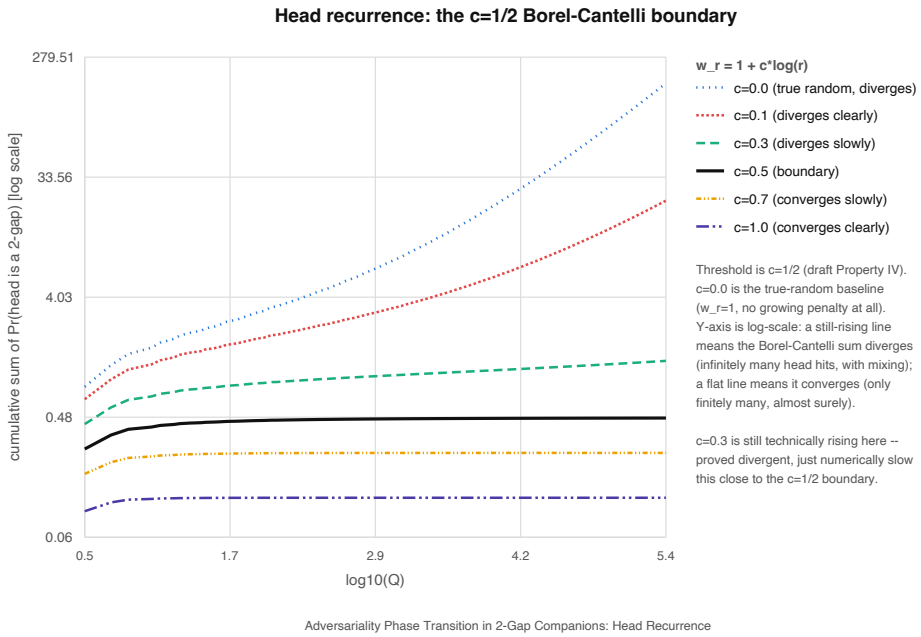


Figure 4: Cumulative sum of Pr(head is a 2-gap) over enumerated primes, log scale: c=0.0 and c=0.1 climb the whole way, c=0.3 climbs slowly, c=0.5 flattens only very slowly at the boundary, and c=0.7 and c=1.0 flatten quickly – the c=1/2 Borel-Cantelli threshold

Equivalently, the robust relative-factor regimes are

$$w_r < (1-\varepsilon)\log r \quad \text{[Square-Window Survival]},$$
$$w_r < \left(\frac{1}{2}-\varepsilon\right)\log r \quad \text{[Head Recurrence]},$$

up to the asymptotically negligible additive random baseline. In terms of the total segment destruction fraction,

$$f_r < (2-\varepsilon)\frac{\log r}{r} \quad \text{[Square-Window Survival]},$$
$$f_r < (1-\varepsilon)\frac{\log r}{r} \quad \text{[Head Recurrence]}. \qquad \text{[Q.E.D.]}$$

These are cumulative asymptotic regimes, not pointwise allowances that reset at each filter. Irregular schedules must be evaluated through $D(Q)$. The complete proof record appears in Appendix A.4.

3.6 Relative-to-Random Phase Diagram

The answer is not a maximum fixed percentage. It is a growth-rate boundary for the realized local damage relative to the random benchmark.

Realized relative factor	Total local destruction	Square windows	Head 2-gaps
$w_r = 1$	$2/r$	Eventually nonempty almost surely	Infinitely often with mixing

Any fixed finite $w_r = w$	$2w/r$	Eventually nonempty almost surely	Infinitely often with mixing
$w_r = 1 + c\log r$, $0 \le c < 1/2$	$2/r + 2c\log r/r$	Eventually nonempty almost surely	Infinitely often with mixing
$w_r = 1 + c\log r$, $1/2 \le c < 1$	$2/r + 2c\log r/r$	Eventually nonempty almost surely	Only finitely often almost surely
$w_r = 1 + c\log r$, $c \ge 1$	$2/r + 2c\log r/r$	Expected population tends to zero	Only finitely often almost surely
$f_r = 1$ at a tracked step	1	That tracked cohort is lost	That candidate is lost

Consequently, there is **no largest finite constant multiple of random**. For square-window survival, the filter may become almost $\log r$ times worse than random; for infinitely recurring head 2-gaps, it may become almost $\frac{1}{2}\log r$ times worse. In total local-destruction terms, the robust sufficient regimes are respectively

$$f_r < (2-\varepsilon)\frac{\log r}{r} \qquad \text{and} \qquad f_r < (1-\varepsilon)\frac{\log r}{r}.$$

These conclusions concern damage realized inside the tracked segment. A small global adversarial budget can still cause $f_r = 1$ if it is allocated with enough target information; the allocation theorem in Section 5 isolates that second axis.

4 Absolute-Share Mixtures

This section uses the same external analytic and probabilistic inputs as the relative-hazard calculation: Prime Number Theorem estimates and prime-sum consequences from Hardy and Wright [1], and Borel-Cantelli style almost-sure conclusions in the form stated in Section 2.1, with the Kochen–Stone reference [2] for the divergent mixed case.

4.1 Adversarial/Random Parent Square-Window Boundary

We now express the general hazard result through an explicit mixture. Each parent is adversarial with share α_r and otherwise random. When the surviving starts follow the spatial-uniformity model of the balanced random companion, a square-safe window has length

$$L_Q \asymp Q^2.$$

The expected mixed population is

$$\begin{aligned} \lambda_Q^{\mathrm{mix}} &= L_Q \delta_Q^{\mathrm{mix}} && \text{[Expected Uniform Occupancy]} \\ &\asymp C\frac{Q^2}{(\log Q)^2} e^{-A(Q)}. && \text{[Substitution]} \end{aligned}$$

Taking logarithms exposes the threshold:

$$\log \lambda_Q^{\mathrm{mix}} = 2\log Q - 2\log\log Q - A(Q) + O(1). \quad \text{[Logarithm]}$$

Therefore, for every fixed $\varepsilon > 0$,

$$\begin{aligned} A(Q) \le (2-\varepsilon)\log Q &\Longrightarrow \lambda_Q^{\mathrm{mix}} \longrightarrow \infty, && \text{[Subcritical Adversarial Budget]} \\ A(Q) \ge (2+\varepsilon)\log Q &\Longrightarrow \lambda_Q^{\mathrm{mix}} \longrightarrow 0. && \text{[Supercritical Adversarial Budget]} \end{aligned}$$

The boundary $A(Q) = 2\log Q + o(\log Q)$ requires its lower-order terms; the $-2\log\log Q$ contribution cannot be discarded there. Under uniform placement, an empty-window estimate has the usual form

$$\Pr(X_Q = 0) \le e^{-\lambda_Q^{\mathrm{mix}}}.$$

Whenever

$$\sum_{Q \text{ prime}} e^{-\lambda_Q^{\mathrm{mix}}} < \infty,$$

the first Borel-Cantelli lemma gives only finitely many empty square windows almost surely. A convenient sufficient condition is

$$\lambda_Q^{\mathrm{mix}} \ge (1+\varepsilon)\log Q \qquad \text{[Q.E.D.]}$$

for all sufficiently large Q. This is stronger than merely requiring $\lambda_Q^{\mathrm{mix}} \to \infty$ and prevents a slow divergent expectation from being mistaken for an eventual-survival theorem. This proves eventual safe-window occupancy inside the spatially uniform mixed companion. Section 8 states the separate conditions needed to transfer the result to the real sieve. The complete proof record appears in Appendix A.5.

4.2 Why a Constant Absolute Adversarial Share Is Locally Fatal

A constant adversarial share sounds mild, but it adds the same positive loss at every filter while the random benchmark keeps shrinking. We therefore expect it to overwhelm local survival. Let one fixed share $0 < \alpha < 1$ be adversarial at every filter. Then

$$\begin{aligned} A(Q) &= -\big(\pi(Q) + O(1)\big)\log(1-\alpha) && \text{[Constant Share]} \\ &\asymp \big[-\log(1-\alpha)\big]\frac{Q}{\log Q}. && \text{[Prime Number Theorem]} \end{aligned}$$

Since $Q/\log Q$ grows faster than $\log Q$, this lies far above the square-window critical budget. Hence

$$\begin{aligned} \lambda_Q^{(\alpha)} &\asymp C\frac{Q^2}{(\log Q)^2}(1-\alpha)^{\pi(Q)} \\ &\longrightarrow 0. && \text{[Exponential Loss Beats Quadratic Growth]} \end{aligned}$$

Thus

every fixed positive per-filter adversarial share is locally fatal [Q.E.D.]

in the repeated-mixture model, even though the complete-period population continues to grow without bound. This is different from applying one adversarial dilution after all random filters have finished. A one-time dilution multiplies the final count by $1-\alpha$ once; the repeated model multiplies it once per prime. Confusing these two experiments reverses the asymptotic conclusion.

4.3 Two Decaying Absolute-Share Families

The useful question is therefore not "what fixed percentage is tolerable?" The useful question is how quickly α_r must decay. In this section and the comparison in Section 5, the displayed schedules hold exactly for all sufficiently large filters. Mere asymptotic equivalence gives weaker remainder control and does not determine the critical cases. The estimates below use the Taylor expansion of $-\log(1-x)$ with summable remainders, together with the prime harmonic estimate and Prime Number Theorem partial-summation consequences; we use Hardy and Wright [1] for these classical analytic number theory inputs.

Reciprocal Decay: $\alpha_r = c/r$

For fixed $c > 0$ and sufficiently large primes,

$$\begin{aligned} A(Q) &= c \sum_{r<Q} \frac{1}{r} + O(1) && \text{[Taylor Expansion; Summable Error]} \\ &= c \log\log Q + O(1). && \text{[Prime Harmonic Sum]} \end{aligned}$$

Therefore

$$e^{-A(Q)} \asymp \frac{1}{(\log Q)^c}$$

and

$$\lambda_Q^{\text{mix}} \asymp C \frac{Q^2}{(\log Q)^{2+c}} \longrightarrow \infty.$$

The window population grows polynomially faster than its logarithmic losses, so the empty-window probabilities are summable under the spatial model.

Logarithmic-Over-Linear Decay: $\alpha_r = c \log r/r$

For a finite initial prefix, define the shares separately so that they remain in $[0, 1)$; this changes only the positive final constant. On the asymptotic tail,

$$\begin{aligned} A(Q) &= c \sum_{r<Q} \frac{\log r}{r} + O(1) && \text{[Taylor Expansion; Summable Error]} \\ &= c \log Q + O(1). && \text{[Prime Number Theorem By Partial Summation]} \end{aligned}$$

Consequently,

$$\begin{aligned} e^{-A(Q)} &\asymp Q^{-c}, \\ \lambda_Q^{\text{mix}} &\asymp C \frac{Q^{2-c}}{(\log Q)^2}. \end{aligned}$$

The square-window phase diagram is therefore

$$\begin{aligned} c < 2 &\Longrightarrow \lambda_Q^{\text{mix}} \longrightarrow \infty, \\ c \geq 2 &\Longrightarrow \lambda_Q^{\text{mix}} \longrightarrow 0. \end{aligned}$$

For $c < 2$, the divergence is polynomial, so the empty-window bound is summable and every sufficiently large square window is nonempty almost surely under the spatial-uniformity premise.

4.4 Adversarial/Random Parent Head Boundary

The head contains only one distinguished position, so it receives no quadratic window reserve. Under uniform head marginals, its occurrence probability is the surviving local density itself. To turn a divergent sum of these probabilities into almost-sure recurrence, we also require independence or a sufficiently strong weak-mixing substitute.

$$\Pr(H_Q) \asymp \delta_Q^{\text{mix}} \asymp \frac{C}{(\log Q)^2} e^{-A(Q)}.$$

Under adequate cross-layer mixing, the second Borel-Cantelli lemma gives

$$\sum_{Q \text{ prime}} \Pr(H_Q) = \infty \Longrightarrow H_Q \text{ occurs infinitely often almost surely.}$$

For $\alpha_r = c/r$,

$$\Pr(H_Q) \asymp \frac{C}{(\log Q)^{2+c}},$$

and the sum over prime Q diverges for every fixed c. Reciprocal decay is therefore compatible with infinitely many head events under mixing. For $\alpha_r = c \log r/r$,

$$\Pr(H_Q) \asymp \frac{C}{Q^c (\log Q)^2}.$$

Using prime density $dQ/\log Q$, the corresponding series has the same convergence behavior as

$$\int^{\infty} \frac{dx}{x^c (\log x)^3}.$$

Therefore

$$\begin{aligned} c < 1 &\Longrightarrow \sum_{Q \text{ prime}} \Pr(H_Q) = \infty, \\ c \geq 1 &\Longrightarrow \sum_{Q \text{ prime}} \Pr(H_Q) < \infty. \end{aligned} \qquad \text{[Q.E.D.]}$$

For $c < 1$, adequate mixing implies infinitely many head events almost surely. For $c \geq 1$, the first Borel-Cantelli lemma implies only finitely many head events almost surely; no independence assumption is needed for that convergent direction. The head threshold $c = 1$ is stricter than the safe-window threshold $c = 2$. There is an intermediate regime

$$1 \leq c < 2$$

in which square-safe windows remain populated almost surely under the spatial model, while head recurrence fails almost surely in the mixed companion.

4.5 Adversarial/Random Parent Phase Diagram

For the representative schedule $\alpha_r = c \log r/r$, the companion separates into three regimes:

Adversarial scale	Global 2-gaps	Square-safe windows	Head recurrence
$0 \leq c < 1$	Persist and grow	Eventually nonempty almost surely	Infinite almost surely, with mixing
$1 \leq c < 2$	Persist and grow	Eventually nonempty almost surely	Only finitely many almost surely
$c \geq 2$	Persist and grow	Mixed expectation tends to zero	Only finitely many almost surely
Fixed $\alpha > 0$	Persist and grow	Mixed expectation tends to zero	Only finitely many almost surely

The table's last two columns are statements inside the spatial model. The global column is unconditional for every balanced companion.

4.6 Why a Fixed Absolute Percentage Gives the Wrong Maximum

Within position-blind repeated mixtures, percentages that do not change with the filter prime have a blunt but secondary answer. If the same absolute adversarial share α is applied at every filter, every $\alpha > 0$ is eventually fatal to the local mixed baseline. In that restricted normalization,

$$\text{maximum sustainable fixed absolute adversarial share} = 0\%.$$

This is not the meaningful answer to "how much worse than random can the filter be?" Random destruction itself shrinks as $2/r$, while fixed $\alpha > 0$ adds a positive floor and makes the relative factor $w_r = 1 + (r-2)\alpha/2$ diverge linearly. The primary answer from Section 3.6 is instead that every fixed finite w survives the companion model defined above, with the first transition only when w_r grows on the order of $\log r$. The zero-percent statement concerns safe-window and head survival under this fixed absolute-share policy. It does not concern the global population, which survives even under 100% adversarial selection. Nonzero adversariality remains supportable when its share decreases with r. For a fixed margin $\varepsilon > 0$, the representative sufficient schedules are

$$\alpha_r \le (2-\varepsilon)\frac{\log r}{r} \quad \text{[Square-Window Regime]},$$
$$\alpha_r \le (1-\varepsilon)\frac{\log r}{r} \quad \text{[Head-Recurrence Regime, With Mixing]}.$$

Ignoring the required strict margin only to display the boundary curves, these become

$$\text{square-window boundary percentage} = 200\frac{\log r}{r}\%,$$
$$\text{head-recurrence boundary percentage} = 100\frac{\log r}{r}\%.$$

Here log is the natural logarithm. Representative values are:

Filter prime r	Square-window boundary	Head-recurrence boundary
101	9.14%	4.57%
1,000	1.38%	0.691%
19,000	0.104%	0.0519%
100,003	0.0230%	0.0115%

These entries are asymptotic boundary values, not independent allowances that reset at each filter. A schedule may temporarily cross a displayed value and remain viable if it spent less adversarial budget earlier; it may also fail despite staying below isolated entries if its cumulative behavior is worse on other filters. The governing quantities remain

$$A(Q) = \sum_{r<Q} -\log(1-\alpha_r)$$

for square windows and

$$\sum_{Q \text{ prime}} \frac{e^{-A(Q)}}{(\log Q)^2}$$

for head recurrence.

4.7 What "Percentage Adversarial" Must Specify

There is no unique mixture until we specify what receives the adversarial label.

Mixture	Choice made	Consequence
Parent level	Each parent is adversarial with probability α_r	Independent branching interpretation
Whole filter	The complete filter is adversarial with probability α_r	Same one-lineage marginal, stronger dependence between parents
One-time final	One adversarial dilution is applied after random filtering	One factor $1-\alpha$; no cumulative phase transition

The calculations in Sections 3–4 concern a repeated share at every filter. Their expectations apply to the first two interpretations because one lineage has the same marginal survival probability. Their almost-sure conclusions do not transfer automatically: a whole-filter choice coordinates all parents and therefore needs its own spatial or mixing argument. The one-time model answers a different question because it applies the loss only once.

5 Allocation and the Protective Parent

5.1 The Same Adversarial Percentage Can Produce Different Outcomes

This property is a capacity comparison. Suppose K parents may use the adversarial policy and L parents contribute a child to the target window. A target-aware allocator can clear the window exactly when its budget covers all the relevant parents:

$$K \geq L.$$

For a fixed adversarial share $\alpha = K/N > 0$, if the relevant fraction $L/N \longrightarrow 0$, then eventually

$$\alpha \geq \frac{L}{N},$$

which is the same condition as $K \geq L$. Thus a fixed adversarial percentage may be enough to suppress the head early and, after the target population becomes sparse enough, remove every 2-gap from the tracked window. The result depends on allocation: a position-blind allocator with the same percentage does not automatically select all L relevant parents. We now derive the complete finite range. The target window is shorter than the old period, so each parent contributes at most one child to it. Let - N be the total number of parents; - R be the set of parents with a child in target region W; - $L = |R|$; - $\mathcal{A}$ be the set of adversarial parents; and - $K = |\mathcal{A}|$. In this allocation experiment every parent outside $\mathcal{A}$ is protective: it preserves its target child. Thus the only target losses come from adversarial parents. A random complement would add further losses and would not satisfy the following survivor identity. The number destroyed is

$$H = |\mathcal{A} \cap R|,$$

and the number surviving is

$$S = L - H.$$

The intersection size obeys the sharp bounds

$$\begin{aligned} H &\le \min(K, L) && \text{[Intersection Cannot Exceed Either Set]}, \\ H &\ge \max(0, K - (N - L)) && \text{[Only } N - L \text{ Irrelevant Parents Exist]}. \end{aligned}$$

Substituting into $S = L - H$ gives

$$\max(0, L - K) \le S \le \min(L, N - K).$$

Both endpoints are attainable. A target-aware allocator spends its budget on R first:

$$S_{\text{targeted}} = \max(0, L - K).$$

A protective allocator spends the same budget on the $N - L$ irrelevant parents first:

$$S_{\text{protective}} = \min(L, N - K).$$

Between these endpoints, a position-blind allocator chooses a uniformly random size-K subset of the N parents. Then

$$H \sim \text{Hypergeometric}(N, L, K)$$

and

$$\begin{aligned} \mathbb{E}[H] &= \frac{KL}{N}, \\ \mathbb{E}[S] &= L\left(1 - \frac{K}{N}\right). \end{aligned}$$

When $K \ge L$, the exact probability of total local destruction is

$$\Pr(S = 0) = \frac{\binom{N-L}{K-L}}{\binom{N}{K}} = \frac{\binom{K}{L}}{\binom{N}{L}}.$$

Thus the same budget can produce complete protection, average proportional loss, or total local destruction. In particular,

$$S_{\text{targeted}} = 0 \iff K \ge L \iff \alpha \ge \frac{L}{N}. \qquad \text{[Q.E.D.]}$$

The three scales should not be confused. At the head, $L = 1$, so one correctly allocated adversarial parent kills the current head candidate. In a sparse tracked window, a fixed share clears the whole window once $K \ge L$. Neither statement erases the complete-period 2-gap population: every targeted parent still leaves $r-2$ other descendants outside the target, so the global recurrence from Section 3.1 continues to grow. Preventing future head or window 2-gaps therefore requires the allocator to repeat the targeted choice at later filters. The complete finite proof appears in Appendix A.6.

5.2 The Protective Parent Policy

The protective parent policy is the local opposite of the adversarial parent policy. It preserves a parent's target child whenever the exact-two deletion rule permits that choice. It does not create extra descendants and cannot change the global recurrence. For parent g, let $T_g(W)$ be the indices of its children in target region W. In the post-crossover regime,

$$|T_g(W)| \le 1.$$

Because $r \geq 5$, at least $r - 1 \geq 4$ child indices lie outside $T_g(W)$. The protective parent policy may therefore choose a harmful pair

$$K_{g,r}^{\text{protective}} \subseteq (\mathbb{Z}/r\mathbb{Z}) \setminus T_g(W), \qquad |K_{g,r}^{\text{protective}}| = 2.$$

The adversarial parent policy instead chooses a pair containing the target index whenever $T_g(W)$ is nonempty. Both policies remove exactly two children, so both leave $r - 2$ descendants globally. Their only difference is local placement:

$$T_g(W) \neq \varnothing \Longrightarrow \text{the protective parent preserves the target child,}$$
$$T_g(W) \neq \varnothing \Longrightarrow \text{the adversarial parent destroys the target child.}$$

The protective parent is an oracle comparison, not a plausible random filter. It is allowed to see the chosen target and place its two deletions elsewhere. Its purpose is to define the protective endpoint of the same balanced family in which the adversarial parent defines the pessimistic endpoint.

5.3 Fixed-Cohort Survival Under Adversarial/Protective Parent Mixing

We next alternate the two target-aware endpoint policies without letting the allocator inspect current positions. Consider N_0 locally relevant lineages from distinct initial parents, following one target child per parent through the chain. These histories do not share subsequent parents. At filter r, every surviving lineage independently becomes an adversarial parent with probability α_r or a protective parent with probability $1 - \alpha_r$. The adversarial policy destroys its target child; the protective policy preserves it. The calculation changes if the allocator may first observe which parents are locally relevant. One lineage survives the complete chain with probability

$$\begin{aligned} P_Q &= \prod_{r<Q} (1 - \alpha_r) && \text{[Survive Every Independent Filter Label]} \\ &= e^{-A(Q)}. && \text{[Definition Of } A(Q)\text{]} \end{aligned}$$

Independence between parent lineages then gives

$$X_Q \sim \text{Binomial}(N_0, P_Q),$$

so

$$\mathbb{E}[X_Q] = N_0 e^{-A(Q)},$$
$$\Pr(X_Q > 0) = 1 - \left(1 - e^{-A(Q)}\right)^{N_0}.$$

For one filter this reduces to

$$X_{k+1} \mid X_k = N \sim \text{Binomial}(N, 1 - \alpha_r),$$

with immediate wipeout probability α_r^N. Population redundancy is therefore useful under blind parent assignment: every relevant lineage must become an adversarial parent in the same transition to erase the cohort. If $\alpha_r = \alpha > 0$ is constant, then

$$P_Q = (1 - \alpha)^{\pi(Q) + O(1)} \longrightarrow 0.$$

Every one of the finite N_0 lineages eventually becomes an adversarial parent with probability one. Hence the fixed cohort becomes extinct almost surely even though every lineage continues to have

$r-2$ descendants elsewhere in the complete period. Compared with the adversarial/random mixture, the adversarial/protective law removes the random factor $1-2/r$:

$$\begin{aligned} s_r^{\text{adversarial/random}} &= (1-\alpha_r)\left(1-\frac{2}{r}\right), \\ s_r^{\text{adversarial/protective}} &= 1-\alpha_r. \end{aligned} \qquad \text{[Q.E.D.]}$$

This improvement is local. It does not overcome a fixed positive adversarial share repeated through infinitely many filters.

5.4 Growing Square Windows Under Adversarial/Protective Parent Mixing

The protective policy removes the balanced-random density penalty by preserving every eligible target child. Suppose the fully protective model supplies $B(Q) \asymp C_0 Q^2$ eligible lineages in the square window, while each history has the conditional survival law $e^{-A(Q)}$. The cumulative adversarial-label probability is then the only marginal local loss. From Section 5.3, each of the $B(Q)$ eligible lineages survives with probability $e^{-A(Q)}$. If their complete target histories are independent, then

$$X_Q^{\text{adversarial/protective}} \sim \text{Binomial}\left(B(Q), e^{-A(Q)}\right)$$

and

$$\begin{aligned} \lambda_Q^{\text{adversarial/protective}} &:= \mathbb{E}[X_Q^{\text{adversarial/protective}}] \\ &= B(Q)e^{-A(Q)} \\ &\asymp C_0 Q^2 e^{-A(Q)}. \end{aligned}$$

Under this additional independence hypothesis, the empty-window probability satisfies

$$\begin{aligned} \Pr(X_Q^{\text{adversarial/protective}} = 0) &= \left(1-e^{-A(Q)}\right)^{B(Q)} \\ &\le e^{-\lambda_Q^{\text{adversarial/protective}}}. \quad [1-x \le e^{-x}] \end{aligned}$$

Independent labels at individual parents do not establish independence of terminal histories that share ancestors. The binomial model here is an additional local experiment, not a consequence of balanced branching alone. For the occupancy conclusions below it suffices instead to assume the blind-placement bound $\Pr(X_Q = 0) \le e^{-\lambda_Q}$ directly; the expectation formula uses only the marginal survival law. Taking logarithms gives the phase boundary

$$\log \lambda_Q^{\text{adversarial/protective}} = 2\log Q - A(Q) + O(1).$$

Hence, for every fixed $\varepsilon > 0$,

$$\begin{aligned} A(Q) \le (2-\varepsilon)\log Q &\Longrightarrow \lambda_Q^{\text{adversarial/protective}} \longrightarrow \infty, \\ A(Q) \ge (2+\varepsilon)\log Q &\Longrightarrow \lambda_Q^{\text{adversarial/protective}} \longrightarrow 0. \end{aligned}$$

In the first regime the expectation grows polynomially, the empty-window probabilities are summable, and the first Borel-Cantelli lemma gives only finitely many empty square windows almost surely. For the representative schedule

$$\alpha_r = c\frac{\log r}{r},$$

we have $A(Q) = c \log Q + O(1)$ and therefore

$$\lambda_Q^{\text{adversarial/protective}} \asymp C_0 Q^{2-c}.$$

Thus

$$c < 2 \Longrightarrow \text{eventually nonempty square windows almost surely,}$$
$$c = 2 \Longrightarrow \text{critical order-one expected population,}$$
$$c > 2 \Longrightarrow \text{expected population tends to zero.}$$

The leading threshold $c = 2$ matches the adversarial/random companion, but the boundary term differs:

$$\lambda_Q^{\text{adversarial/random}} \asymp C \frac{Q^{2-c}}{(\log Q)^2}, \qquad \text{[Q.E.D.]}$$
$$\lambda_Q^{\text{adversarial/protective}} \asymp C_0 Q^{2-c}.$$

At $c = 2$, the random mixture tends to zero while the protective mixture retains only an order-one expectation, still insufficient for eventual almost-sure nonemptiness.

5.5 Head Recurrence Under Adversarial/Protective Parent Mixing

At the head, the protective policy can preserve an eligible lineage but cannot create one. Let b_Q be its availability probability and suppose $b_Q \geq b > 0$ for all sufficiently large Q. Conditional on availability, the lineage must avoid every adversarial assignment in its chain. Independence or adequate weak mixing between head events then supplies the recurrence step.

$$\Pr(H_Q) = b_Q e^{-A(Q)}.$$

The lower bound on b_Q makes the recurrence criterion equivalent, up to positive constants, to

$$\sum_{Q \text{ prime}} e^{-A(Q)}.$$

Under adequate cross-layer mixing, the second Borel-Cantelli lemma gives

$$\sum_{Q \text{ prime}} e^{-A(Q)} = \infty \Longrightarrow H_Q \text{ occurs infinitely often almost surely.}$$

If the series converges, the first Borel-Cantelli lemma gives only finitely many head events almost surely without any independence premise. For

$$\alpha_r = c \frac{\log r}{r},$$

we have $e^{-A(Q)} \asymp Q^{-c}$. The prime-head series therefore behaves like

$$\sum_{Q \text{ prime}} \frac{1}{Q^c}.$$

The prime harmonic series diverges at $c = 1$, while the series converges for $c > 1$. Hence

$$c \leq 1 \Longrightarrow \text{infinitely many head events almost surely, with mixing,}$$
$$c > 1 \Longrightarrow \text{only finitely many head events almost surely.}$$

The boundary differs from adversarial/random mixing. There the balanced-random head density contributes $(\log Q)^{-2}$:

$$\Pr(H_Q^{\text{adversarial/random}}) \asymp \frac{1}{Q^c(\log Q)^2},$$
$$\Pr(H_Q^{\text{adversarial/protective}}) \asymp \frac{b_Q}{Q^c}. \qquad \text{[Q.E.D.]}$$

At $c = 1$, the adversarial/random prime series converges, while the adversarial/protective series diverges. Thus the protective parent policy changes the inclusion of the critical boundary, even though both mixtures have the same leading threshold scale. For the gentler schedule $\alpha_r = c/r$, the occurrence probability is comparable to $(\log Q)^{-c}$ and the sum over prime heads diverges for every fixed finite c.

5.6 Parent-Mixture Comparison

Under their respective spatial premises, the two position-blind mixtures have the following asymptotic behavior:

Adversarial schedule	Adversarial/ random square window	Adversarial/ protective square window	Adversarial/ random head	Adversarial/ protective head
Fixed $\alpha > 0$	Expectation tends to zero	Expectation tends to zero	Finitely many almost surely	Finitely many almost surely
$\alpha_r = c/r$	Eventually nonempty almost surely	Eventually nonempty almost surely	Infinite with mixing	Infinite with mixing
$\alpha_r = c\log r/r$, $c < 1$	Eventually nonempty almost surely	Eventually nonempty almost surely	Infinite with mixing	Infinite with mixing
$\alpha_r = c\log r/r$, $c = 1$	Eventually nonempty almost surely	Eventually nonempty almost surely	Finitely many almost surely	Infinite with mixing
$\alpha_r = c\log r/r$, $1 < c < 2$	Eventually nonempty almost surely	Eventually nonempty almost surely	Finitely many almost surely	Finitely many almost surely
$\alpha_r = c\log r/r$, $c = 2$	Expectation tends to zero	Order-one expectation	Finitely many almost surely	Finitely many almost surely
$\alpha_r = c\log r/r$, $c > 2$	Expectation tends to zero	Expectation tends to zero	Finitely many almost surely	Finitely many almost surely

The protective parent policy removes the balanced-random $(\log Q)^{-2}$ loss. This does not change the leading square-window threshold $c = 2$, because the quadratic window dominates logarithmic factors away from the boundary. It does change the boundary behavior and, most visibly, includes $c = 1$ on the recurrent side of the head transition. Every entry assumes position-blind adversarial labels. A target-aware allocator is governed by Section 5.1 instead and may erase the head with one correctly placed adversarial label regardless of this table's percentage regime.

6 Allocation Mechanisms and Local Damage

6.1 Mechanism Families

An adversarial share becomes meaningful only after we say how policies are assigned to parents. A useful comparison begins with a position-blind allocator and then adds positional information in controlled steps. We write P for a protective parent and A for an adversarial parent.

Mechanism	Adversarial share	Position information	What it measures
Independent parent coin	Random around α_r	None	Simplest branching law
Exact-quota shuffle	Exactly K/N	None	Canonical finite-population baseline
Shuffled alternation, $P, A, P, A, \ldots$	Fixed by pattern	None	Balanced deterministic labels after shuffling
Block-balanced shuffle	Fixed inside each block	None	Sensitivity to local clustering
Random cyclic mask	Fixed by pattern	None	Sensitivity to periodic allocation
Position-blind hash	Random around α_r	None	Reproducible parent coins
Delayed adversary	Exactly K/N	Previous layer	Persistence of positional information
Noisy ranking	Exactly K/N	Tunable	Transition from blind to targeted allocation
Perfect adversary	Exactly K/N	Current layer	Worst-case endpoint

The exact-quota shuffle is the primary null model because it fixes the budget without using position. Shuffled patterns and block balance test whether local clustering changes the result. A delayed allocator can use only the previous layer, while noisy ranking assigns weight $e^{-\beta d_g}$ to a parent at distance d_g from the target. Thus $\beta = 0$ is uniform allocation and large β approaches the perfect adversary. Every mechanism assigns one policy to a parent; the chosen policy still removes exactly two of that parent's children. Alternating labels cannot restore a target child destroyed at an earlier filter, and an unshuffled periodic pattern may lock to the sieve geometry. We therefore compare mechanisms through their realized local damage and targeting strength rather than through the scheduled percentage alone.

6.2 Targeting and Local Hazard

The primary observed state space has two coordinates:

$$(w_r, \theta_r) = (\text{damage relative to random}, \text{realized targeting strength}).$$

The first says how much total damage the tracked segment actually received. The second says how concentrated the controllable adversarial-label budget was relative to the locally relevant parents. The scheduled share $\alpha_r = K_r/N_r$ remains an experimental input, but it is not itself the local damage.

Normalized Targeting Strength

For one nondegenerate transition, retain the notation from Section 5.1 and define

$$\begin{aligned} H_{\min} &= \max(0, K - (N - L)), \\ H_0 &= \frac{KL}{N}, \\ H_{\max} &= \min(K, L). \end{aligned}$$

These are the protective minimum, uniform-random mean, and adversarial maximum numbers of locally relevant parents hit. When $H_{\min} < H_0 < H_{\max}$, normalize the realized hit count H by

$$\theta(H) = \begin{cases} \dfrac{H - H_0}{H_0 - H_{\min}}, & H \le H_0, \\ \dfrac{H - H_0}{H_{\max} - H_0}, & H \ge H_0. \end{cases}$$

Then

$$\begin{aligned} \theta = -1 &\iff H = H_{\min} \quad \text{[Protective Endpoint]}, \\ \theta = 0 &\iff H = H_0 \quad \text{[Uniform Benchmark]}, \\ \theta = 1 &\iff H = H_{\max} \quad \text{[Targeted Endpoint]}. \end{aligned}$$

The center H_0 is an expectation and need not be an integer, so a finite realization need not attain $\theta = 0$ exactly. It remains the neutral reference point. Degenerate cases with a zero denominator should report the raw tuple (N, L, K, H) instead of assigning a synthetic score. In every case the local survivor count remains the exact observable

$$S = L - H.$$

The score measures realized placement, not hostile intent. A random shuffle may occasionally produce positive θ, and a nominal adversary with poor information may produce negative θ.

Realized Local Hazard

Let T_r be the total number of locally relevant target children destroyed by the complete filter, including both random-baseline and adversarial-label destruction. Define

$$f_r^{\text{local}} = \frac{T_r}{L_r}, \qquad w_r^{\text{local}} = \frac{rT_r}{2L_r}, \qquad L_r > 0.$$

In the pure adversarial/protective assignment, an adversarial label destroys its target and a protective label preserves it, so $T_r = H_r$. In the adversarial/random assignment, T_r also contains the random branch's $2/r$ baseline. For blind adversarial/random labels,

$$\mathbb{E}[f_r^{\text{local}}] \approx \alpha_r + (1 - \alpha_r)\frac{2}{r}.$$

A perfect adversary can make $f_r^{\text{local}} = 1$ even when the global budget $\alpha_r = K_r/N_r$ is tiny, provided its budget and information cover the local target. For a fixed nested cohort, the observed cumulative hazard is

$$\widehat{D}(Q) = \sum_{r<Q} -\log\left(1 - f_r^{\text{local}}\right),$$

whenever every factor is positive. If one transition has $f_r^{\text{local}} = 1$, the tracked local cohort is extinct and $\widehat{D}(Q)$ is effectively infinite from that point. This diagnostic generalizes $A(Q)$ and includes the random baseline rather than counting only the excess adversarial-label loss. The separation between scheduled α_r, realized w_r^{local}, and targeting score θ_r measures respectively policy budget, total relative damage, and the value of positional information. When the locally relevant population is redefined at every transition rather than following one cohort, $\widehat{D}(Q)$ is only a cumulative diagnostic; it is not an exact survival exponent for a single population. We can therefore read the earlier phase calculations with three levels of input: - $A(Q)$ is the scheduled budget under the blind-allocation model; - $\widehat{D}(Q)$ records observed damage along a specified cohort; and - θ_r records how strongly the controllable budget targeted the segment. Only the first has a closed form from α_r alone. The probability phase diagram uses conditional hazards $D(Q)$ and w_r. Applying it to the observed scores requires a theorem relating those scores to candidate marginals.

6.3 Comparing the Allocation Mechanisms

The mechanisms differ in how much they know about the target, so we compare them with the same strike budget and the same target region. Uniform shuffling is the neutral reference. Block balance and delayed information show whether dependence alone changes the result. Noisy ranking moves continuously toward the perfectly targeted endpoint.

Allocation	Information about the target	Role in the comparison
Uniform exact quota	None	Random baseline
Block-balanced quota	None, but locally dependent	Clustering test
Delayed allocation	Previous layer only	Memory test
Noisy ranking	Partial current information	Intermediate targeting
Perfect allocation	Complete current information	Adversarial endpoint

For each transition, the essential observation is the tuple

$$(N_r, L_r, K_r, H_r, T_r, w_r^{\text{local}}, \theta_r).$$

It records the total and locally relevant parents, the available adversarial budget, the number of relevant parents selected, the total local destruction, the damage relative to random, and the targeting score. Exact-quota companions also retain J_r, $u_r = J_r/N_r^{\text{strike}}$, and the cumulative sums

$$\sum_{r<Q} u_r \qquad \text{and} \qquad \sum_{r<Q} \left(u_r^2 + \frac{u_r}{N_r^{\text{strike}}} \right),$$

because matching one finite strike count does not establish the cumulative conditions used in Section 7.1. The real modular filter has no assigned parent policy, but the same local observations still apply. Its hit count can be compared with the protective minimum, uniform mean, and adversarial maximum from Section 6.2. Across successive heads, $D(Q)$ then shows whether the arithmetic placement remains near the random companion or accumulates damage like an informed allocator. This is a comparison of behavior, not an attribution of intent.

7 Exact-Quota Companion Processes

7.1 Exact CRT Quotas With Random Locations

The random parent model fixes two harmful copy indices per parent. A separate statistical companion instead retains the exact number of accepted strikes supplied by a chosen CRT population and randomizes only their locations. This preserves the real count while removing the arithmetic targeting information, but it does not preserve two losses per 2-gap parent or the balanced global recurrence. Exact quotas create dependence within one layer, but they do not change the one-position survival scale when allocated uniformly on the conditioned eligible population. The head result also uses persistent availability and cross-layer mixing, while the square-window result uses blind placement and a quadratic eligible supply. At filter r, let U_r contain N_r eligible values and let the CRT quota be J_r, with $0 \le J_r \le N_r - 2$. The exact-quota random parent model chooses one uniformly random size-J_r subset of U_r as its strike set. For a specified 2-gap whose two endpoints belong to U_r, both endpoints survive precisely when every strike is selected from the other $N_r - 2$ values. Therefore

$$\begin{aligned} s_r &= \frac{\binom{N_r-2}{J_r}}{\binom{N_r}{J_r}} && \text{[Uniform Exact-Quota Choice]} \\ &= \frac{(N_r - J_r)(N_r - J_r - 1)}{N_r(N_r - 1)}. && \text{[Factorial Simplification]} \end{aligned}$$

Write the strike fraction as

$$u_r := \frac{J_r}{N_r}.$$

The exact formula gives

$$\log s_r = -2u_r + O\left(u_r^2 + \frac{u_r}{N_r}\right).$$

This separates the exact finite quota from the cumulative condition needed for recurrence. Assume that along the conditioned chain to head Q,

$$\begin{aligned} \sum_{r<Q} u_r &= \log\log Q + O(1), && \text{[CRT-Rate Cumulative Quota]} \\ \sum_{r<Q} \left(u_r^2 + \frac{u_r}{N_r}\right) &= O(1). && \text{[Summable Finite-Population Error]} \end{aligned}$$

The complete-period CRT benchmark $u_r = 1/r$ satisfies these conditions. A different local CRT quota must be checked against them; preserving a numerical strike count alone does not make the conclusion automatic. The prime-head summation below uses the Prime Number Theorem in the form cited from Hardy and Wright [1], and the infinite-recurrence conclusion uses the Borel-Cantelli/Kochen–Stone mixing premise stated in Section 2.1 [2]. Multiplying the exact

without-replacement factors gives

$$\begin{aligned} P_{\text{quota}}(Q) &= \prod_{r<Q} s_r && \text{[Survive Every Filter]} \\ &= \exp\left(\sum_{r<Q} \log s_r\right) && \text{[Product To Sum]} \\ &= \exp\left(-2\sum_{r<Q} u_r + O(1)\right) && \text{[Summable Error]} \\ &\asymp \frac{C}{(\log Q)^2}. && \text{[Cumulative Quota Condition]} \end{aligned}$$

Thus the exact-quota companion has the same one-head survival order as the random parent model. It is not an independent Bernoulli filter inside one layer; it is a uniform shuffle conditioned on the exact CRT strike count. Suppose the distinguished head candidate is eligible with conditional probability at least $b_0 > 0$, uniformly for all sufficiently large prime heads, and that this availability is compatible with the quota-survival experiment. Then

$$\begin{aligned} b_0 P_{\text{quota}}(Q) \le \Pr(H_Q) &\le P_{\text{quota}}(Q), && \text{[Uniform Availability Bounds]} \\ \Pr(H_Q) &\asymp \frac{1}{(\log Q)^2}. && \text{[Quota Survival Asymptotic]} \end{aligned}$$

Consequently,

$$\begin{aligned} \sum_{Q \text{ prime}} \Pr(H_Q) &\asymp \sum_{Q \text{ prime}} \frac{1}{(\log Q)^2} \\ &= \infty. && \text{[Prime Number Theorem]} \end{aligned}$$

Under independence or an adequate cross-layer mixing condition, the second Borel-Cantelli lemma yields

$$\Pr(H_Q \text{ occurs infinitely often}) = 1. \qquad \text{[Q.E.D.]}$$

This is an almost-sure theorem, not a guarantee for every random realization. The set of realizations with only finitely many head hits has probability zero, but it is not logically empty. For square-safe windows, assume $B(Q) \asymp C_0 Q^2$ eligible starts and the same blind-placement empty-window premise used by the random parent model. Then

$$\begin{aligned} \lambda_{\text{quota}}(Q) &= B(Q) P_{\text{quota}}(Q) && \text{[Expected Surviving Starts]} \\ &\asymp C_0 \frac{Q^2}{(\log Q)^2} \longrightarrow \infty. && \text{[Quadratic Supply Dominates]} \end{aligned}$$

If

$$\Pr(W_Q \text{ is empty}) \le e^{-\lambda_{\text{quota}}(Q)},$$

the empty probabilities are summable over prime Q. The first Borel-Cantelli lemma then gives only finitely many empty square windows almost surely. This eventual-window statement is stronger than the twin-prime-style target: an unbounded sequence of successful windows, or infinitely many head hits, is already sufficient for infinitely many distinct certificates. For consecutive primes $p < q$, the real accepted-strike count in the next safe window is proved in Gap Dynamics §9.1 [4]:

$$A(p,q) = \pi\left(\left\lfloor \frac{q^2-1}{p} \right\rfloor\right) - \pi(p-1).$$

Using this local quota as $J_r = A(p, q)$ in the random-location companion is well defined, but its fractions $u_r = J_r/N_r$ must still satisfy the displayed cumulative conditions for the head proof above. The complete-period count [4] supplies the global density; neither exact count determines local placement in the real sieve.

7.2 Biased Exact Quotas and the Logarithmic Skew Frontier

The neutral exact-quota companion treats every eligible value symmetrically. We now keep the same quota J_r while making 2-gap endpoints proportionally more likely to receive a harmful strike. This asks how much positional preference the quota can carry before the almost-sure head conclusion changes. We use a group-exchangeable allocation law, assume double strikes on one endpoint pair have quadratic order, and retain the availability, placement, and mixing premises of Section 3.5 and Section 7.1. Let $E_r \subseteq U_r$ be the eligible values that are endpoints of locally relevant 2-gaps, and define

$$x_r := \frac{|E_r|}{N_r}, \qquad u_r := \frac{J_r}{N_r}.$$

Stipulate a group-exchangeable size-J_r allocation law in which every endpoint has marginal inclusion probability p_r^E, every ordinary value has marginal inclusion probability p_r^O, and the endpoint preference ratio is

$$\beta_r := \frac{p_r^E}{p_r^O} \geq 1.$$

We require that these marginals are feasible:

$$0 \leq p_r^O, p_r^E \leq 1.$$

This condition, together with the fixed-size law, is an assumption on the group-exchangeable quota construction rather than a consequence of the ratio β_r alone. The fixed quota forces the average marginal inclusion probability to equal u_r. Therefore

$$\begin{aligned} x_r p_r^E + (1 - x_r) p_r^O &= u_r && \text{[Exact Quota]} \\ p_r^O &= \frac{u_r}{1 + (\beta_r - 1)x_r} && \text{[Substitution]} \\ p_r^E &= \frac{\beta_r u_r}{1 + (\beta_r - 1)x_r}. && \text{[Simplification]} \end{aligned}$$

Define the quota-normalized preference

$$\kappa_r^{\text{eff}} := \frac{\beta_r}{1 + (\beta_r - 1)x_r}.$$

For one 2-gap, assume the probability that both endpoints are struck is $O((p_r^E)^2)$. Its destruction fraction then satisfies

$$\begin{aligned} f_r &= 2p_r^E - \Pr(\text{both endpoints are struck}) && \text{[Inclusion-Exclusion]} \\ &= 2u_r\kappa_r^{\text{eff}} + O\left((u_r\kappa_r^{\text{eff}})^2\right). && \text{[Normalized Preference]} \end{aligned}$$

To use the displayed linear term cumulatively, we additionally require $p_r^E = u_r\kappa_r^{\text{eff}} \to 0$ and

$$\sum_{r<\infty} \left(u_r\kappa_r^{\text{eff}}\right)^2 < \infty.$$

Without this error control, the per-filter expansion does not determine the asymptotic cumulative hazard. The raw weight β_r and the effective destruction skew are not identical: the exact quota must take probability away from ordinary values when it gives more probability to endpoints. For the complete-period eligible-sieve population, let V_r be its accepted-value count and T_r its 2-gap count. After filter 3, 2-gaps have disjoint endpoints, so $|E_r| = 2T_r$ and the displayed product estimate is the standard Mertens-type product asymptotic cited from Hardy and Wright [1]:

$$\begin{aligned} x_r &= \frac{2T_r}{V_r} \\ &= 2 \prod_{3 \le p < r} \frac{p-2}{p-1} \\ &= \Theta\left(\frac{1}{\log r}\right). \qquad \text{[Mertens-Type Product Estimate]} \end{aligned}$$

The prior $\Theta((\log r)^{-2})$ density would be the density among raw integers, not among the eligible values U_r used in this quota. For the illustrative sharper asymptotics

$$x_r \sim \frac{C}{\log r}, \qquad \beta_r \sim b \log r,$$

we instead obtain

$$\kappa_r^{\text{eff}} = \frac{b}{1+bC} \log r \, (1 + o(1)).$$

Thus a logarithmic raw preference changes its leading effective coefficient. A sparse relation $x_r = O((\log r)^{-2})$ may be imposed as an additional assumption for a different local target population, but it is not the complete-period eligible-sieve density. For the phase theorem, measure skew by the realized effective factor

$$\kappa_r := \frac{f_r}{2/r}.$$

This is the same quantity called w_r in the general hazard analysis. Its cumulative survival is defined once $2\kappa_r < r$. Every regime below satisfies this inequality for all sufficiently large filters; the finite prefix is absorbed into a positive constant. Thus

$$P_\kappa(Q) = \prod_{r<Q} \left(1 - \frac{2\kappa_r}{r}\right) = e^{-D_\kappa(Q)},$$

where

$$D_\kappa(Q) := \sum_{r<Q} -\log\left(1 - \frac{2\kappa_r}{r}\right).$$

If $\kappa_r = \kappa < \infty$ is fixed, then

$$\begin{aligned} D_\kappa(Q) &= 2\kappa \log\log Q + O(1) \quad \text{[Prime Harmonic Sum]} \\ P_\kappa(Q) &\asymp \frac{C_\kappa}{(\log Q)^{2\kappa}}. \qquad \text{[Exponentiation]} \end{aligned}$$

The sum of this probability over prime heads diverges for every finite κ. With persistent head availability and adequate cross-layer mixing,

every fixed finite proportional skew gives infinitely many head hits almost surely.

Thus there is no finite constant-skew maximum. The first transition appears when effective skew grows logarithmically. Set

$$\kappa_r = 1 + c \log r, \qquad c \geq 0.$$

Then

$$D_c(Q) = 2 \log\log Q + 2c \log Q + O(1) \quad \text{[Prime-Sum Asymptotics]}$$
$$P_c(Q) \asymp \frac{C_c}{Q^{2c}(\log Q)^2}. \quad \text{[Exponentiation]}$$

For prime heads, the occurrence series has the same convergence behavior as

$$\int^{\infty} \frac{dx}{x^{2c}(\log x)^3}.$$

Therefore

$$\begin{aligned} c < \frac{1}{2} &\Longrightarrow \text{infinitely many head hits almost surely, with mixing,} \\ c \geq \frac{1}{2} &\Longrightarrow \text{only finitely many head hits almost surely.} \end{aligned} \qquad \text{[Q.E.D.]}$$

In this explicitly normalized family, the equality $c = 1/2$ is on the failure side because the remaining factor $(\log Q)^{-2}$ makes the boundary prime series converge. If β_r, rather than effective κ_r, is specified at its raw equality boundary, quota normalization changes lower-order terms and the cumulative series $D_\kappa(Q)$ must be evaluated directly. The robust head-safe frontier is therefore

$$\kappa_r \leq 1 + \left(\frac{1}{2} - \varepsilon\right) \log r \quad \Longrightarrow \quad \text{head recurrence almost surely, with mixing.}$$

For square windows, quadratic supply gives the larger robust frontier

$$\kappa_r \leq 1 + (1 - \varepsilon) \log r \quad \Longrightarrow \quad \text{eventual square-window occupancy almost surely.}$$

Hence the intermediate range between approximately $(1/2) \log r$ and $\log r$ preserves square windows but not infinitely recurring head hits. The window statement is stronger than necessary: the twin-prime-style conclusion needs only infinitely many successful windows. For an irregular skew schedule, the head criterion is

$$\sum_{Q \text{ prime}} e^{-D_\kappa(Q)} = \infty,$$

together with persistent availability and adequate mixing. A pointwise skew percentage cannot replace this cumulative test.

8 Relation to the Real Sieve

For a real 2-gap $(a, a+2)$ and incoming prime r, the harmful copies are not chosen freely. They are fixed by

$$K_{a,r}^{\text{real}} = \{-aM^{-1}, -(a+2)M^{-1}\} \pmod{r}.$$

Different parents are coupled through this single arithmetic rule. The real filter has neither independent policy coins nor freely allocated protective and adversarial labels.

It nevertheless has directly measurable local destruction $\widehat{f}_r$, a relative factor $\widehat{w}_r = r\widehat{f}_r/2$, and, for a fixed cohort, observed hazard $\widehat{D}(Q)$. These are the empirical quantities of Section 3.2, not probabilities assigned to the deterministic sieve.

Assigning it an effective policy share α_r requires an additional declared companion benchmark. Assessing its positional concentration requires the separate hit count and targeting normalization from Section 6.2.

The relative factor defined in Section 3.2 provides a finite-transition destruction normalization. Section 6.2 adds the allocation question: for the same global destruction budget, how close is the realized local hit count to the uniform mean or targeted maximum?

The companion diagrams identify what a transfer theorem would need to control:

- the observed relative damage $\widehat{w}_r$ and fixed-cohort hazard $\widehat{D}(Q)$, together with a separate theorem relating them to candidate conditional hazards;
- the exact-quota fractions u_r and their cumulative deviation from $\sum_{r<Q} 1/r$;
- raw endpoint preference β_r, quota-normalized effective skew κ_r^{eff}, and cumulative skew hazard $D_\kappa(Q)$;
- any scheduled or effective policy budget $A(Q)$ used for a particular companion specialization;
- the availability and abundance premises for the chosen target; and
- a deterministic discrepancy bound comparing the real head indicators I_Q with the divergent companion reference weights ρ_Q, as formalized in Section 10.

8.1 Finite Empirical Comparison With Random

We can compare the real sieve with the random companion at two square-window scales. The first comparison counts the 2-gaps already present in each sequence's own window. For head h, let $G_{\mathrm{real}}(h)$ be the number of real 2-gap starts in $[h, h^2)$. The random companion expectation is

$$E_{\mathrm{random}}(h) = (h^2 - h)\frac{1}{2}\prod_{3\le r<h}\left(1 - \frac{2}{r}\right).$$

The $c = 1$ square-window frontier adds the logarithmic excess hazard from Section 3.5:

$$E_{c=1}(h) = E_{\mathrm{random}}(h)\prod_{7\le r<h}\left(1 - \frac{2\log r}{r-2}\right).$$

The figure compares these two expectations with the real count in every fully covered sequence window.

Across 188 heads from 3 through 1129, the mean ratio $G_{\mathrm{real}}/E_{\mathrm{random}}$ is 0.967. At the largest covered head it is 0.947: the real window contains 10,056 2-gaps, compared with a random expectation of 10,616.

The corresponding $c = 1$ frontier expectation is only 0.0845. Over this finite range, the real square-window population follows the random scale and remains far above the square-window failure frontier. The figure is generated by the per-sequence frontier calculation from the per-sequence survivor data.

The second comparison isolates one transition. For consecutive primes $p < q$, let G_p be the pre-filter 2-gap population in $[q, q^2)$ and let H_p be the number destroyed when filter p is installed. The observed fraction and its relative factor are

$$f_p^{\mathrm{real}} := \frac{H_p}{G_p}, \qquad w_p^{\mathrm{real}} := \frac{p f_p^{\mathrm{real}}}{2}.$$

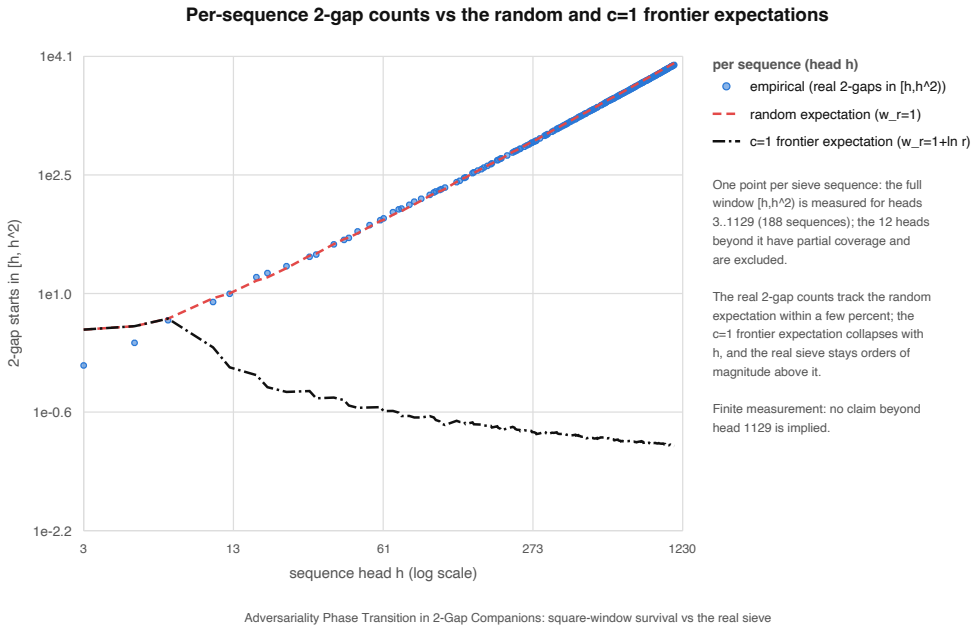


Figure 2: Real per-sequence square-window 2-gap counts compared with the random expectation and the c=1 square-window frontier

The chart compares f_p^{real} with the random rate $2/p$ and the $c = 1$ square-window boundary $2(1 + \log p)/p$.

Among 187 distinct measured transitions from $p = 3$ through $p = 19{,}429$, 186 lie below the random rate and the remaining transition, $p = 3$, equals it. Ninety-five transitions destroy no 2-gap in the measured window.

From $p \geq 1000$, the largest observed relative factor is

$$w_p^{\text{real}} = 0.0523.$$

Thus the measured real transition is not slightly more destructive than the random filter; on these windows it is substantially less destructive. The figure is generated by the per-transition frontier calculation from the dense and sparse transition data. Zero-destruction transitions are displayed on the chart's 10^{-7} floor so that they remain visible on a logarithmic axis.

The two findings are compatible. The first chart measures the population left after all earlier filters; the second isolates the next filter acting on a new window. Neither dataset follows one fixed cohort through every filter below a single head, so their changing-window fractions cannot be multiplied into one cumulative hazard.

The complete modular cycle supplies an exact reference in two complementary views. The first is local to one filter: it asks what fraction of the expanded cyclic population that filter destroys. The second is cumulative: it asks how much of a normalized starting population remains after those one-filter fractions are compounded. The two figures below deliberately mirror these two steps of the calculation.

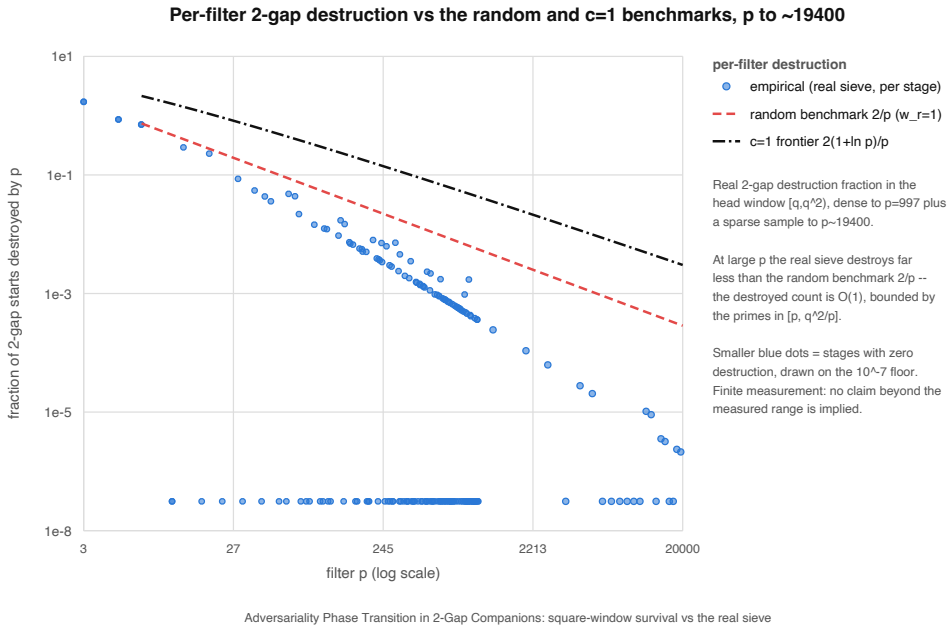


Figure 3: Real per-transition square-window 2-gap destruction compared with the random rate and the c=1 square-window boundary

Exact Per-Filter Destruction

If T old cyclic 2-gaps are expanded through a new prime r, there are rT copies and exactly two harmful copy indices per parent. Consequently,

$$\begin{aligned} H_r^{\text{cycle}} &= 2T && \text{[Two Harmful Copy Classes]}, \\ f_r^{\text{cycle}} &= \frac{H_r^{\text{cycle}}}{rT} = \frac{2}{r} && \text{[Substitution]}, \\ w_r^{\text{cycle}} &= 1, && \text{[By Definition]}, \\ D_{\text{cycle}}(R) - D_{\text{random}}(R) &= 0. && \text{[Termwise Equality; Q.E.D.]} \end{aligned}$$

Thus the full-cycle destruction fraction equals the neutral benchmark as a count identity. This does not say that the harmful positions are independently random.

The figure below makes the identity visible alongside the $c = 1$ reference on the valid range $29 \le r \le 251$. Its simplicity is the point: the overlap between the exact-cycle and neutral curves is the graphical form of the algebra above, while the separated $c = 1$ curve shows the scale of the hypothetical logarithmic worsening.

The chart is strong precisely because the reader can verify the stated relationship without additional interpretation. The figure is generated by the full-cycle destruction calculation. The underlying two-class result is proved in Gap Dynamics §6.1 [4].

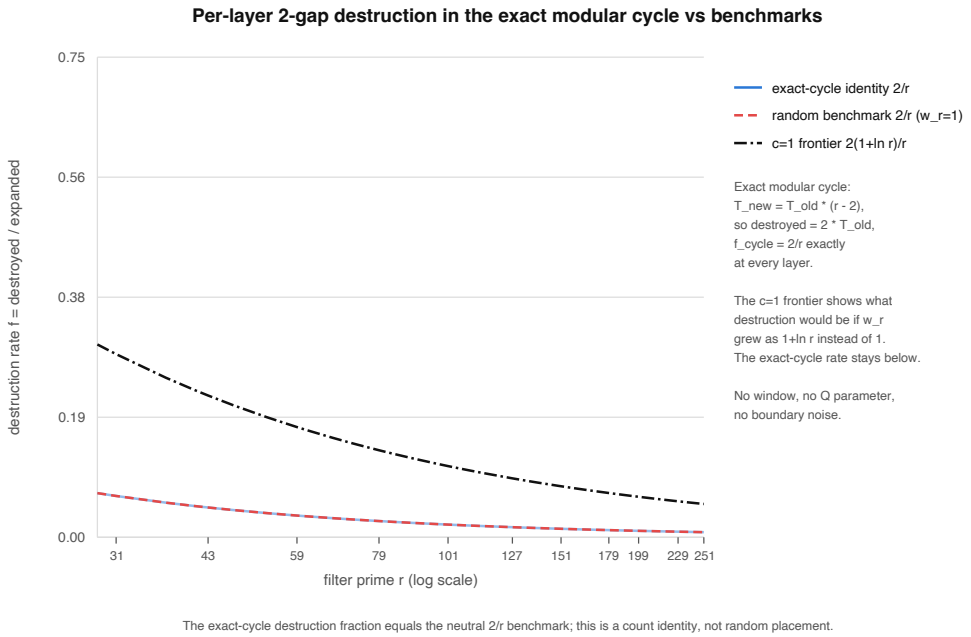


Figure 4: Exact full-cycle 2-gap destruction fraction compared with the neutral rate and the c=1 reference

Cumulative Survival Consequence

The first diagram makes the proved one-filter equality visible. Compounding those same factors gives the normalized complete-cycle survival reference

$$P_{\text{cycle}}(29, R) = \prod_{29 \le p \le R} \left(1 - \frac{2}{p}\right),$$
$$P_{c=1}(29, R) = \prod_{29 \le p \le R} \left(1 - \frac{2(1 + \log p)}{p}\right).$$

At $R = 251$, these normalized products are 0.3733 and 0.003676, respectively, a ratio of about 102.

The second figure therefore shows the cumulative consequence that the first figure cannot show by itself: the repeated per-filter separation from the $c = 1$ schedule compounds into a separation of more than two orders of magnitude over the plotted range.

The anchor 29 is the first plotted prime and keeps every $c = 1$ factor in $(0, 1)$. Changing a finite anchor changes the normalizing constants, not the asymptotic exponents. This is a reference comparison, not evidence of head recurrence. The figure is generated by the full-cycle survival calculation.

Read together, the two full-cycle figures give a direct progression from the one-filter identity to its cumulative effect.

A separate fixed-cohort experiment then asks the next question: what remains true when a finite window cuts through partial cycles? It follows every 2-gap start initially present in $[Q, Q^2)$ through all filters $r < Q$.

Exact set comparisons at $Q = 17$ and $Q = 101$ confirm that this explicit cohort agrees layer by

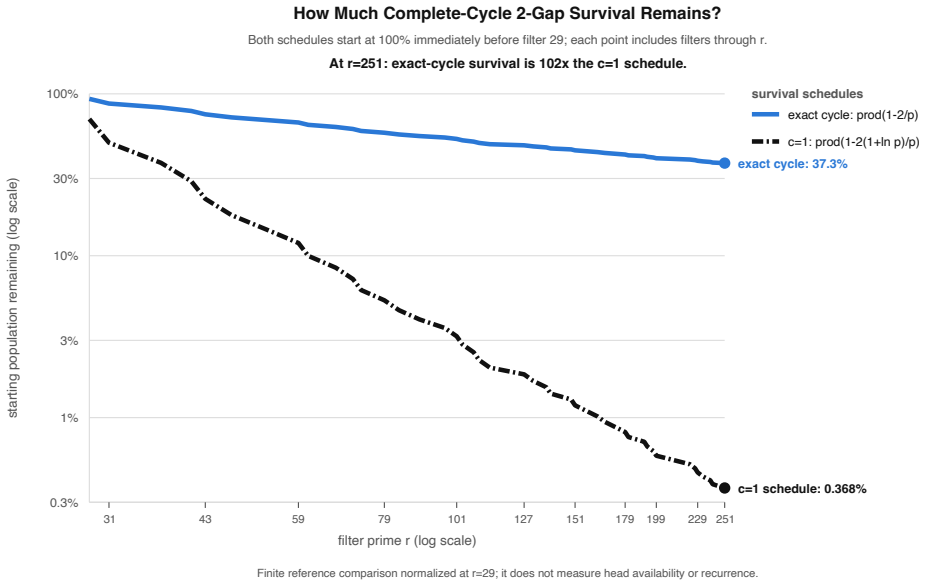


Figure 5: Normalized full-cycle 2-gap survival under the exact per-filter law compared with the c=1 schedule

layer with the maintained Reading A lineage. For

$$c_{\mathrm{eff}}(r) := \frac{\widehat{D}_{\mathrm{real}}(r) - D_{\mathrm{random}}(r)}{2\log r},$$

the four runs $Q \in \{17, 101, 251, 503\}$ give signed values between -0.0353 and 0.00908. The largest positive value is 0.00907 at $Q = 251$; among the two larger runs, all absolute values are at most 0.00908.

Because the complete-cycle excess is exactly zero, the deviations arise when the fixed interval cuts partial cycles. Their size and placement remain finite empirical facts; partial-cycle localization can contain the arithmetic difficulty that the complete-cycle identity does not see. The figure is generated by the fixed-lineage hazard calculation from the dedicated fixed-cohort CSVs. Its value is a robustness check: even in non-aligned finite windows the boundary deviations are small relative to the $c = 1/2$ and $c = 1$ comparison scales. None of these measurements follows the distinguished pair at the head. Transferring the head phase diagram still requires an arithmetic theorem that controls coherent CRT-coupled placement together with persistent availability and cross-layer dependence.

9 Limitations

The balanced companions track descendants of existing 2-gaps rather than a coherent randomized sequence of integers. They do not model non-2 gaps, gap mergers, or shared endpoint effects.

Their adversary is stronger than the real filter because it may select harmful copies separately for each parent. The protective parent policy is also an oracle: it knows the current target and moves both deletions elsewhere. Neither endpoint is a description of the real filter.

The adversarial/random window theorem assumes spatial uniformity. The adversarial/protective window theorem instead assumes a quadratic protective supply $B(Q)$. Its head theorem assumes an eligible protective-parent head lineage with availability bounded below. None of these premises follows from the exact-two choice alone.

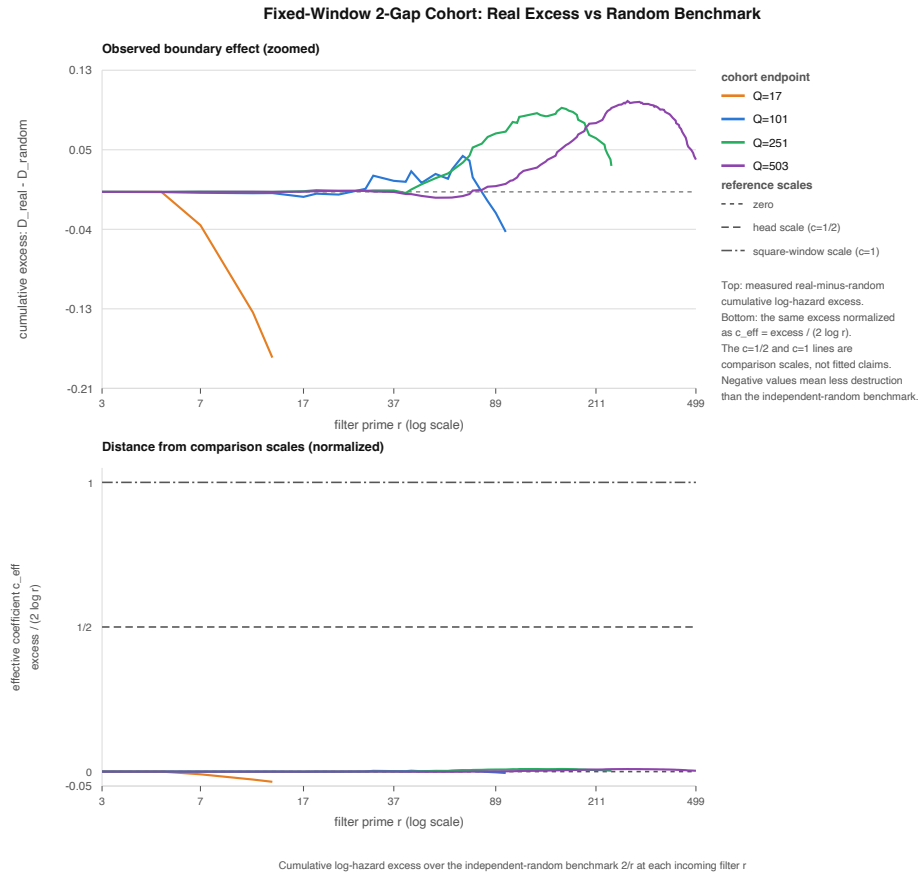


Figure 6: Cumulative hazard of fixed-window 2-gap cohorts for four finite Q values

The exact-CRT-quota/random-location theorem preserves strike counts but discards their deterministic arithmetic locations. Its recurrence conclusion requires the cumulative quota conditions in Section 7.1, persistent head availability, and cross-layer mixing.

The exact local formula $A(p, q)$ does not establish those premises merely by being exact.

The biased-quota theorem additionally assumes a group-exchangeable allocation law and quadratic-order double hits on one endpoint pair. Its raw weight β_r is not the same as effective destruction skew κ_r: exact-quota normalization can change its leading coefficient for the eligible-sieve population. A raw equality case must therefore be decided through $D_\kappa(Q)$ and the quota feasibility conditions.

The use of expected population also has a strict boundary. An expectation tending to infinity does not alone prove nonempty windows; the summable empty-window bound supplies that step under uniform placement. Likewise, a divergent head-event series does not alone prove infinite recurrence; adequate cross-layer independence or mixing is required. Exact quotas, whole-filter coins, block balance, delayed information, and noisy ranking create different dependencies and cannot inherit one another's almost-sure conclusions merely because they share the same marginal budget or destruction rate.

The empirical comparison in Section 8.1 is finite. Two datasets use a different square window at each measured stage; the fixed-cohort dataset instead follows all initial 2-gap starts in one $[Q, Q^2)$ window.

The latter yields a coherent window hazard, but its deviation from the exact full-cycle rate

comes from partial-cycle localization and the cohort is not the distinguished head pair. None of the data therefore substitutes for persistent availability, cross-layer mixing, or an arithmetic proof below the $c = 1/2$ head frontier.

We have therefore proved a phase diagram for the defined companions. The remaining question is whether the real sieve occupies one of these regimes.

10 Conclusion

Balanced 2-gap companions make global persistence deliberately uninformative: every parent always leaves $r-2$ children, so the complete-period population grows under protective, random, adversarial, and mixed selection alike. The local distinction is carried by conditional destruction relative to random, its cumulative hazard, and allocation. Finite real-sieve data instead supply observed cohort fractions, which require a separate transfer theorem. The proof begins with the exact global recurrence and then replaces population counting by cumulative local hazard:

$$\begin{aligned} N_{k+1} &= (r_k - 2)N_k, \\ w_r &= \frac{r f_r}{2}, \\ D(Q) &= \sum_{r<Q} -\log(1 - f_r), \\ P(Q) &= e^{-D(Q)}. \end{aligned}$$

The main answer is that there is no maximum finite constant factor worse than random. If $w_r = w < \infty$, local survival decays only as $(\log Q)^{-2w}$; quadratic square-window supply still dominates, and the head probability series still diverges. Under the stated spatial and mixing premises, square windows are eventually nonempty and head 2-gaps recur infinitely often for every fixed finite w. The first nontrivial boundary occurs when worsening grows with the filter. Under Section 7.1's cumulative quota/error conditions and the relevant spatial premises, the neutral exact-quota companion recovers the random survival scale. The biased exact-quota companion recovers the following general-hazard frontiers when its conditional effective skew follows the displayed schedule and its placement, availability, and mixing premises hold:

$$\begin{aligned} w_r = 1 + c\log r, \quad c < 1 &\Longrightarrow \text{eventual square-window occupancy}, \\ w_r = 1 + c\log r, \quad c < \frac{1}{2} &\Longrightarrow \text{infinitely recurring head 2-gaps, with mixing}. \end{aligned}$$

The allocation theorem explains why a percentage alone cannot locate a process in this phase diagram. Uniform, protective, and targeted allocation can apply the same adversarial budget and produce different local damage. The quantity that enters the theorem is therefore the conditional local hazard, not the policy label by itself. The complete-period comparison is exact: every new filter destroys the fraction $2/r$ of expanded cyclic 2-gaps, so its cumulative excess over the neutral benchmark is zero. The two full-cycle diagrams expose the two parts of this statement separately. The destruction diagram makes the per-filter identity immediate; the survival diagram shows what compounding that identity does and how strongly it separates from the $c = 1$ schedule. Neither diagram is weakened by being a direct rendering of the calculation: their value is that the local equality and cumulative consequence can each be checked visually in the representation best suited to it. The finite square-window measurements then describe the remaining localization question. Through head 1129, the observed population has mean ratio 0.967 to the random expectation, and

the measured one-step window rates through filter 19,429 are at or below $2/r$. Fixed cohorts for $Q = 17, 101, 251, 503$ have signed effective coefficients between -0.0353 and 0.00908; these are finite partial-cycle measurements around the exact cycle law. The measurements remain far from the $c = 1$ window-failure scale, but they do not locate the distinguished head pair relative to the $c = 1/2$ recurrence frontier. The real sieve question is now expressible as a cumulative comparison rather than a vague claim of random or adversarial behavior: measure the observed cohort destruction $\widehat{f}_r$ generated by the CRT-coupled harmful indices, normalize it to $\widehat{w}_r$, and establish whether those observations control the relevant candidate hazards and placement near the head. Here is one explicit sufficient deterministic counting criterion. Let $I_Q \in \{0, 1\}$ indicate that the real sieve has a 2-gap at prime head Q. Let $\rho_Q > 0$ be the companion reference weight obtained from the below-frontier cumulative hazard and the stated availability bound, and define

$$R(X) := \sum_{\substack{Q \le X \\ Q \text{ prime}}} \rho_Q.$$

Assume

$$R(X) \longrightarrow \infty, \qquad \text{[Divergent Reference Mass]}$$

$$\sum_{\substack{Q \le X \\ Q \text{ prime}}} (I_Q - \rho_Q) = o(R(X)). \qquad \text{[Deterministic Discrepancy Bound]}$$

Then

$$\sum_{\substack{Q \le X \\ Q \text{ prime}}} I_Q = R(X) + o(R(X)) \longrightarrow \infty. \qquad \text{[Q.E.D.]}$$

This is an unproved sufficient condition for the real CRT sieve. It is stronger than recurrence itself because it asserts an asymptotic counting formula. The reference mass must diverge, rather than every individual filter satisfying one pointwise bound. "Infinitely often" does not mean that every sufficiently large head is a 2-gap. Proving this condition would prove infinitely many head 2-gaps and therefore the twin-prime conjecture; it does not follow from the hazard estimates in this article.

11 Future Work

The companion theorems reduce the real-sieve question to measurable arithmetic inputs. The first direction is to estimate the real local hazard $D_\kappa(Q)$ from CRT-coupled filter locations and compare it with the head frontier without assigning hostile intent to individual filters. The second is to replace the stochastic mixing premise with a deterministic decorrelation or discrepancy theorem strong enough to transfer a divergent head-event sum to the real sequence. The third is to test exact-quota, delayed, block-balanced, and noisy-ranking companions on the same transitions, reporting both raw endpoint preference and quota-normalized effective skew. These directions are deliberately separate. A finite experiment can identify which companion resembles the observed sieve, but it cannot establish the infinite transfer theorem. Conversely, a discrepancy theorem must control where the real CRT strikes land, not only how many strikes or global 2-gaps exist.

References

[1] G. H. Hardy and E. M. Wright. *An Introduction to the Theory of Numbers*. Oxford University Press, 6th edition, 2008. https://doi.org/10.1093/oso/9780199219858.001.0001.

[2] Simon Kochen and Charles Stone. A note on the Borel–Cantelli lemma. *Illinois Journal of Mathematics*, 8(2):248–251, 1964. https://doi.org/10.1215/ijm/1256059668.

[3] Thiago Henrique Ramos da Mata. Formal Verification of Sieve Sequence Stages and Their Transitions. https://doi.org/10.5281/zenodo.22955782, 2026. Zenodo:10.5281/zenodo.22955782.

[4] Thiago Henrique Ramos da Mata. Structural Properties and Signed Boundaries of 2-Gaps in Sieve Sequences. https://doi.org/10.5281/zenodo.22955786, 2026. Zenodo:10.5281/zenodo.22955786.

A Appendix A. Selected Companion Proof Records

The body develops every result in its mathematical context. This appendix collects six core companion-process proof records with their premises and conclusions; it is a selected reference, not a complete catalog of the body.

A.1 Global Persistence Is Independent of Allocation

Once the initial population is nonzero, this result is unconditional with respect to allocation inside every balanced companion. Let $N_k = |\mathcal{G}_k|$ be the complete-period 2-gap population before installing prime r_k, with $N_0 > 0$. Every parent produces r_k copies and loses exactly two, regardless of where those two removals occur. Therefore

$$\begin{aligned} N_{k+1} &= \sum_{g \in \mathcal{G}_k} (r_k - 2) && \text{[Exactly Two Copies Removed Per Parent]} \\ &= (r_k - 2) N_k. && \text{[Simplification]} \end{aligned}$$

Iterating the recurrence gives

$$\begin{aligned} N_k &= N_0 \prod_{i<k} (r_i - 2) && \text{[Iteration]} \\ &> 0 && [N_0 > 0;\ r_i \geq 5] \\ &\longrightarrow \infty. && \text{[Every Factor Is At Least 3]} \end{aligned} \qquad \text{[Q.E.D.]}$$

Thus allocation may eliminate 2-gaps from the head or a tracked window, but it cannot exhaust the complete-period population while the exact-two removal rule is preserved.

A.2 Cumulative Local-Hazard Law

Follow one initially eligible candidate through successive filters. Let f_r be its conditional destruction probability given earlier survival and initial eligibility, and assume $0 \leq f_r < 1$. Define

$$w_r := \frac{r f_r}{2}, \qquad D(Q) := \sum_{r<Q} -\log(1 - f_r).$$

The conditional survival factors multiply by the probability chain rule, without requiring independent filter events:

$$\begin{aligned} P(Q) &= \prod_{r<Q}(1-f_r) && \text{[Survive Every Filter]} \\ &= \exp\left(\sum_{r<Q}\log(1-f_r)\right) && \text{[Product To Sum]} \\ &= e^{-D(Q)}. && \text{[Definition Of } D(Q)\text{]} \end{aligned} \qquad \text{[Q.E.D.]}$$

For the random benchmark $w_r = 1$, the prime harmonic sum gives the standard prime-harmonic estimate cited from Hardy and Wright [1]:

$$\begin{aligned} D_{\text{random}}(Q) &= \sum_{r<Q} -\log\left(1-\frac{2}{r}\right) \\ &= 2\log\log Q + O(1), \\ P_{\text{random}}(Q) &\asymp \frac{C}{(\log Q)^2}. \end{aligned}$$

This identity determines the candidate's survival probability from its conditional hazards. For a fixed nested finite cohort, the analogous product of observed fractions gives its final-to-initial size ratio. Neither version supplies window abundance, head availability, or cross-layer mixing. A lethal step eliminates the tracked candidate or cohort, not every later candidate.

A.3 Every Fixed Finite Worsening Factor Survives

Let $w \geq 0$ be fixed and suppose $f_r = 2w/r$ for all sufficiently large filters. A finite prefix changes only the positive leading constant. Since the quadratic Taylor remainder is summable over primes, and using the prime harmonic estimate from Hardy and Wright [1], Appendix A.2 gives

$$\begin{aligned} D_w(Q) &= \sum_{r<Q} -\log\left(1-\frac{2w}{r}\right) && \text{[Definition Of } D(Q)\text{]} \\ &= 2w\sum_{r<Q}\frac{1}{r} + O(1) && \text{[Taylor Expansion; Summable Remainder]} \\ &= 2w\log\log Q + O(1). && \text{[Prime Harmonic Sum]} \end{aligned}$$

Therefore

$$P_w(Q) \asymp \frac{C_w}{(\log Q)^{2w}}.$$

Assume first that a square window supplies $B(Q) \asymp C_0 Q^2$ eligible lineages and that their placement satisfies the blind empty-window bound. Its expected surviving population is

$$\lambda_w(Q) \asymp C_0 \frac{Q^2}{(\log Q)^{2w}} \longrightarrow \infty.$$

This polynomial growth makes the empty-window probabilities summable, so only finitely many square windows are empty almost surely. For a distinguished head whose baseline availability is bounded below,

$$\Pr(H_Q) \asymp \frac{C_w}{(\log Q)^{2w}}.$$

The sum over prime heads diverges for every finite w. With adequate cross-layer mixing, head 2-gaps therefore recur infinitely often almost surely. Hence there is no finite constant-factor maximum worse than random. ■

A.4 Logarithmically Growing Worsening Has Two Thresholds

Retain the supply, availability, placement, and mixing premises of Appendix A.3, and set

$$w_r = 1 + c \log r, \qquad c \geq 0.$$

Then $f_r = 2/r + 2c \log r/r$. Prime summation and the summable Taylor remainder give the following estimate, again using the Prime Number Theorem and partial summation consequences cited from Hardy and Wright [1]:

$$\begin{aligned} D_c(Q) &= \sum_{r<Q} -\log(1 - f_r) && \text{[Definition Of } D(Q)\text{]} \\ &= 2 \sum_{r<Q} \frac{1}{r} + 2c \sum_{r<Q} \frac{\log r}{r} + O(1) && \text{[Substitution; Summable Remainder]} \\ &= 2 \log\log Q + 2c \log Q + O(1). && \text{[Prime-Sum Asymptotics]} \end{aligned}$$

Consequently,

$$P_c(Q) \asymp \frac{C_c}{Q^{2c}(\log Q)^2}.$$

For a quadratic square-window supply,

$$\lambda_c(Q) \asymp C_0 \frac{Q^{2-2c}}{(\log Q)^2}.$$

Thus $c < 1$ gives eventual square-window occupancy almost surely under the blind-placement premise, while $c \geq 1$ makes the expected population tend to zero. For the head, the prime occurrence series has the same convergence behavior as

$$\int^{\infty} \frac{dx}{x^{2c}(\log x)^3}.$$

The integral diverges for $c < 1/2$ and converges for $c \geq 1/2$. Therefore

$$\begin{aligned} c < 1 &\implies \text{eventual square-window occupancy almost surely,} \\ c < \frac{1}{2} &\implies \text{infinitely recurring head 2-gaps almost surely, with mixing.} \end{aligned} \qquad \text{[Q.E.D.]}$$

The intermediate range $1/2 \leq c < 1$ preserves square windows but not infinitely recurring head events. Irregular schedules must be evaluated by the cumulative hazard $D(Q)$ rather than by isolated pointwise values.

A.5 Adversarial/Random Square-Window Boundary

At filter r, let a parent be adversarial with share α_r and random otherwise. For one locally relevant lineage,

$$1 - f_r = (1 - \alpha_r)\left(1 - \frac{2}{r}\right).$$

Define the cumulative adversarial-share hazard

$$A(Q) := \sum_{r<Q} -\log(1-\alpha_r).$$

The random survival density contributes $(\log Q)^{-2}$, while the repeated adversarial share contributes $e^{-A(Q)}$. If the square-safe window has length $L_Q \asymp Q^2$ and surviving starts obey the spatial-uniformity model, then

$$\begin{aligned} \lambda_Q^{\mathrm{mix}} &= L_Q \delta_Q^{\mathrm{mix}} && \text{[Expected Uniform Occupancy]} \\ &\asymp C \frac{Q^2}{(\log Q)^2} e^{-A(Q)}. && \text{[Cumulative Survival]} \end{aligned}$$

Taking logarithms gives

$$\log \lambda_Q^{\mathrm{mix}} = 2\log Q - 2\log\log Q - A(Q) + O(1).$$

Therefore, for every fixed $\varepsilon > 0$,

$$\begin{aligned} A(Q) \le (2-\varepsilon)\log Q &\Longrightarrow \lambda_Q^{\mathrm{mix}} \longrightarrow \infty, && \text{[Subcritical Budget]} \\ A(Q) \ge (2+\varepsilon)\log Q &\Longrightarrow \lambda_Q^{\mathrm{mix}} \longrightarrow 0. && \text{[Supercritical Budget]} \end{aligned}$$

At the exact boundary, the term $-2\log\log Q$ must be retained. Under uniform placement,

$$\Pr(X_Q = 0) \le e^{-\lambda_Q^{\mathrm{mix}}}.$$

Hence the summability condition

$$\sum_{Q \text{ prime}} e^{-\lambda_Q^{\mathrm{mix}}} < \infty$$

implies that only finitely many square windows are empty almost surely. A convenient sufficient condition is $\lambda_Q^{\mathrm{mix}} \ge (1+\varepsilon)\log Q$ for all sufficiently large Q. ■

A.6 Local Survivor Allocation Range

Consider a target window shorter than the old period, so each parent contributes at most one target child. Let N be the number of parents, let R be the set of L relevant parents, and let $\mathcal{A}$ be the size-K set receiving adversarial treatment. Every parent outside $\mathcal{A}$ is protective and preserves its target child. The number of target children destroyed and surviving are

$$H = |\mathcal{A} \cap R|, \qquad S = L - H.$$

The intersection cannot exceed either set, and at most $N - L$ adversarial labels can be placed outside R. Hence

$$\begin{aligned} H &\le \min(K, L), && \text{[Intersection Upper Bound]} \\ H &\ge \max(0, K - (N-L)). && \text{[Irrelevant-Parent Capacity]} \end{aligned}$$

Substitution into $S = L - H$ gives the sharp survivor range

$$\max(0, L-K) \le S \le \min(L, N-K).$$

Both endpoints are attainable. A target-aware allocator selects relevant parents first, while a protective allocator assigns the adversarial labels to irrelevant parents first:

$$S_{\text{targeted}} = \max(0, L - K),$$
$$S_{\text{protective}} = \min(L, N - K).$$

If $\mathcal{A}$ is instead a uniformly random size-K subset, then

$$H \sim \text{Hypergeometric}(N, L, K),$$

so

$$\mathbb{E}[H] = \frac{KL}{N},$$
$$\mathbb{E}[S] = L\left(1 - \frac{K}{N}\right).$$

When $K \geq L$, uniform allocation clears the target with probability

$$\Pr(S = 0) = \frac{\binom{N-L}{K-L}}{\binom{N}{K}} = \frac{\binom{K}{L}}{\binom{N}{L}}.$$

Writing $\alpha = K/N$, the targeted endpoint becomes

$$S_{\text{targeted}} = 0 \iff K \geq L \iff \alpha \geq \frac{L}{N}. \qquad \text{[Q.E.D.]}$$

At the head, $L = 1$, so one correctly placed adversarial label destroys the current candidate. If $L/N \longrightarrow 0$ in a tracked window, every fixed $\alpha > 0$ eventually has enough capacity to clear that window. Appendix A.1 still applies: the targeted parents leave $r-2$ descendants outside the window, so complete-period growth continues.